

Introduction to Computer Vision

Kaveh Fathian

Assistant Professor

Computer Science Department

Colorado School of Mines


Lecture 17

Face Detection in Images



- Slide windows of different sizes across image
- At each location, match window to face model


Face Detection Framework




- **Features:** which features represent faces well?
- **Classifier:** How to construct a face model & efficiently classify features as face or not?

What Are Good Features

- **Interest points:** edges, corners, SIFT, HOG?




- **Facial components:** templates?



Characteristics of Good Features

- Discriminate face & non-face



- Extremely fast to compute
Need to evaluate millions of windows in an image

Haar Features



Input Image



$$\left[\begin{array}{c} \square \blacksquare H_A \\ \blacksquare \square H_B \\ \square \blacksquare H_C \\ \blacksquare \square H_D \\ \vdots \end{array} \right]$$

Haar Filters

$$= \left[\begin{array}{c} V_A[i,j] \\ V_B[i,j] \\ V_C[i,j] \\ V_D[i,j] \\ \vdots \end{array} \right]$$

Harr Features
 $f[i,j]$

Computing a Haar Feature




White = 1
Black = -1

Response to filter H_A at location (i, j) :

$$V_A[i, j] = \sum_m \sum_n I[m - i, n - j] H_A[m, n]$$

$$V_A[i, j] = \sum (\text{pixel intensities in white area}) - \sum (\text{pixel intensities in black area})$$

Discriminative Ability of Haar Features



- Haar features are sensitive to the *directionality* of patterns


Discriminative Ability of Haar Features

- Using only 2 Haar features on a 24x24 pixel window:




- Can yield 100% face detection rate, and 50% false positive rate

Detecting Faces at Different Scales




- Compute Haar features at different scales to detect faces of different sizes

Classifier for Face Detection




- Given features for a window, how to decide whether it contains a face or not?



- Haar feature f (a vector) at a pixel is a point in n-D space, $f \in \mathbb{R}^n$
- Can use any classifier engine: nearest neighbor, SVM, neural network, etc.

Haar Feature: Computation Cost



Value = $\sum(\text{pixel intensities in white area}) - \sum(\text{pixel intensities in black area})$

Computation cost = $(N \times M - 1)$ additions per pixel, per filter, per scale


Can we do better?

Requirements for a Face Detector



- Sliding window = tens of thousands of location/scale evaluations
 - One megapixel image has $\sim 10^6$ pixels, and a comparable number of candidate face locations
- Faces are rare: 0-10 per image
 - For computational efficiency, spend as little time as possible on the non-face windows.
- For 1 Mpix image, to avoid having a false positive in every image, our false positive rate must be less than 10^{-6}

Viola-Jones Detector




- A seminal approach to real-time object detection
 - P.Viola and M. Jones. "Rapid object detection using a boosted cascade of simple features." CVPR 2001. ~30K citations!
 - P.Viola and M. Jones. "Robust real-time face detection." IJCV 57(2), 2004.
- Training is slow, but detection is very fast

Key Ideas:

- **Haar features**
- **Integral images** for fast feature computation
- **Boosting** for feature selection
- Attentional **cascade** for fast non-face window rejection


Integral Images for Fast Feature Evaluation

- The *integral image* computes a value at each pixel (x,y) that is the sum of *all* pixel values above and to the left of (x,y) , inclusive.
- This can quickly be computed in one pass through the image.
- ‘Summed area table’




$$I_{\Sigma}(x, y) = \sum_{\substack{x' \leq x \\ y' \leq y}} i(x', y')$$

Computing the Integral Image



Computing the Integral Image




- Cumulative row sum: $s(x, y) = s(x-1, y) + i(x, y)$
- Integral image: $ii(x, y) = ii(x, y-1) + s(x, y)$

Python: `ii = np.cumsum(i)`


Computing Sum within a Rectangle

- Let R be a desired rectangle
- A,B,C,D are the values of the integral image at the corners of R




Computing Sum within a Rectangle

- Let R be a desired rectangle
- A, B, C, D are the values of the integral image at the corners of R




Computing Sum within a Rectangle

- Let R be a desired rectangle
- A, B, C, D are the values of the integral image at the corners of R




Computing Sum within a Rectangle

- Let R be a desired rectangle
- A, B, C, D are the values of the integral image at the corners of R



Computing Sum within a Rectangle

- Let R be a desired rectangle
- A, B, C, D are the values of the integral image at the corners of R




Computing Sum within a Rectangle





- Let R be a desired rectangle
- A,B,C,D are the values of the integral image at the corners of R
- The sum of original image values within the rectangle can be computed as:

$$\text{sum} = A - B - C + D$$

**Only 3 additions are required
for any size of rectangle!**



Computing a Haar Feature




$$V_A[i, j] = \sum(\text{pixel intensities in white area}) - \sum(\text{pixel intensities in black area})$$

Computing a Haar Feature




Integral Image





Number of Features

- Individual Haar features are ‘weak classifiers’
 - Jargon: ‘feature’ and ‘classifier’ are used interchangeably. Also, ‘learner’ and ‘filter’.
- But, what if we combine thousands of them?



Number of Features




- For a 24x24 detection region, the number of possible rectangle features is $\sim 160,000$!
- At test time, it is impractical to evaluate the entire feature set
- Can we learn a ‘strong classifier’ using just a small **subset** of all possible features?

Boosting for Feature Selection

Initially, weight each training example equally

Weight = size of point




Boosting for Feature Selection

In each boosting round:

Find the weak classifier that achieves the lowest *weighted* training error.

Raise the weights of training examples misclassified by current weak classifier.

Round 1:




Boosting Illustration

In each boosting round:

Find the weak classifier
that achieves the lowest
weighted training error.

Raise the weights of
training examples
misclassified by current
weak classifier.

Round 1:




Boosting Illustration

In each boosting round:

Find the weak classifier
that achieves the lowest
weighted training error.

Raise the weights of
training examples
misclassified by current
weak classifier.

Round 2:




Boosting Illustration

In each boosting round:

Find the weak classifier
that achieves the lowest
weighted training error.

Raise the weights of
training examples
misclassified by current
weak classifier.

Round 2:



Weights Increased


Boosting Illustration

In each boosting round:

Find the weak classifier
that achieves the lowest
weighted training error.

Raise the weights of
training examples
misclassified by current
weak classifier.

Round 3:




Weak Classifier 3

Boosting Illustration

- Compute final classifier as linear combination of all weak classifier.
- Weight of each classifier is directly proportional to its accuracy.
- Exact formulas for re-weighting and combining weak learners depend on the boosting scheme (e.g., AdaBoost)

Round 3:



AdaBoost: Y. Freund and R. Schapire, A short introduction to boosting, Journal of Japanese Society for Artificial Intelligence, 14(5):771-780, September, 1999.

Feature Selection with Boosting

- Create a large pool of features (160K)
- Select discriminative features that work well together

$$h(\mathbf{x}) = \text{sign} \left(\sum_{j=1}^M \alpha_j h_j(\mathbf{x}) \right)$$

Final strong learner
 window
 Weak learner
 Learner weight

- “Weak learner” = feature + threshold + ‘polarity’

$$h_j(\mathbf{x}) = \begin{cases} -s_j & \text{if } f_j < \theta_j \\ s_j & \text{otherwise} \end{cases}$$

value of rectangle feature
 threshold
 ‘polarity’ = black or white region flip $s_j \in \pm 1$

- Choose weak learner that minimizes error on weighted training set, then reweight

1. Input the positive and negative training examples along with their labels $\{(\mathbf{x}_i, y_i)\}$, where $y_i = 1$ for positive (face) examples and $y_i = -1$ for negative examples.

2. Initialize all the weights to $w_{i,1} \leftarrow \frac{1}{N}$, where N is the number of training examples. (Viola and Jones (2004) use a separate N_1 and N_2 for positive and negative examples.)

3. For each training stage $j = 1 \dots M$:
- Renormalize the weights so that they sum up to 1 (divide them by their sum).
 - Select the best classifier $h_j(\mathbf{x}; f_j, \theta_j, s_j)$ by finding the one that minimizes the weighted classification error

$$e_j = \sum_{i=0}^{N-1} w_{i,j} e_{i,j}, \quad (14.3)$$

$$e_{i,j} = 1 - \delta(y_i, h_j(\mathbf{x}_i; f_j, \theta_j, s_j)). \quad (14.4)$$

For any given f_j function, the optimal values of (θ_j, s_j) can be found in linear time using a variant of weighted median computation (Exercise 14.2).

- (c) Compute the modified error rate β_j and classifier weight α_j ,

$$\beta_j = \frac{e_j}{1 - e_j} \quad \text{and} \quad \alpha_j = -\log \beta_j. \quad (14.5)$$

- (d) Update the weights according to the classification errors $e_{i,j}$

$$w_{i,j+1} \leftarrow w_{i,j} \beta_j^{1-e_{i,j}}, \quad (14.6)$$

i.e., downweight the training samples that were correctly classified in proportion to the overall classification error.


4. Set the final classifier to

$$h(\mathbf{x}) = \text{sign} \left[\sum_{j=0}^{m-1} \alpha_j h_j(\mathbf{x}) \right]. \quad (14.7)$$

AdaBoost pseudocode Szeliski P.665


Boosting for Face Detection

- First two features selected by boosting:





- This feature combination can yield 100% recall and 50% false positive rate

Attention Cascade for Fast Detection




- Fast classifiers early in cascade which reject many negative examples but detect almost all positive examples
- Slow classifiers later, but most examples don't get there


Attentional Cascade



- Chain classifiers that are progressively more complex
- Minimize **false positive** rates at each stage, not absolute error



Training the Cascade



- Set target detection and false positive rates for each stage
- Keep adding features to the current stage until its target rates have been met
 - Need to lower boosting threshold to maximize detection (as opposed to minimizing total classification error)
 - Test on a *validation set*
- If the overall false positive rate is not low enough, then add another stage
- Use false positives from current stage as the negative training examples for the next stage

Viola Jones Results

- 38 stages with
 - 6061 total Haar features used out of 160K features candidates
- Training time: “weeks” on 466 MHz Sun workstation
- Average of 10 features evaluated per window on test set
- On 700 Mhz Pentium III processor, process 384x288 pixel image in 0.067 sec
 - 15 Hz
 - 15 times faster than previous detector of comparable accuracy (Rowley et al., 1998)

