

# **Introduction to Computer Vision**

---

**Kaveh Fathian**

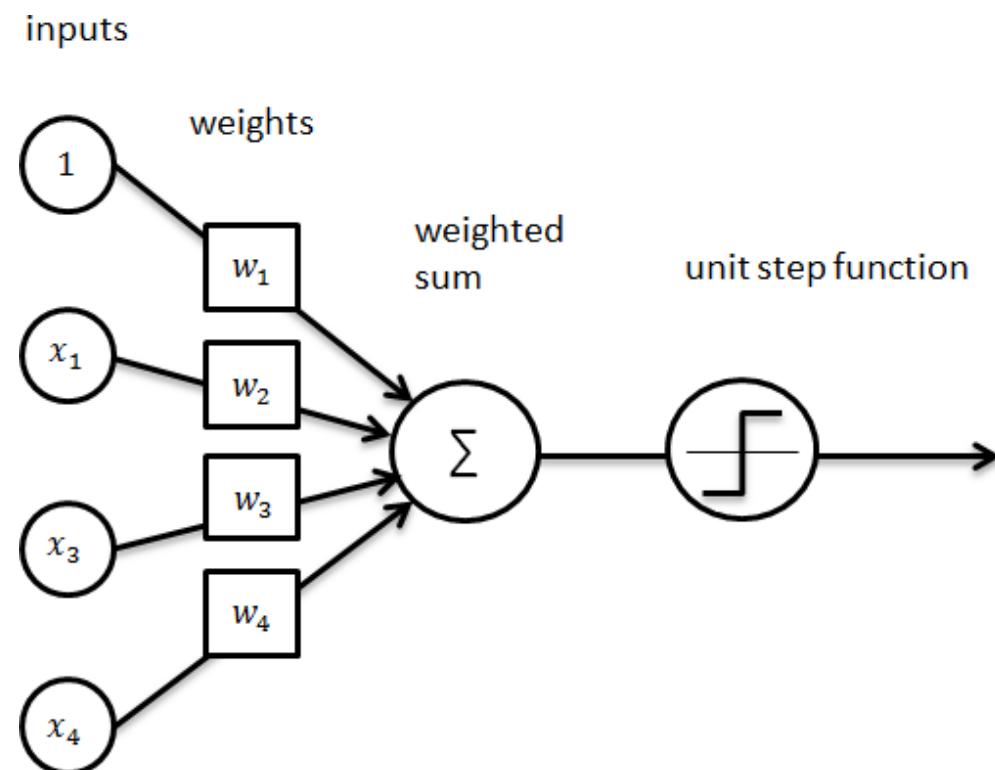
Assistant Professor

Computer Science Department

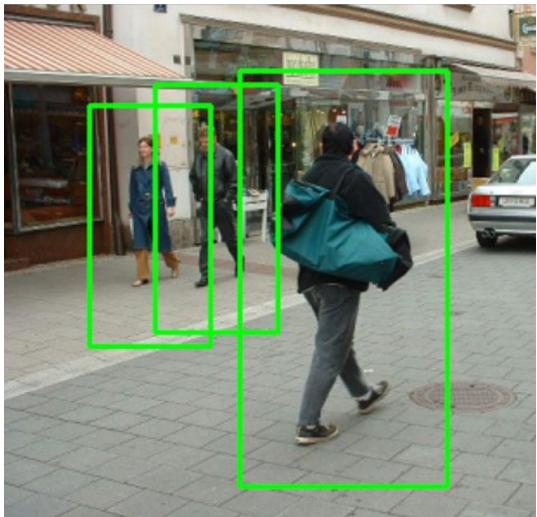
Colorado School of Mines

**Lecture 18**

# Introduction to Neural Networks



# Hand-Engineered Methods



Recognition:

Image formation (+database+labels)

Captured+manual.

Filtering (gradients/transforms)

Hand designed.

Feature points  
(saliency+description)

Hand designed.

Dictionary building  
(compression/quantization)

Hand designed.

Classifier (decision making)

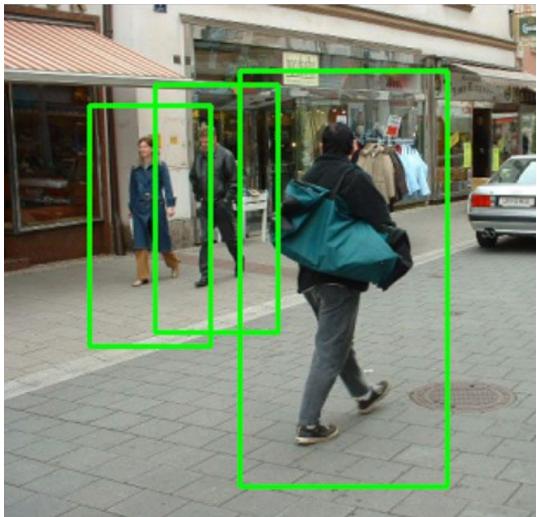
Learned.

Classification

Object  
Detection

Segmentation

# Learning-Based Methods



Recognition:

Image formation (+database+labels)

Filtering (gradients/transforms)

Feature points  
(saliency+description)

Dictionary building  
(compression/quantization)

Classifier (decision making)

Classification

Object  
Detection

Segmentation

Captured+manual.

Learned.

Learned.

Learned.

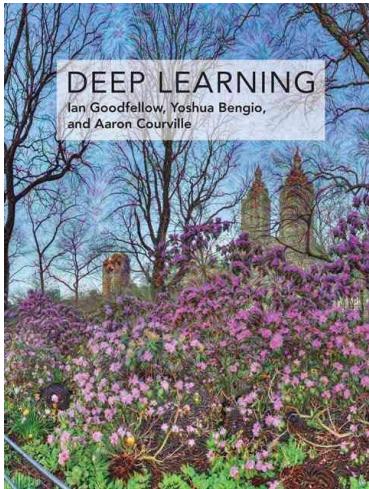
Learned.

End to end  
learning

# Deep learning

- Modeling the visual world is incredibly complicated. We need high-capacity models.
- In the past, we didn't have enough data to fit these models. But now we do!
- We want a class of **high-capacity models** that are **easy to optimize**:

## Deep neural networks!



<http://www.deeplearningbook.org/>

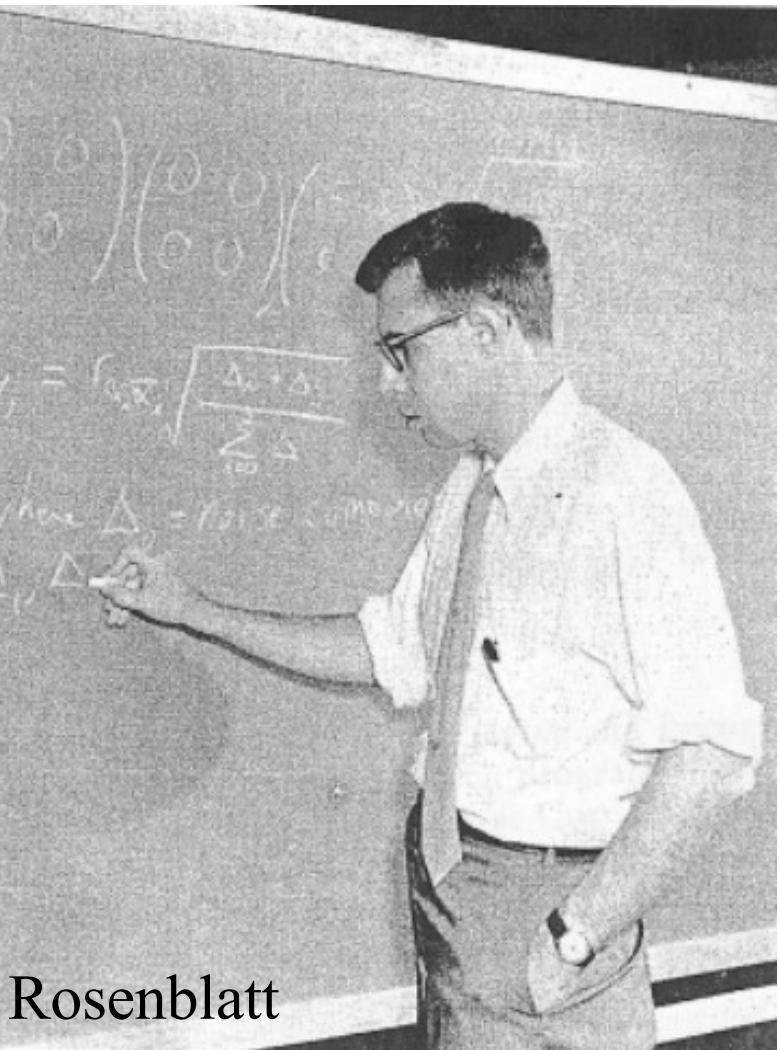
By Ian Goodfellow, Yoshua Bengio and Aaron Courville

November 2016

# A brief history of Neural Networks

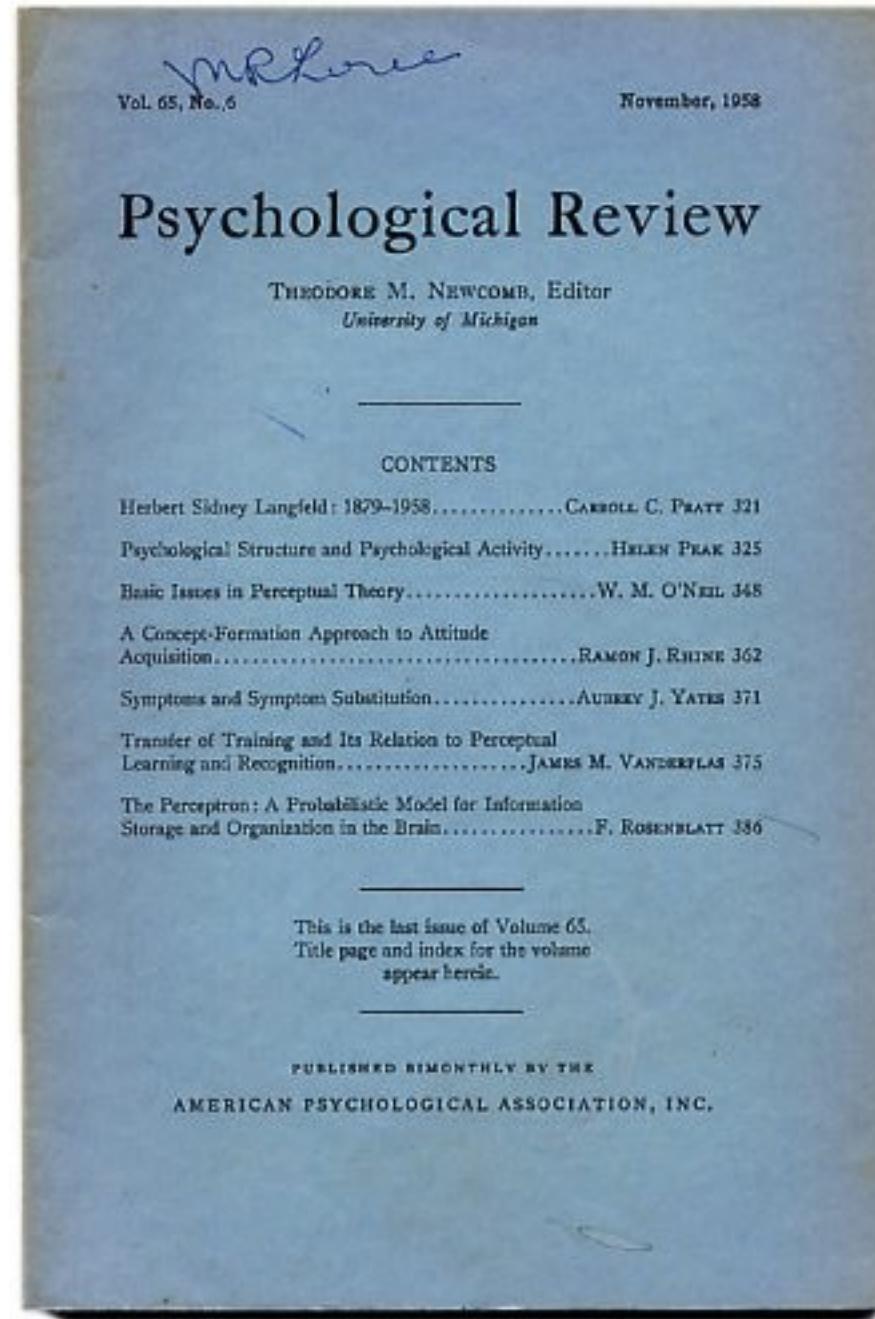


# Perceptrons, 1958



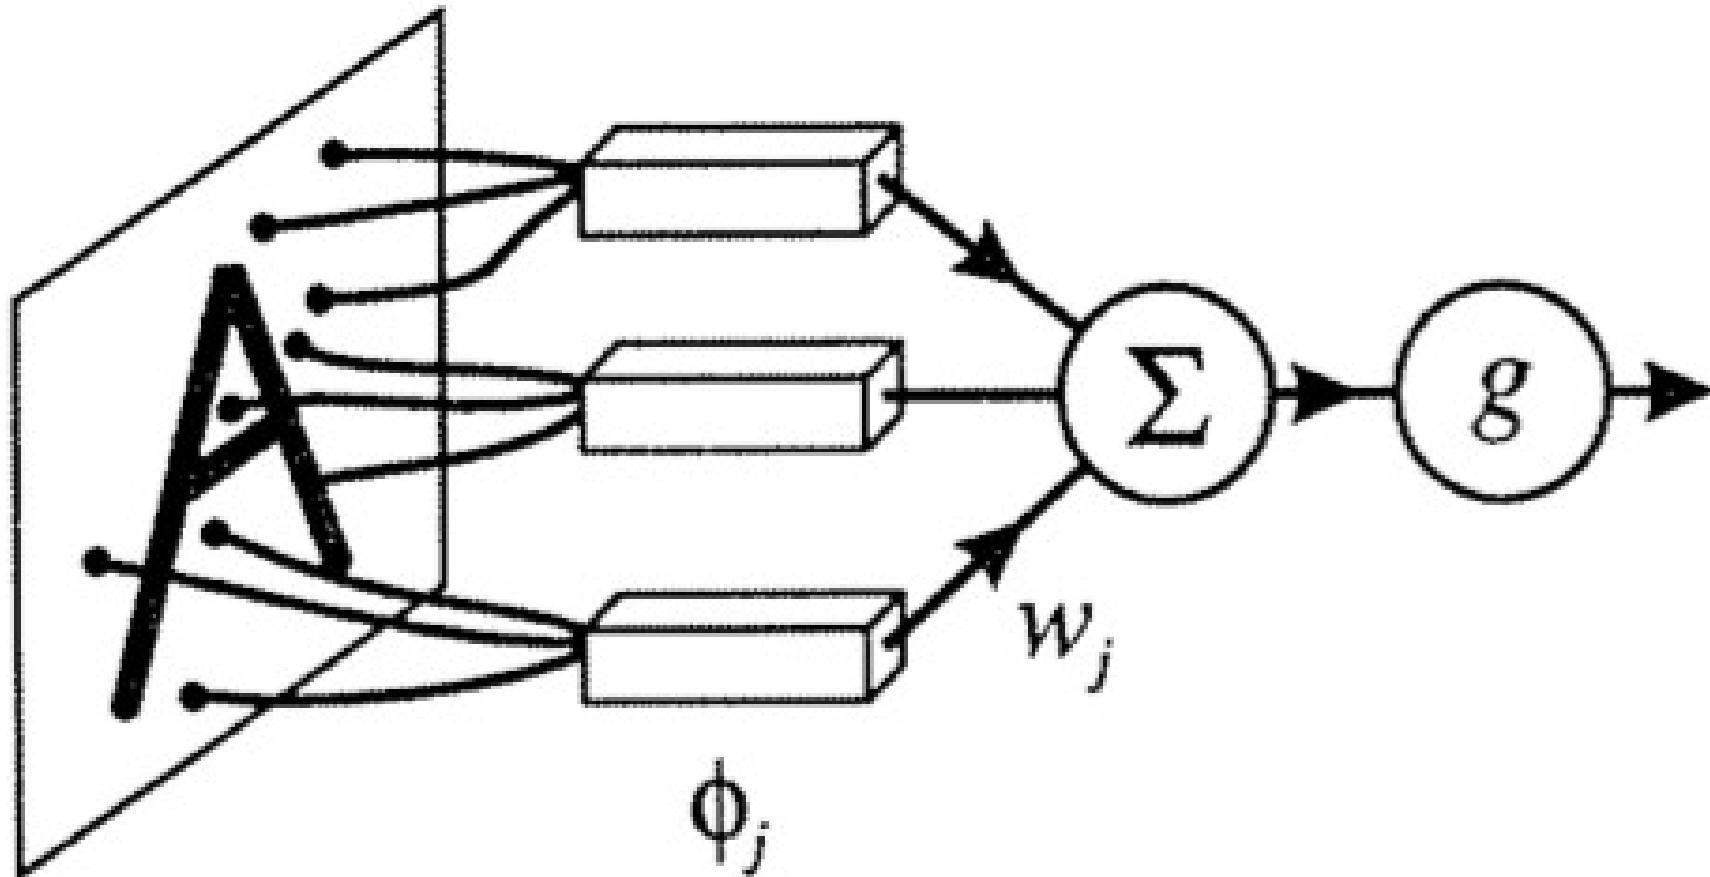
Rosenblatt

[http://www.ecse.rpi.edu/homepages/nagy/PDF\\_chrono/2011\\_Nagy\\_Pace\\_FR.pdf](http://www.ecse.rpi.edu/homepages/nagy/PDF_chrono/2011_Nagy_Pace_FR.pdf). Photo by George Nagy



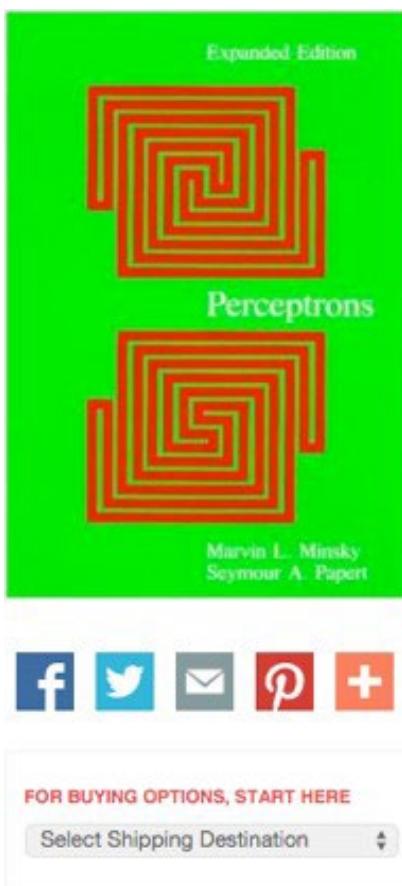
<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.335.3398&rep=rep1&type=pdf>

# Perceptrons, 1958





# Minsky and Papert, Perceptrons, 1972



## Perceptrons, expanded edition

An Introduction to Computational Geometry

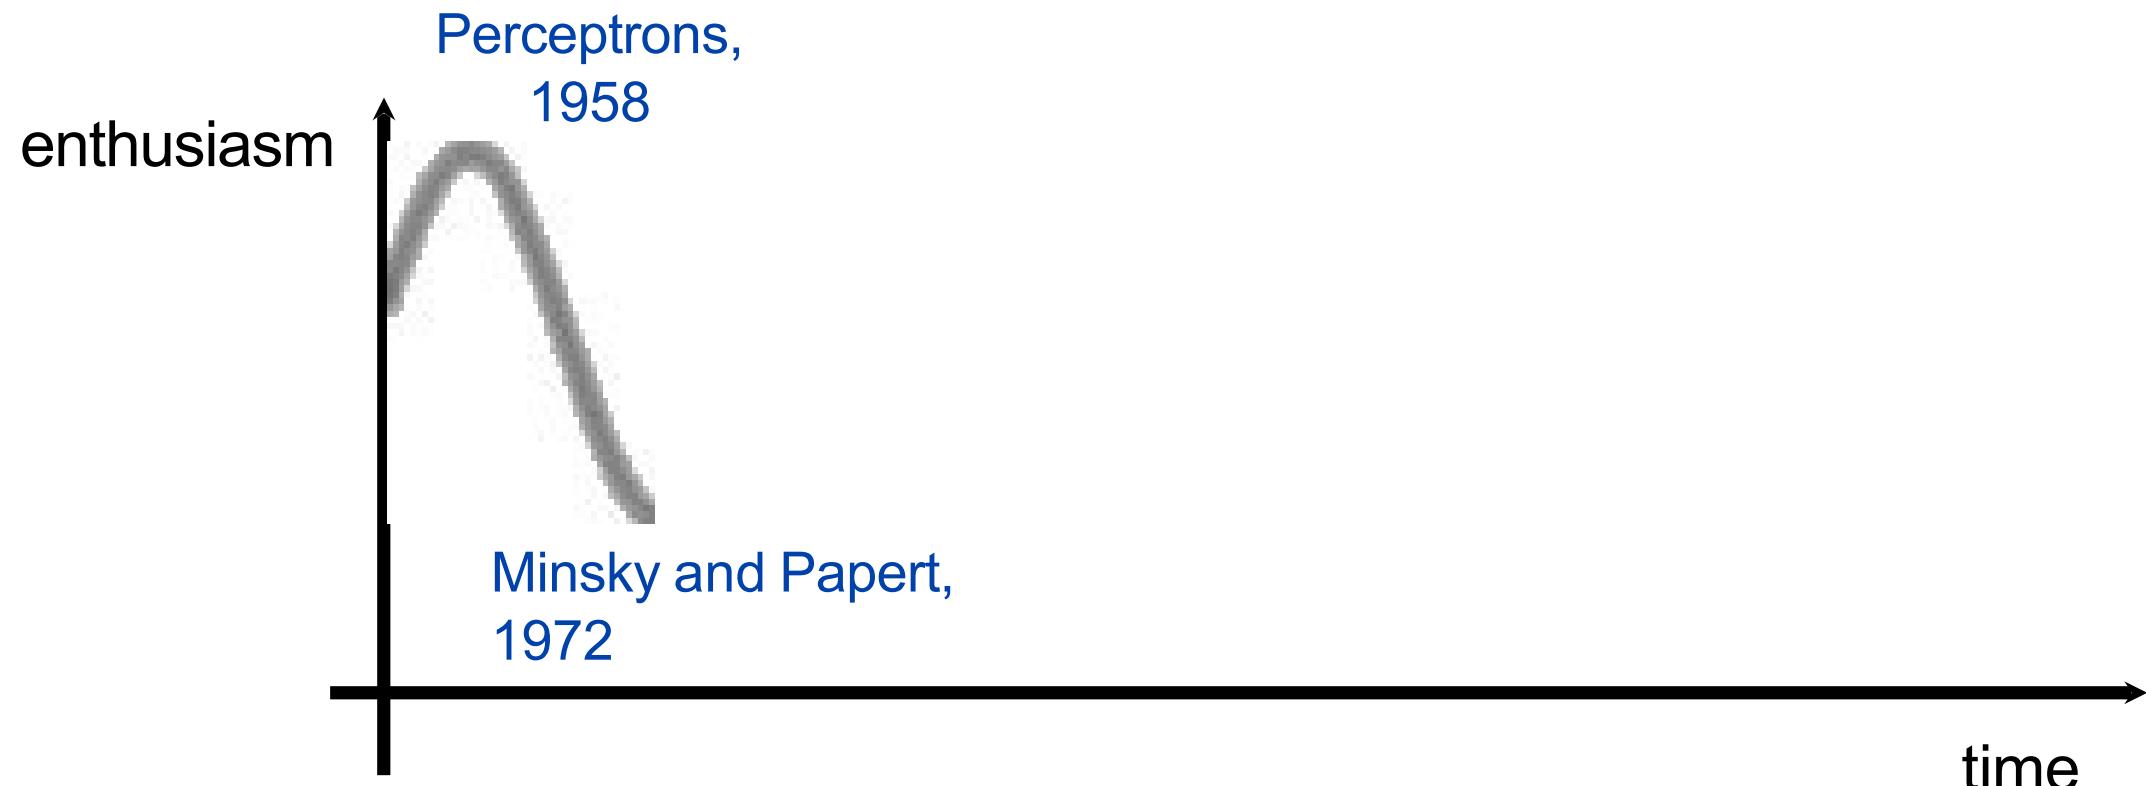
By [Marvin Minsky](#) and [Seymour A. Papert](#)

### Overview

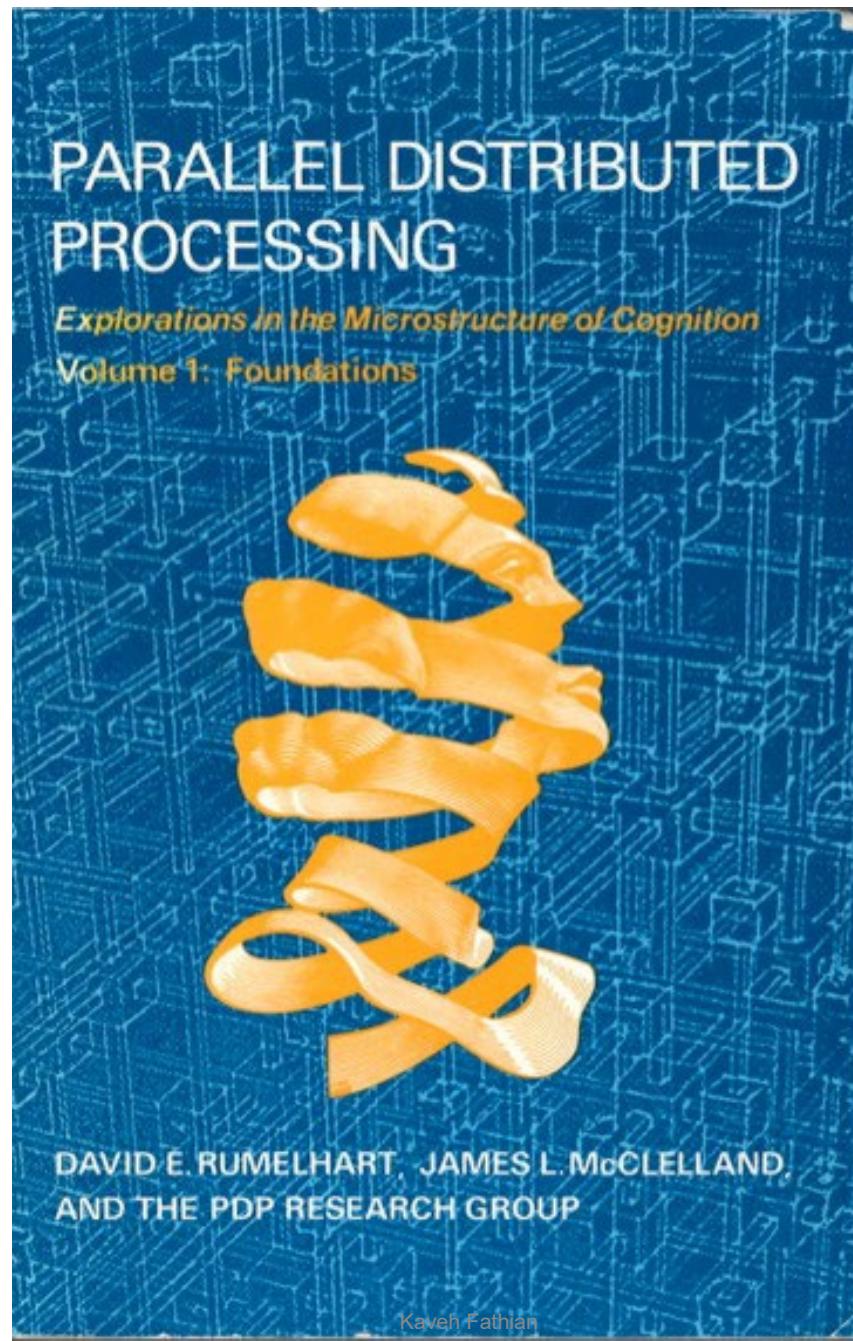
*Perceptrons* - the first systematic study of parallelism in computation - has remained a classical work on threshold automata networks for nearly two decades. It marked a historical turn in artificial intelligence, and it is required reading for anyone who wants to understand the connectionist counterrevolution that is going on today.

Artificial-intelligence research, which for a time concentrated on the programming of von Neumann computers, is swinging back to the idea that intelligence might emerge from the activity of networks of neuronlike entities. Minsky and Papert's book was the first example of a mathematical analysis carried far enough to show the exact limitations of a class of computing machines that could seriously be considered as models of the brain. Now the new developments in mathematical tools, the recent interest of physicists in the theory of disordered matter, the new insights into and psychological models of how the brain works, and the evolution of fast computers that can simulate networks of automata have given *Perceptrons* new importance.

Witnessing the swing of the intellectual pendulum, Minsky and Papert have added a new chapter in which they discuss the current state of parallel computers, review developments since the appearance of the 1972 edition, and identify new research directions related to connectionism. They note a central theoretical challenge facing connectionism: the challenge to reach a deeper understanding of how "objects" or "agents" with individuality can emerge in a network. Progress in this area would link connectionism with what the authors have called "society theories of mind."

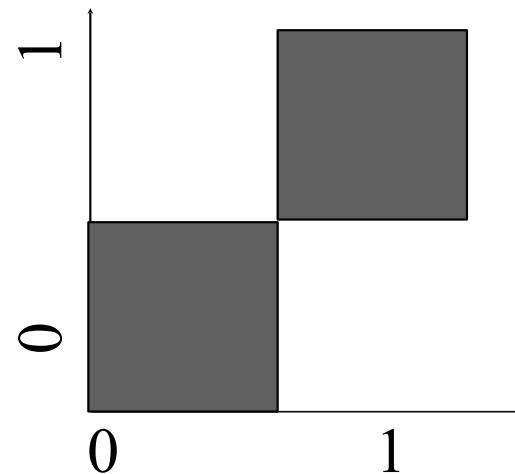


# Parallel Distributed Processing (PDP), 1986

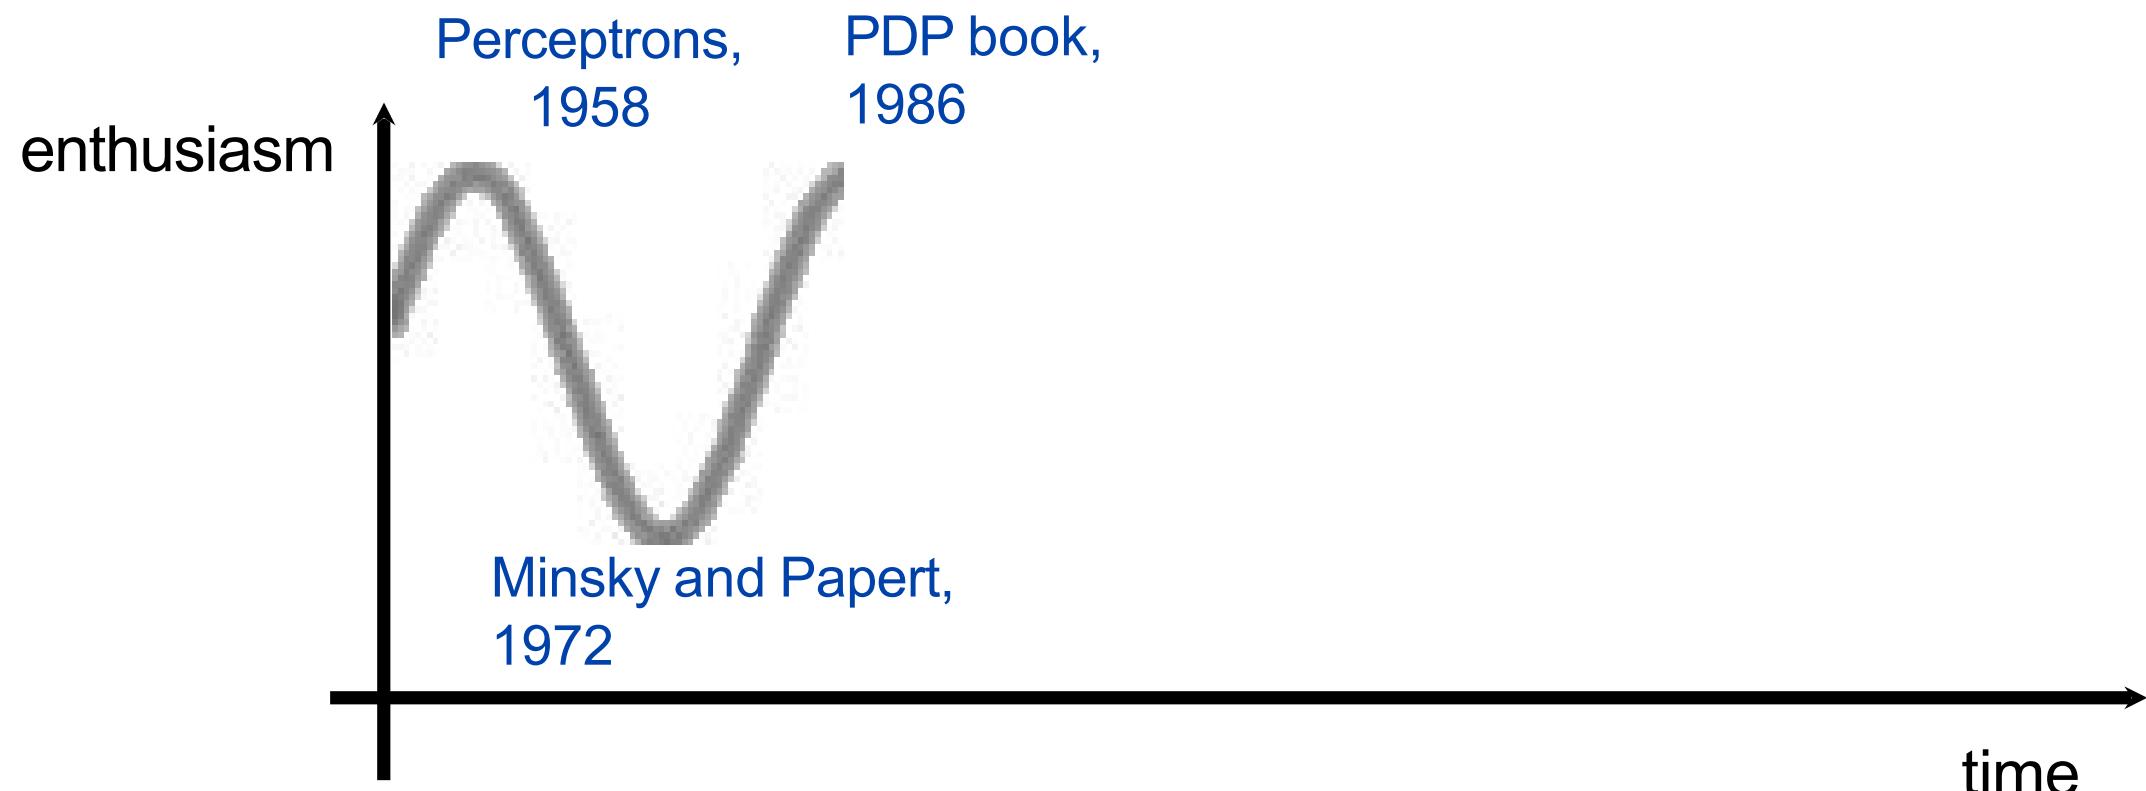


# XOR problem

Inputs		Output
0	0	0
1	0	1
0	1	1
1	1	0



- PDP authors pointed to the backpropagation algorithm as a breakthrough.
- This allowed multi-layer neural networks to be trained.
- Functions that a multi-layer network can represent but a single-layer network cannot: XOR function.



# LeCun conv nets, 1998

PROC. OF THE IEEE, NOVEMBER 1998

7

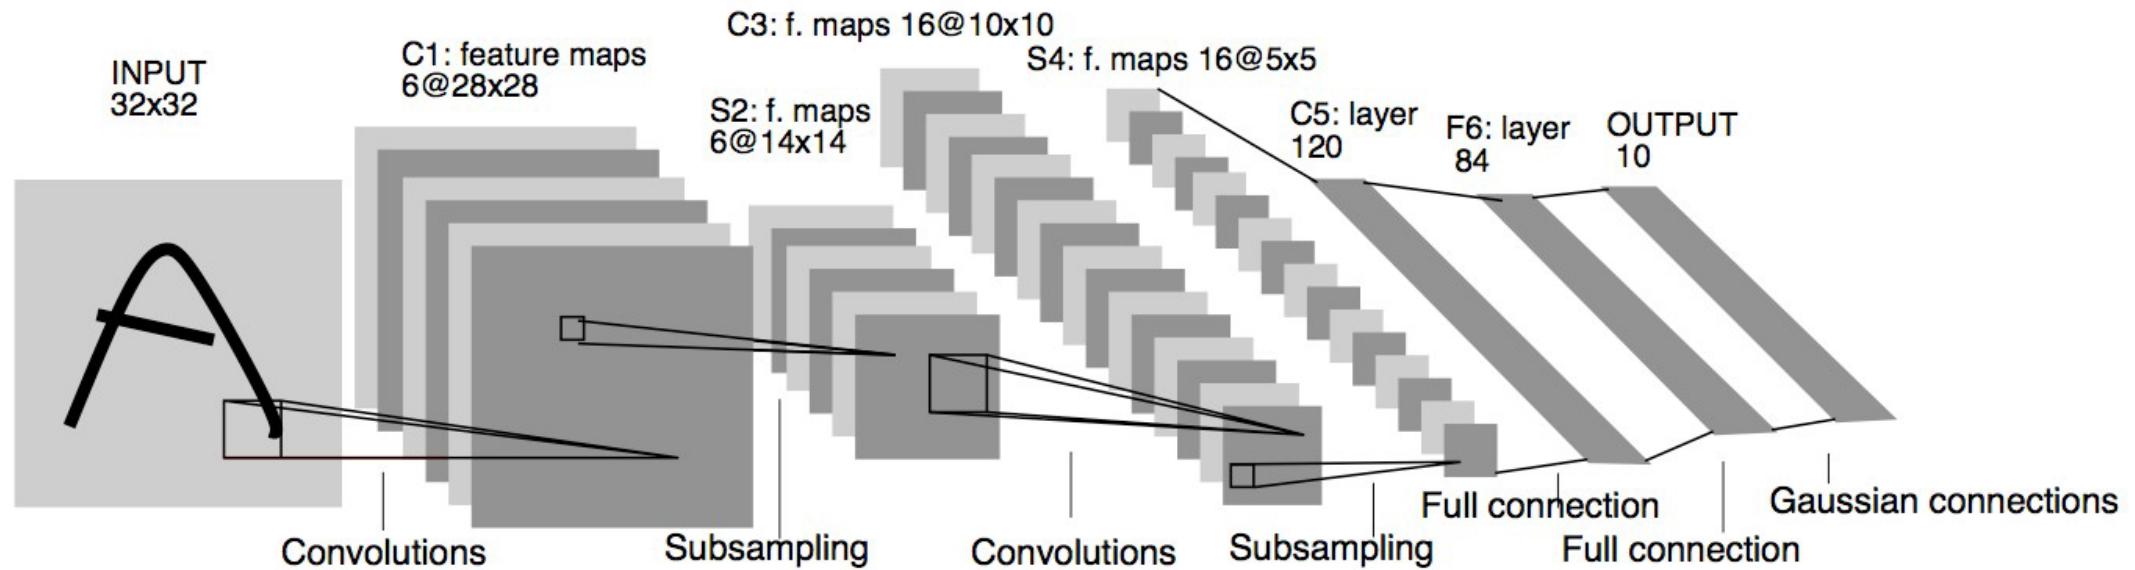


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Demos: <http://yann.lecun.com/exdb/lenet/index.html>

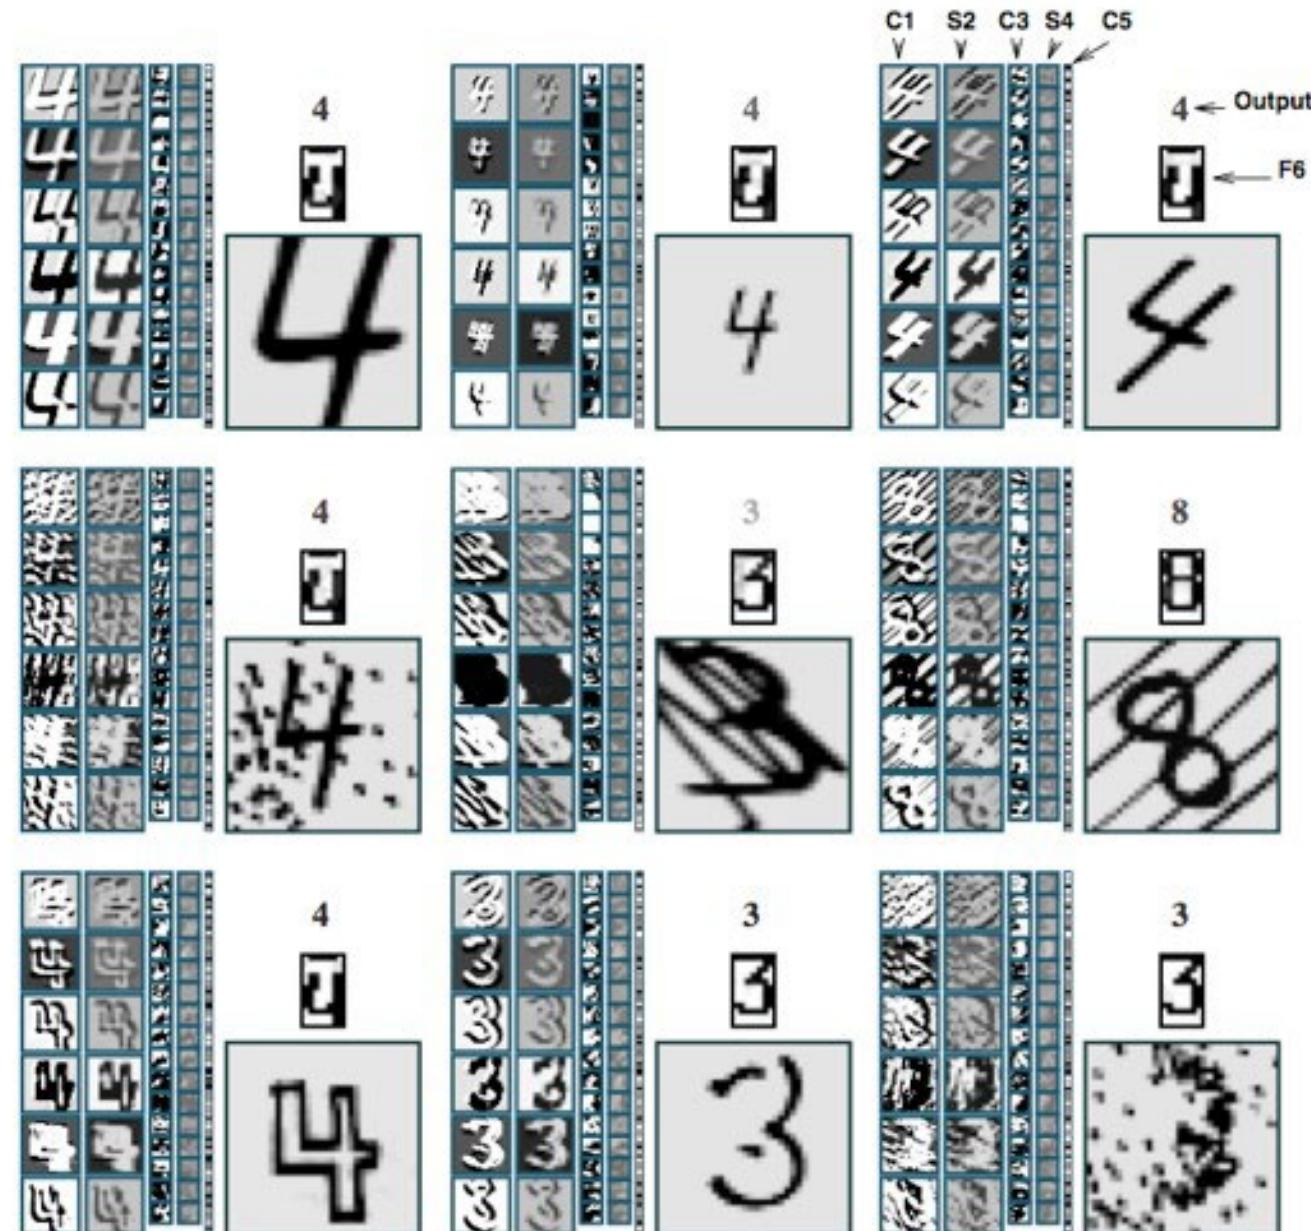
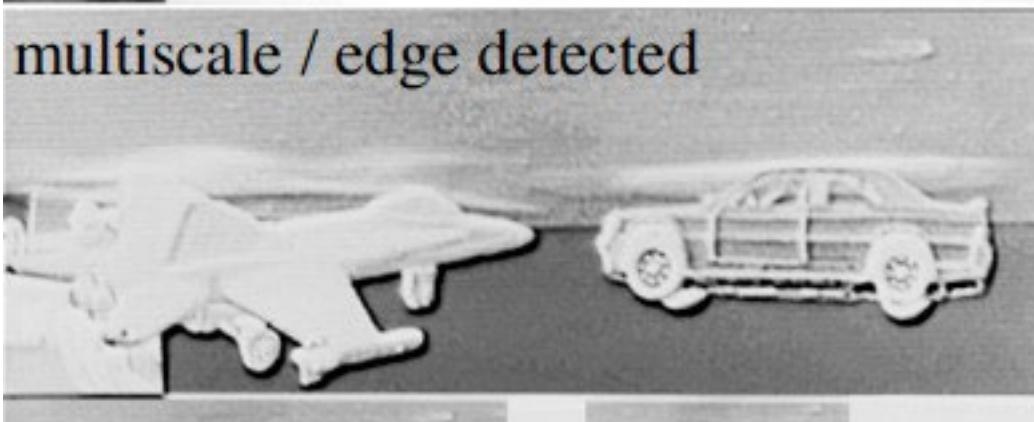


Fig. 13. Examples of unusual, distorted, and noisy characters correctly recognized by LeNet-5. The grey-level of the output label represents the penalty (lighter for higher penalties).

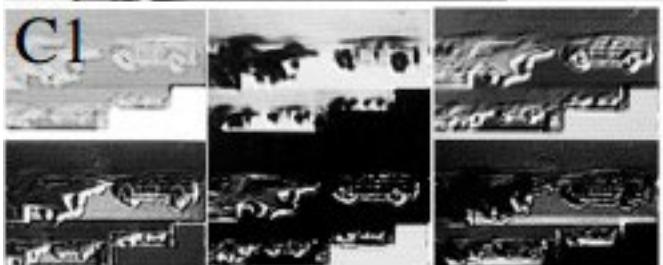
input



multiscale / edge detected



C1



C2



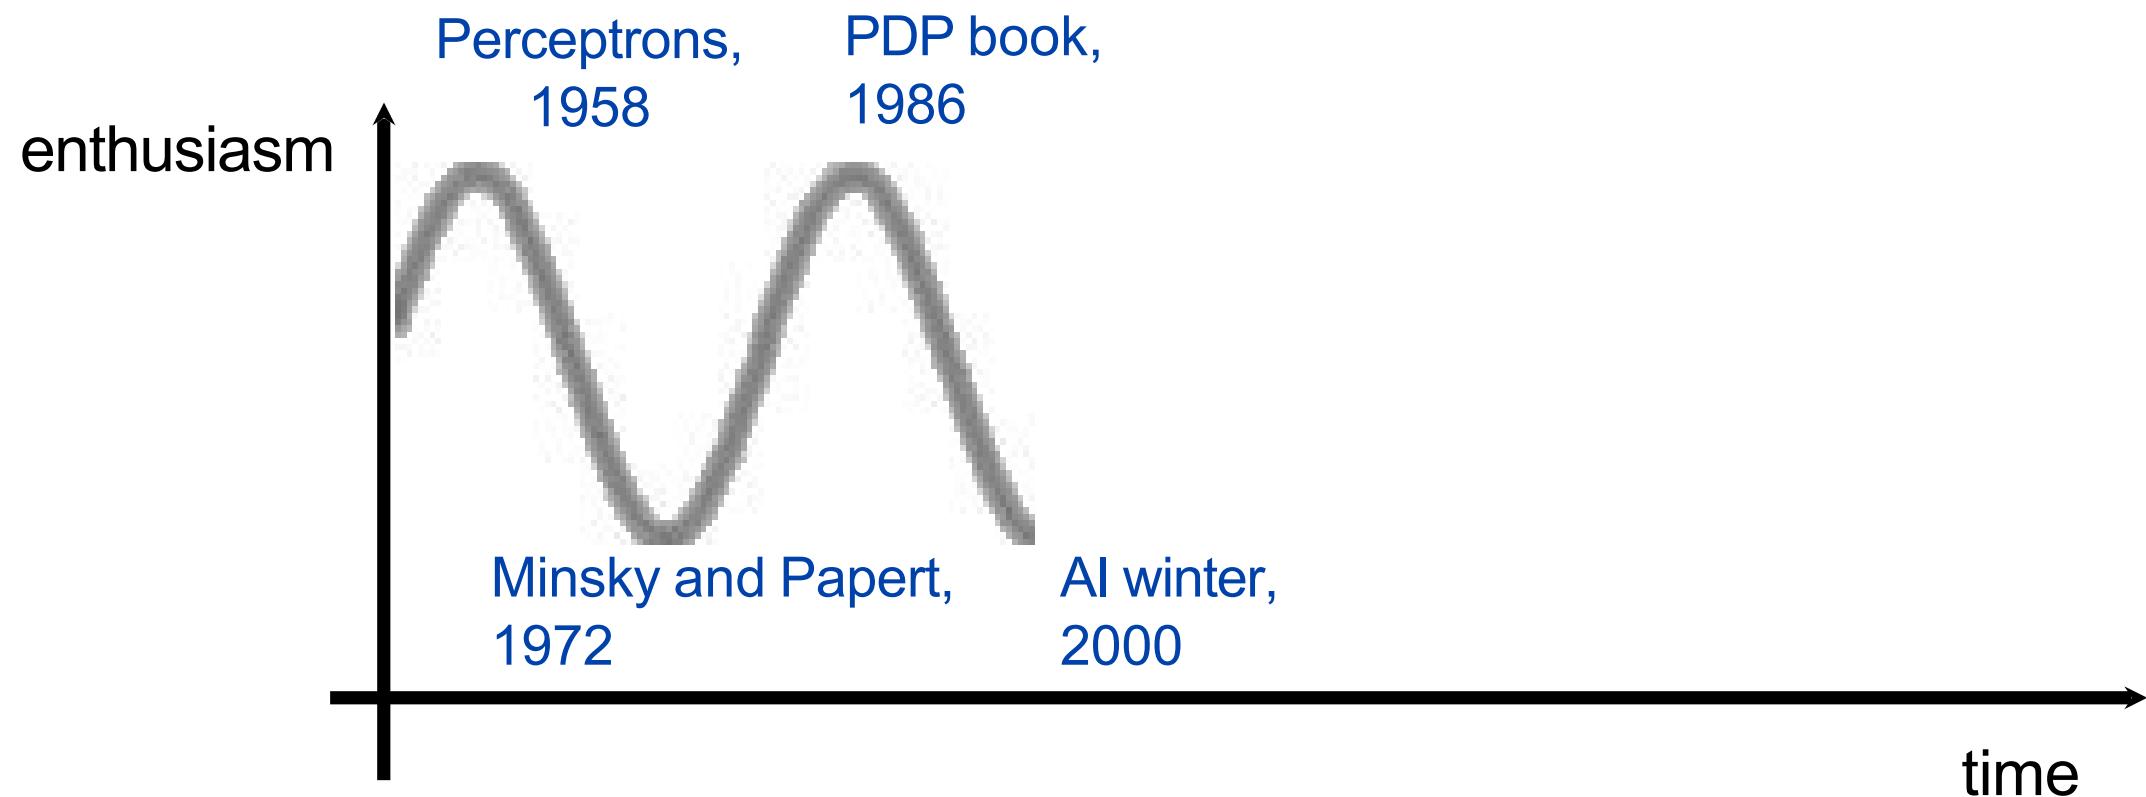
<http://pub.clement.farabet.net/ecvw09.pdf>

Neural networks to  
recognize  
handwritten digits?  
**yes**

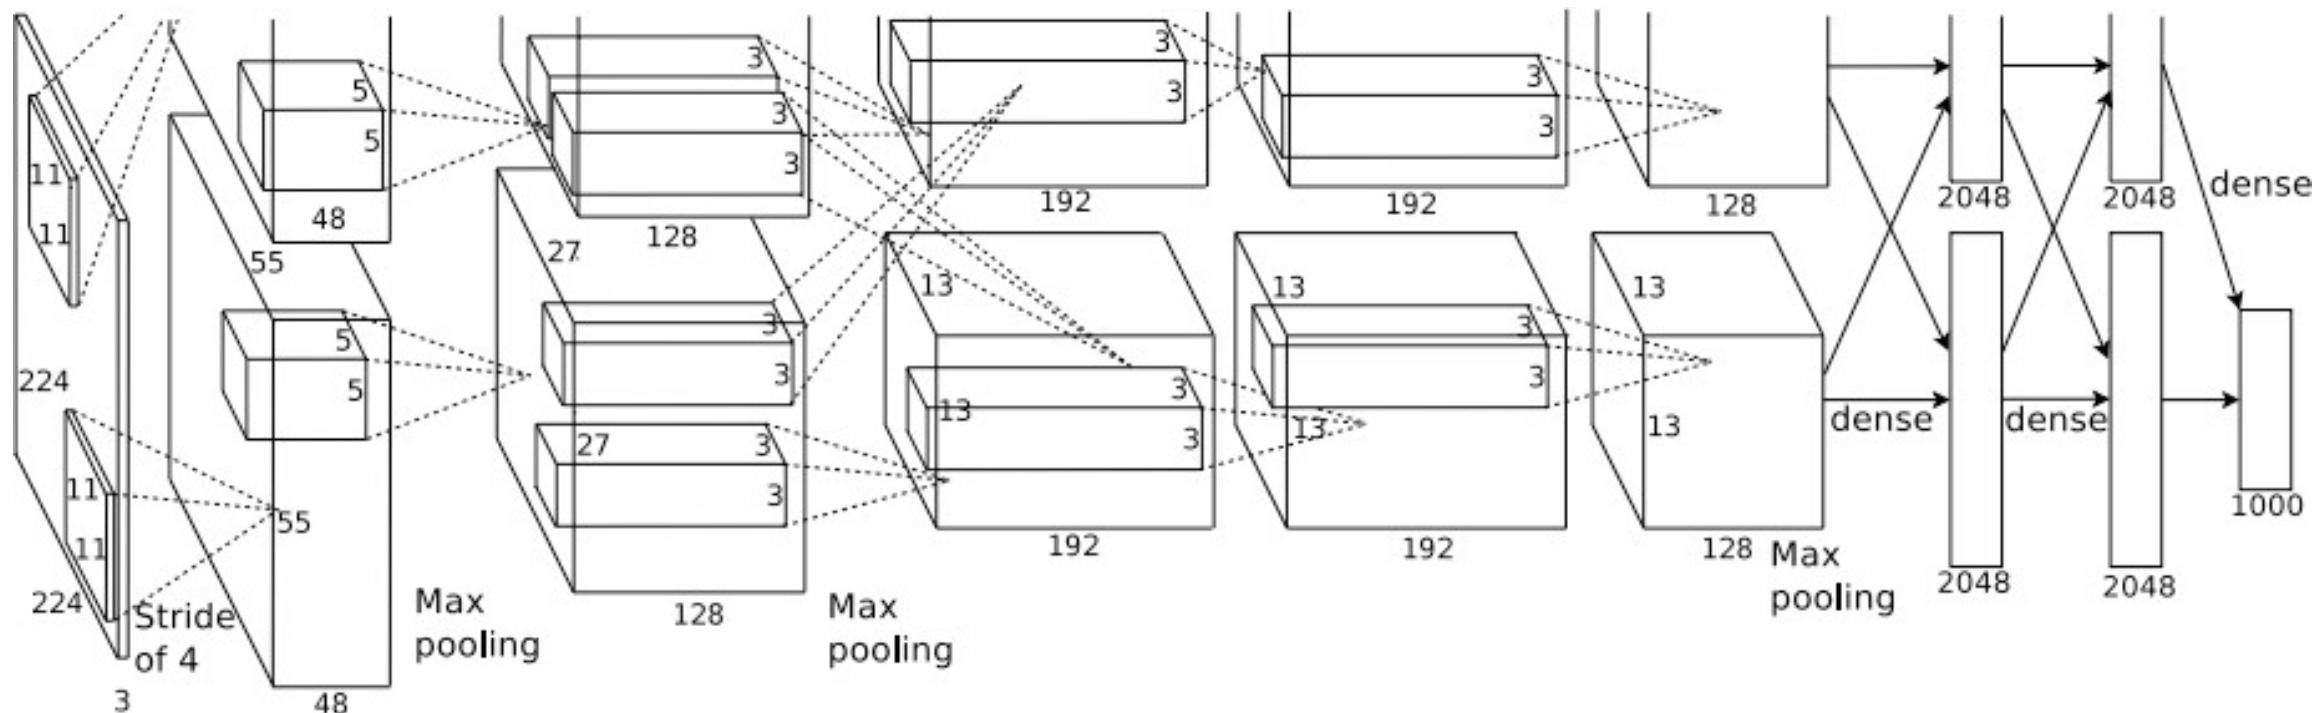
Neural networks for  
tougher problems?  
**not really**

# Neural Information Processing Systems 2000

- Neural Information Processing Systems (NeurIPS), is the premier conference on machine learning.
- Evolved from an interdisciplinary conference to a machine learning conference.
- For the 2000 conference:
  - Title words predictive of paper acceptance: “Belief Propagation” and “Gaussian”.
  - Title words predictive of paper rejection: “Neural” and “Network”.



# Krizhevsky, Sutskever, & Hinton, NeurIPS2012 “Alexnet”



# Krizhevsky, Sutskever, & Hinton, NeurIPS 2012



mite

container ship

motor scooter

leopard

mite	container ship	motor scooter	leopard
black widow	lifeboat	go-kart	jaguar
cockroach	amphibian	moped	cheetah
tick	fireboat	bumper car	snow leopard
starfish	drilling platform	golfcart	Egyptian cat



grille

mushroom

cherry

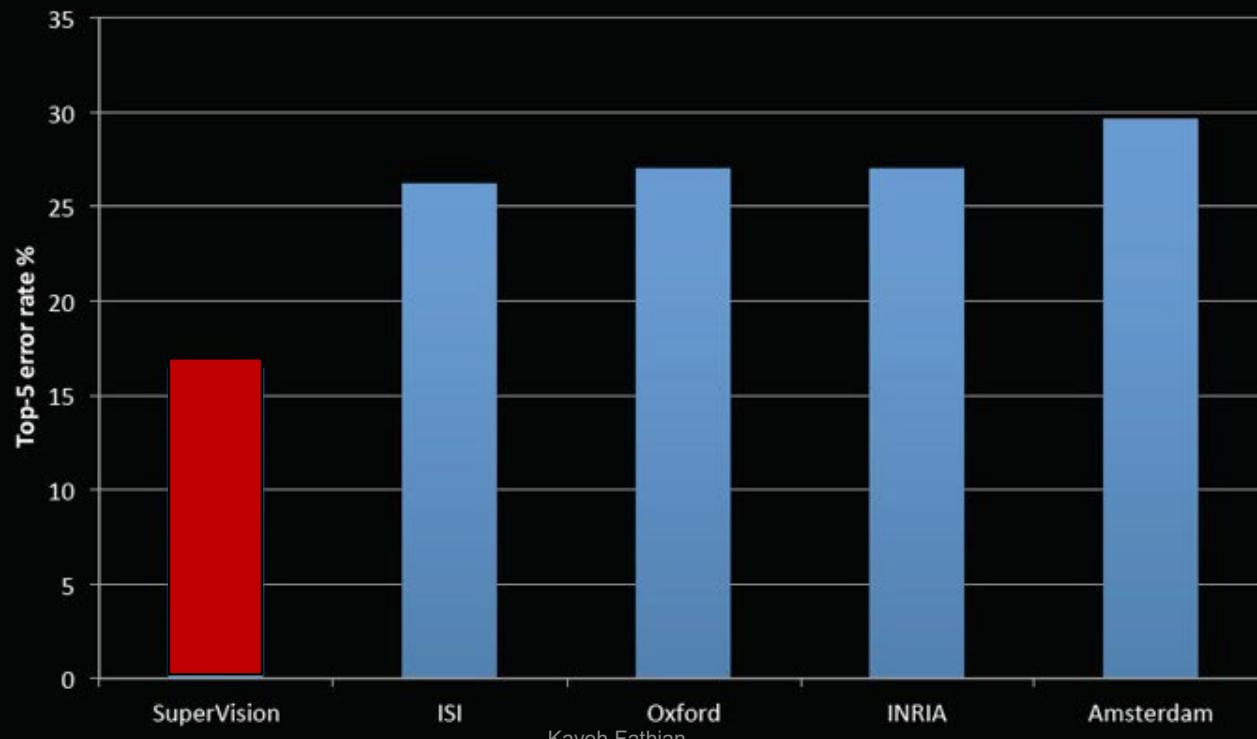
Madagascar cat

convertible	agaric	dalmatian	squirrel monkey
grille	mushroom	grape	spider monkey
pickup	jelly fungus	elderberry	titi
beach wagon	gill fungus	ffordshire bulterrier	indri
fire engine	dead-man's-fingers	currant	howler monkey

# ImageNet Classification 2012

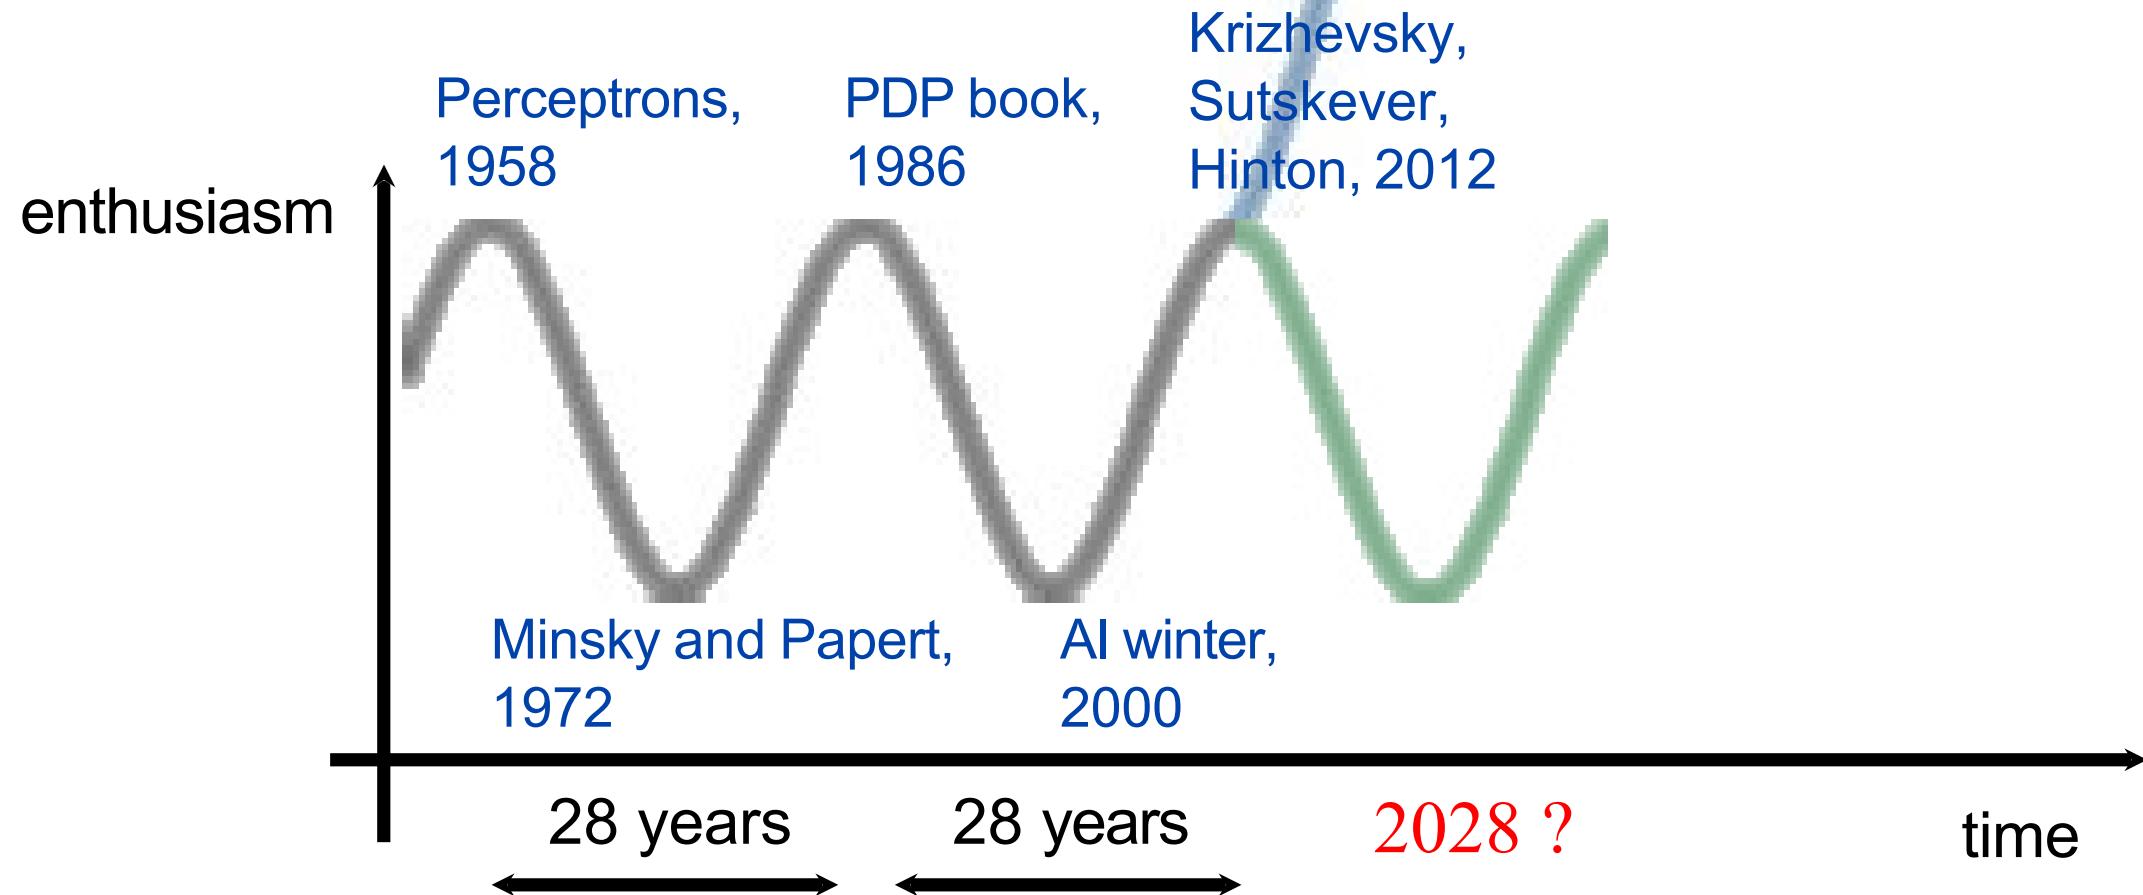
.....

- Krizhevsky et al. -- 16.4% error (top-5)
- Next best (non-convnet) – 26.2% error

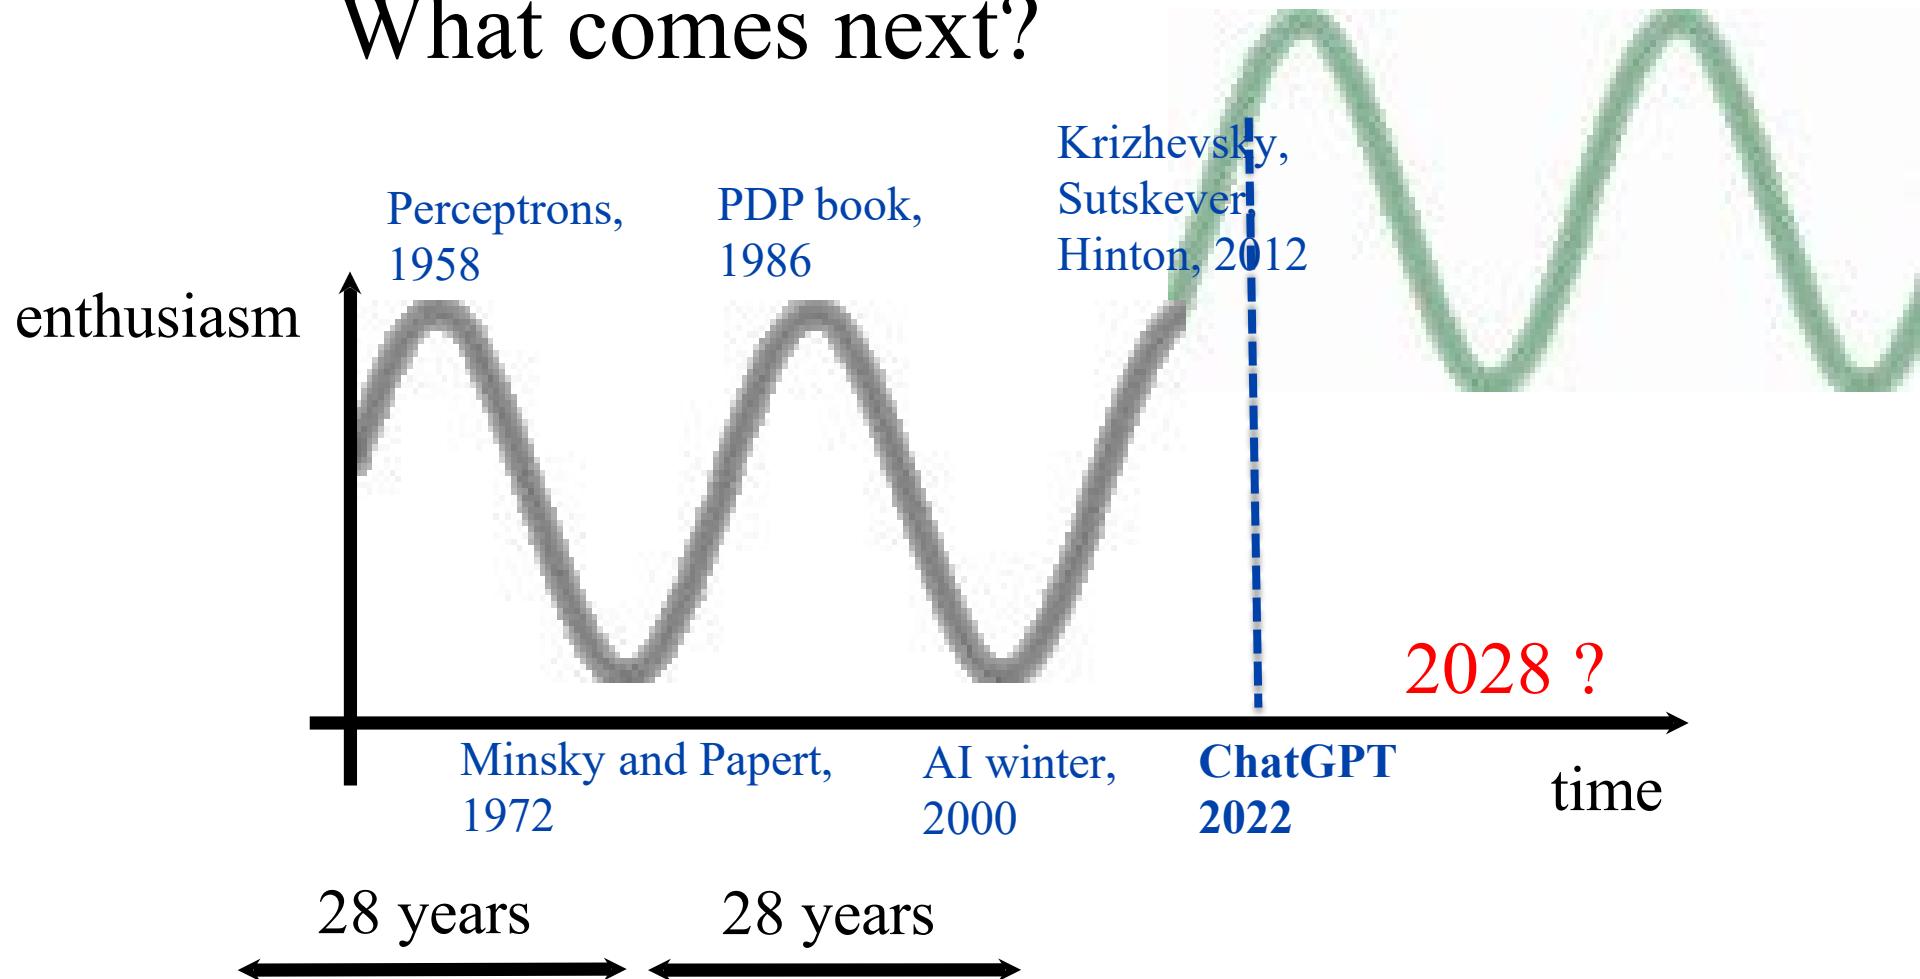




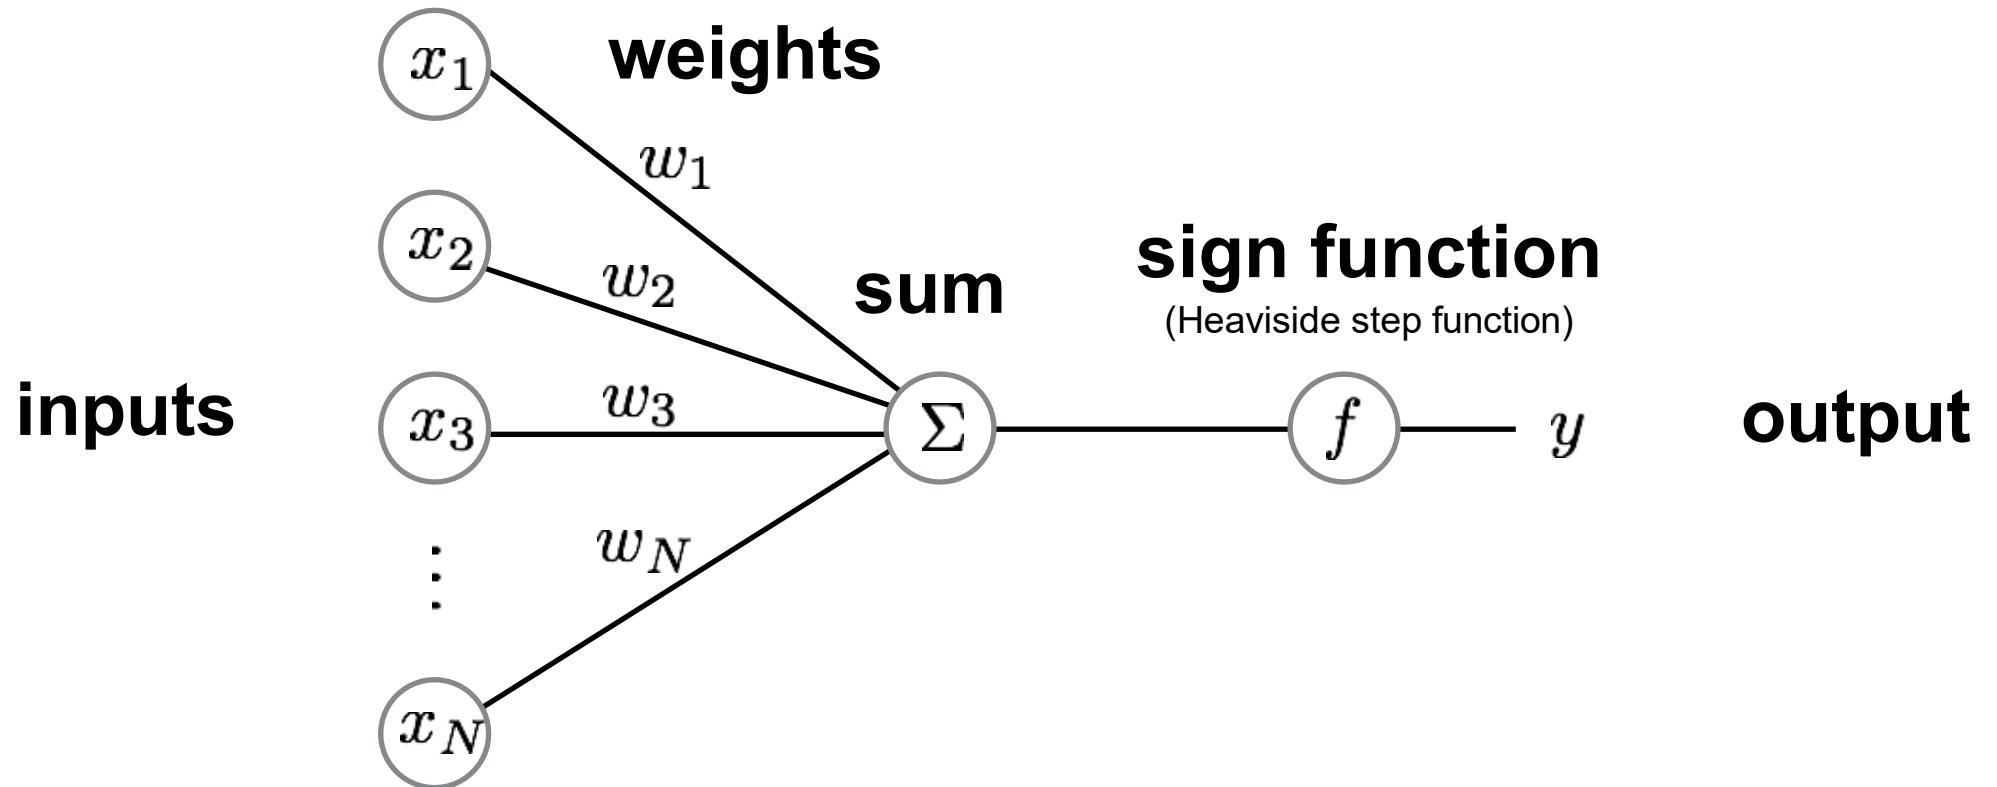
# What comes next?



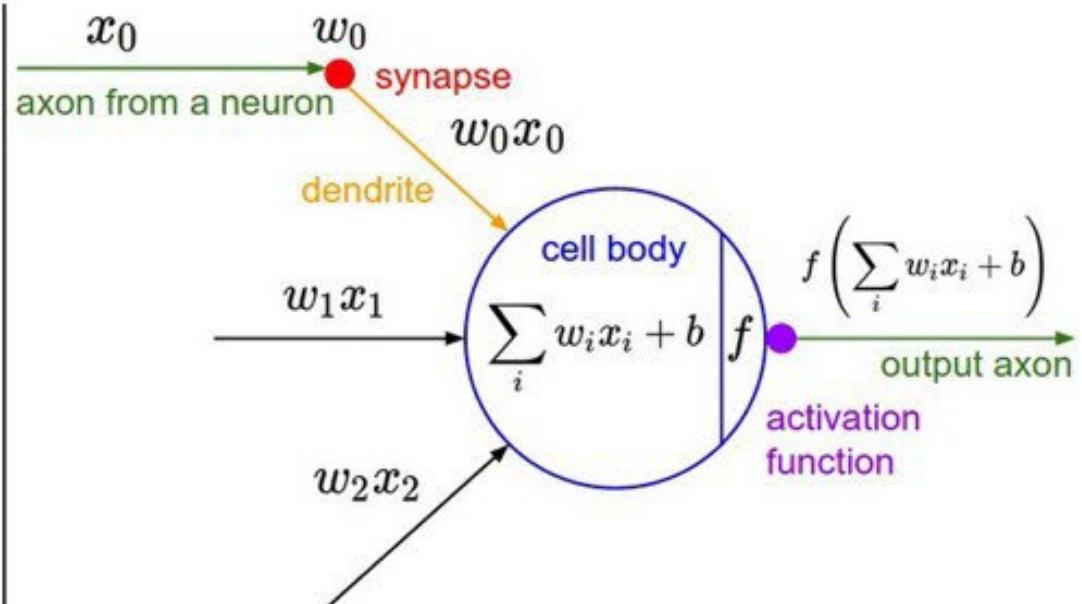
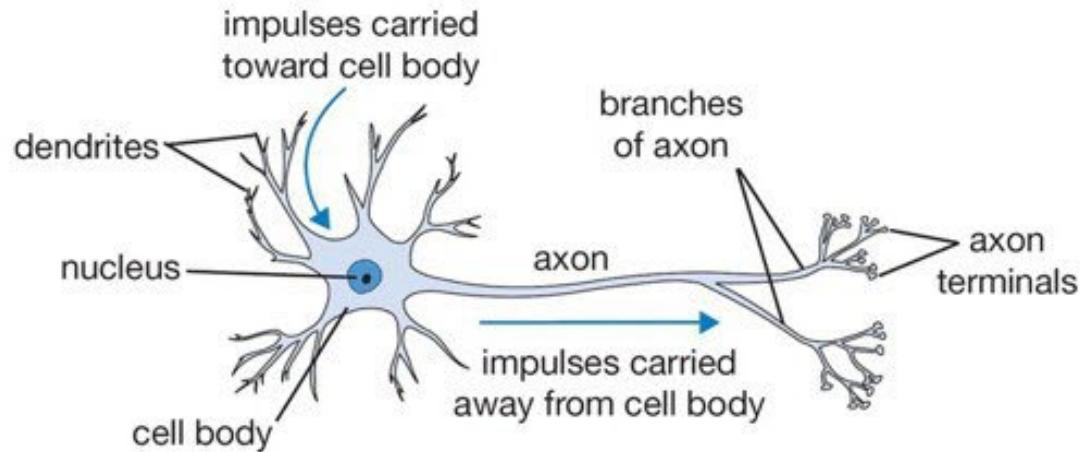
# What comes next?



# The Perceptron



# Aside: Inspiration from Biology



A cartoon drawing of a biological neuron (left) and its mathematical model (right).

- Neural nets/perceptrons are **loosely** inspired by biology.
- But they certainly are **not** a model of how the brain works, or even how neurons work. ➔ cf. spiking neural networks

# Perceptron Algorithm

- 1: **function** PERCEPTRON ALGORITHM
- 2:      $\mathbf{w}^{(0)} \leftarrow \mathbf{0}$
- 3:     **for**  $t = 1, \dots, T$  **do**
- 4:         RECEIVE( $\mathbf{x}^{(t)}$ )                       $\mathbf{x} \in \{0, 1\}^N$  N-d binary vector
- 5:          $\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$               perceptron is just one line of code!  
                    sign of zero is +1
- 6:         RECEIVE( $y^t$ )                       $y \in \{1, -1\}$
- 7:          $w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$

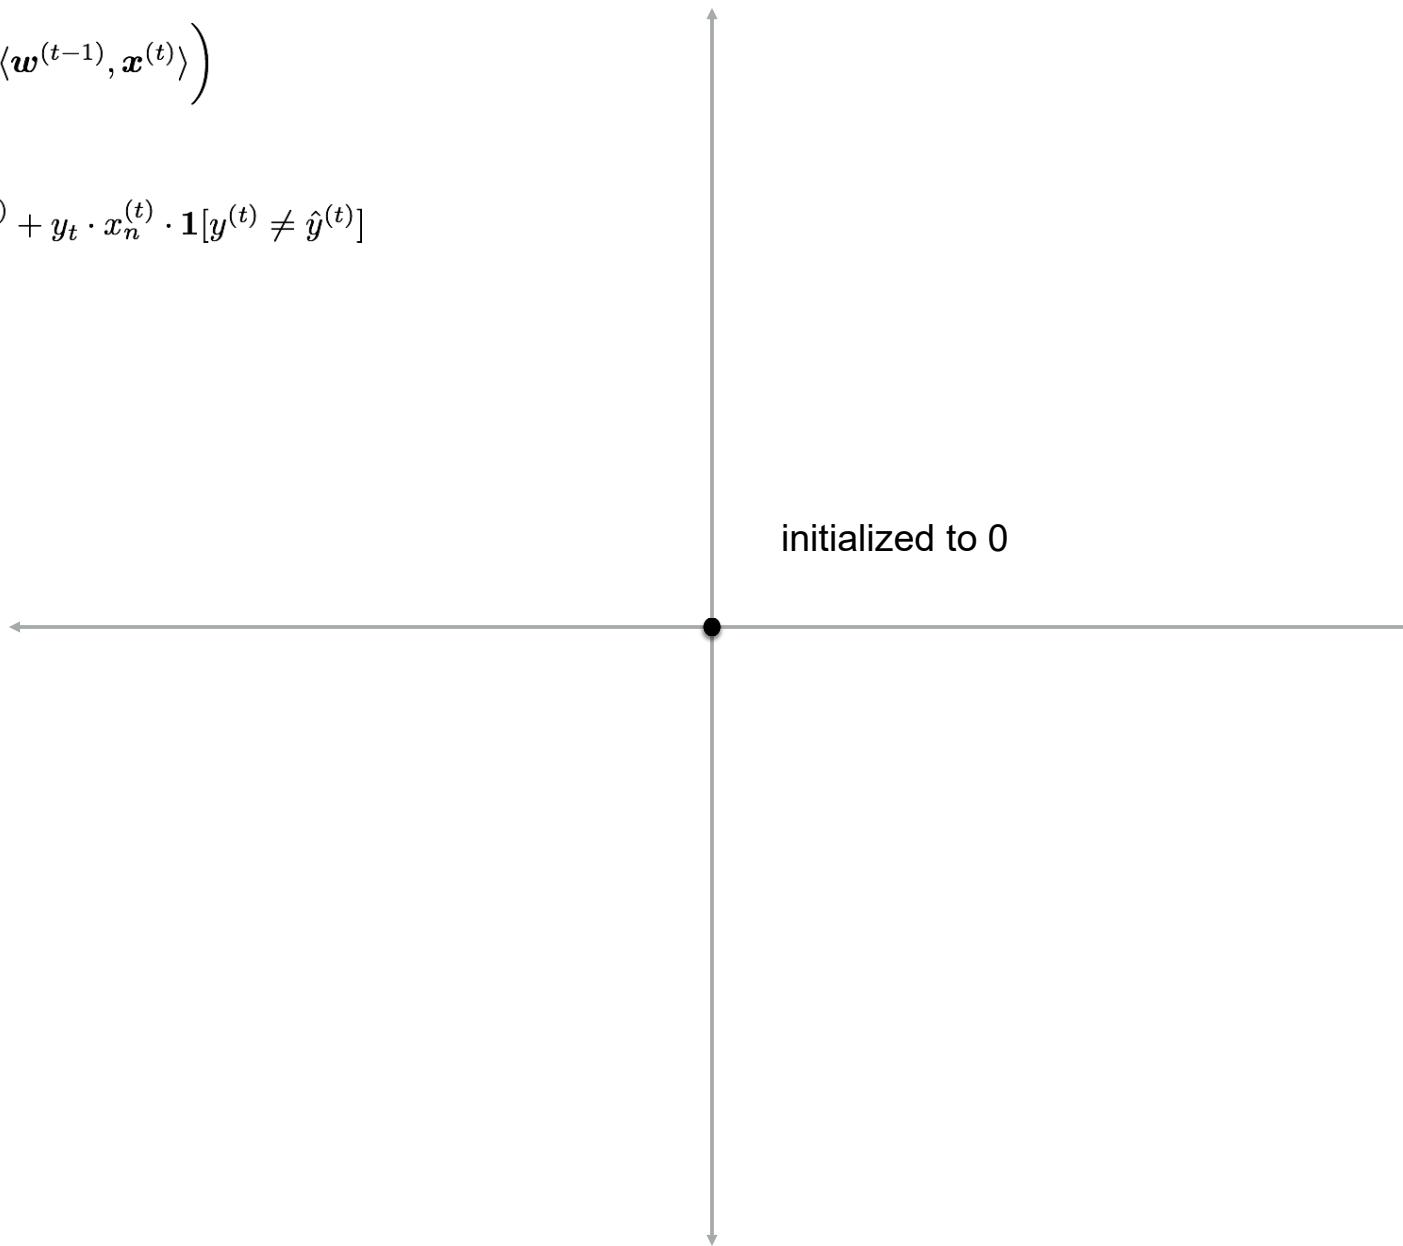
RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

initialized to 0

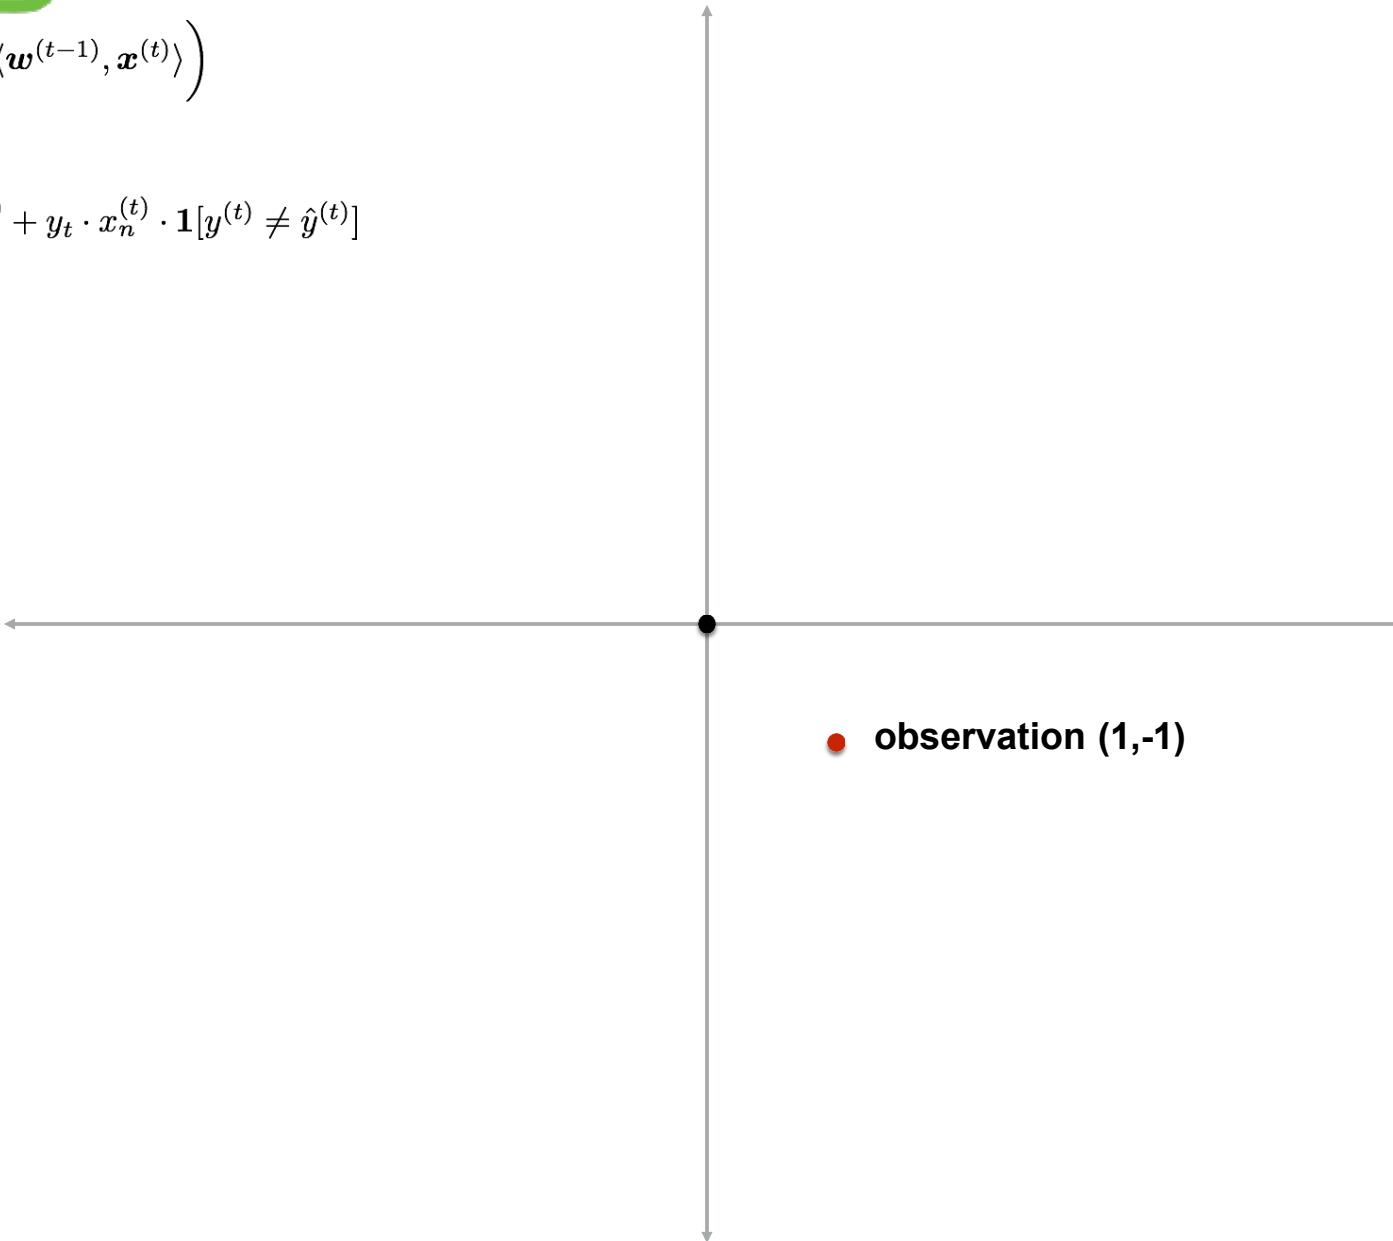


RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

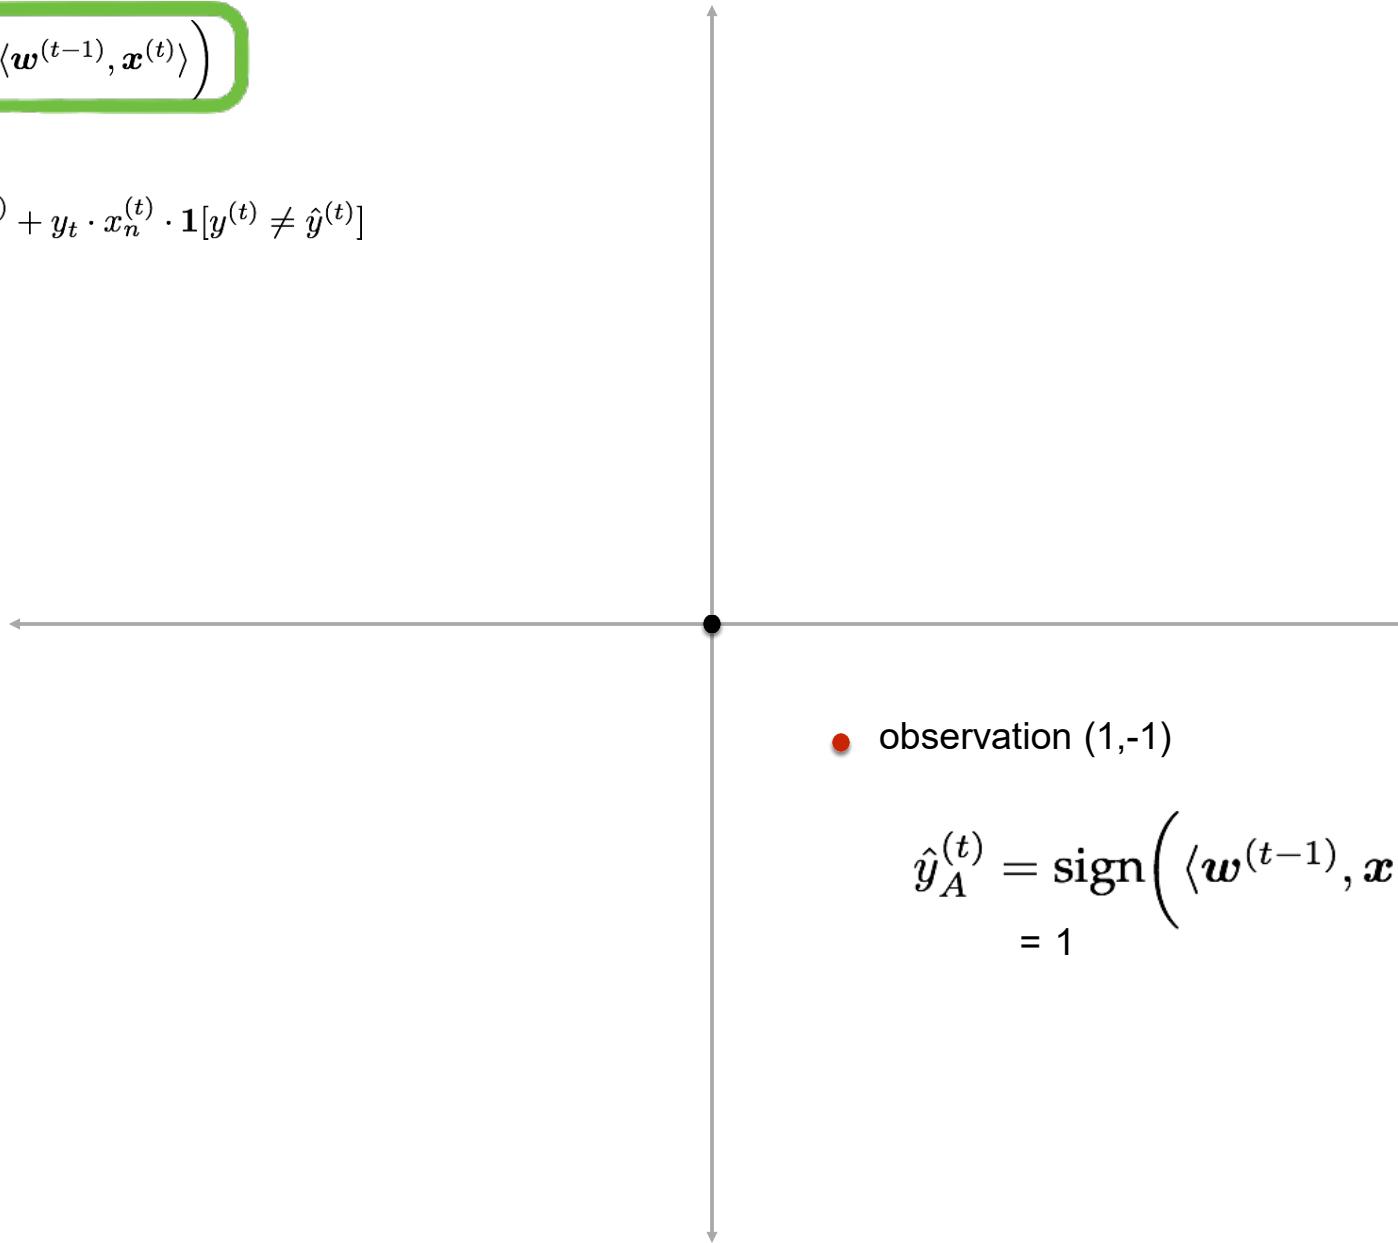


RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$



● observation (1,-1)

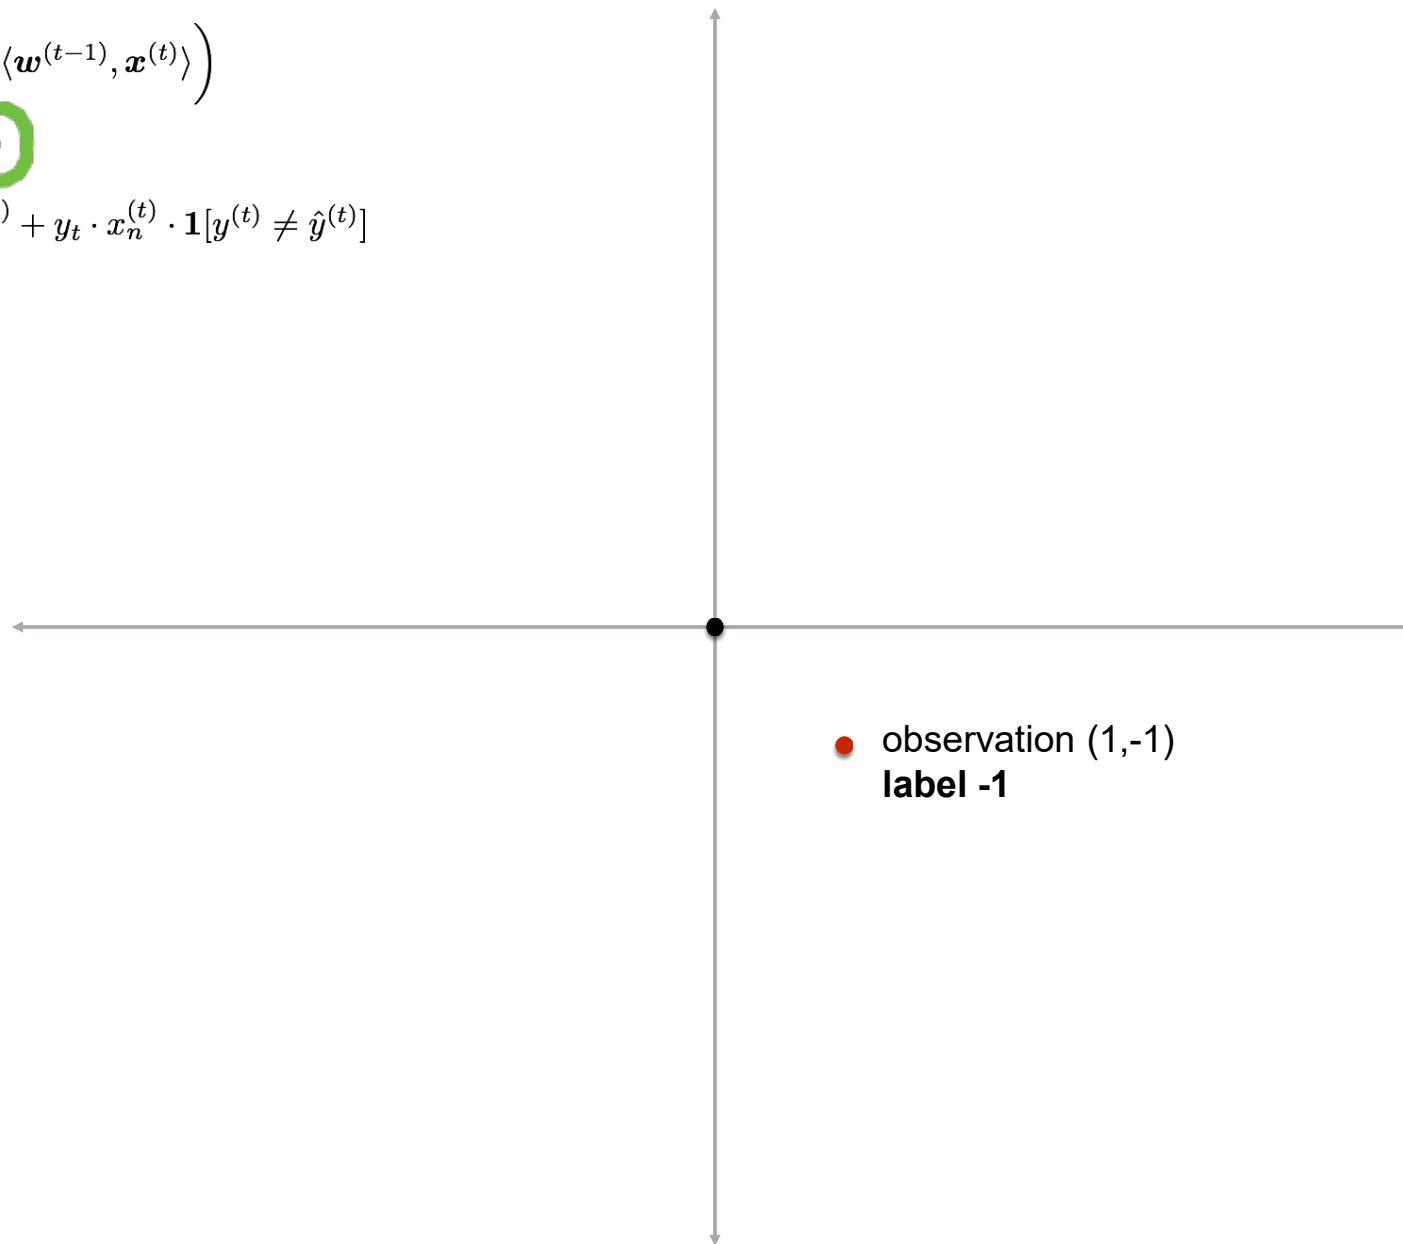
$$\begin{aligned}\hat{y}_A^{(t)} &= \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right) \\ &= 1\end{aligned}$$

RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$



RECEIVE( $\mathbf{x}^{(t)}$ )

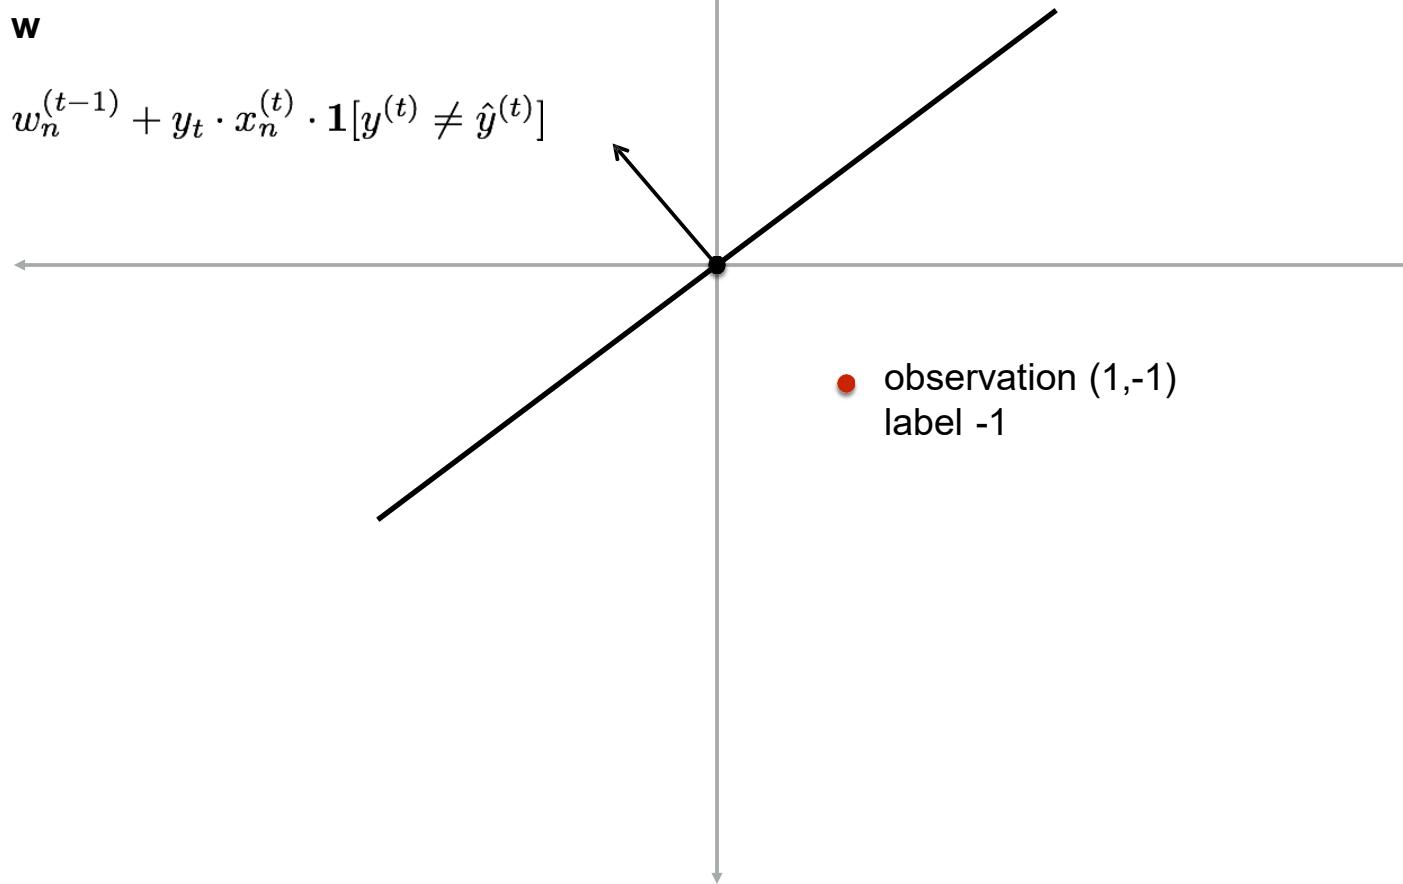
$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

**update w**

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$



RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

**update w**

no match!

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

(-1,1)

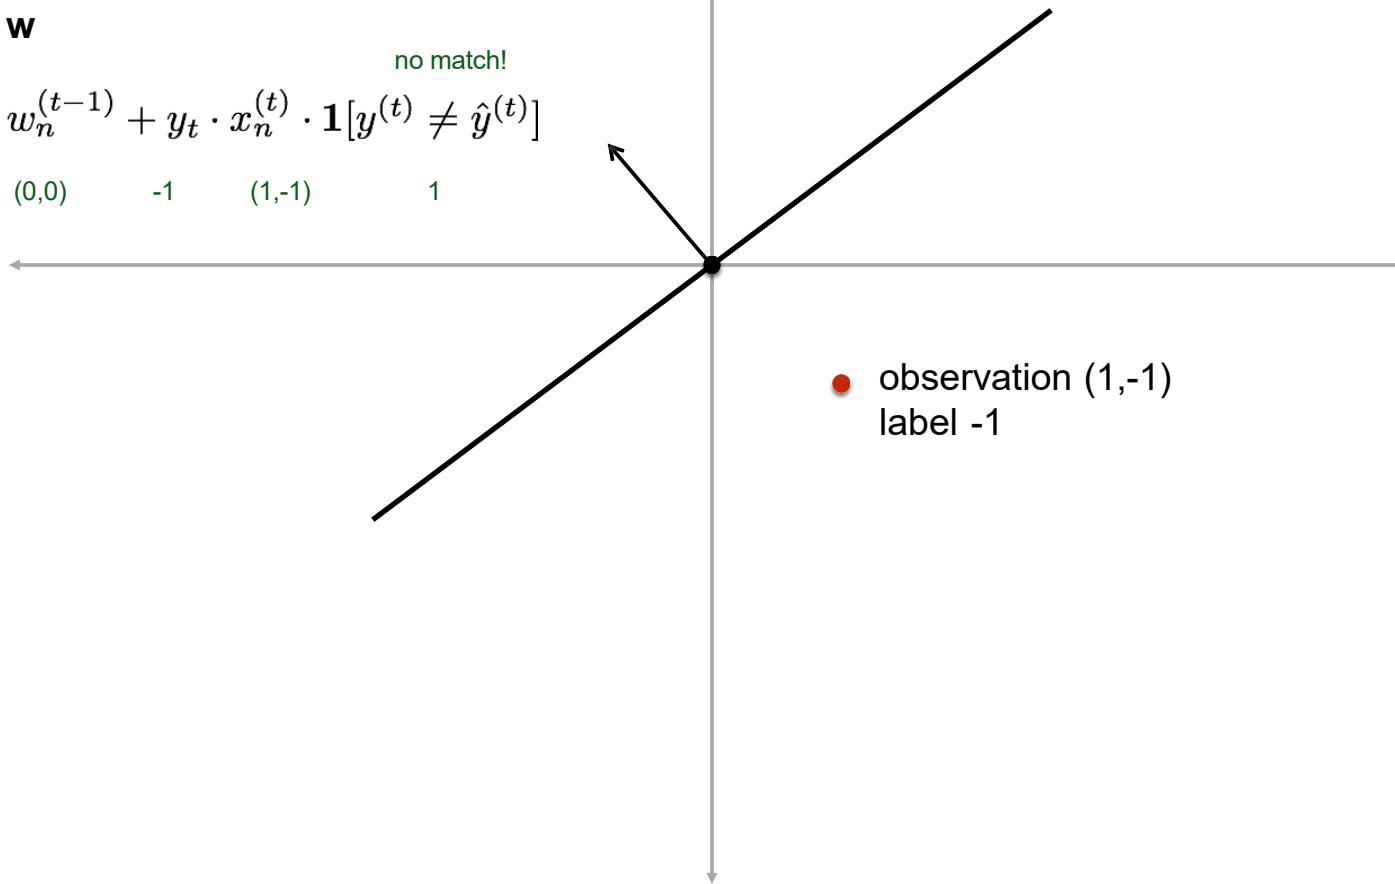
(0,0)

-1

(1,-1)

1

● observation (1,-1)  
label -1



RECEIVE( $\mathbf{x}^{(t)}$ )

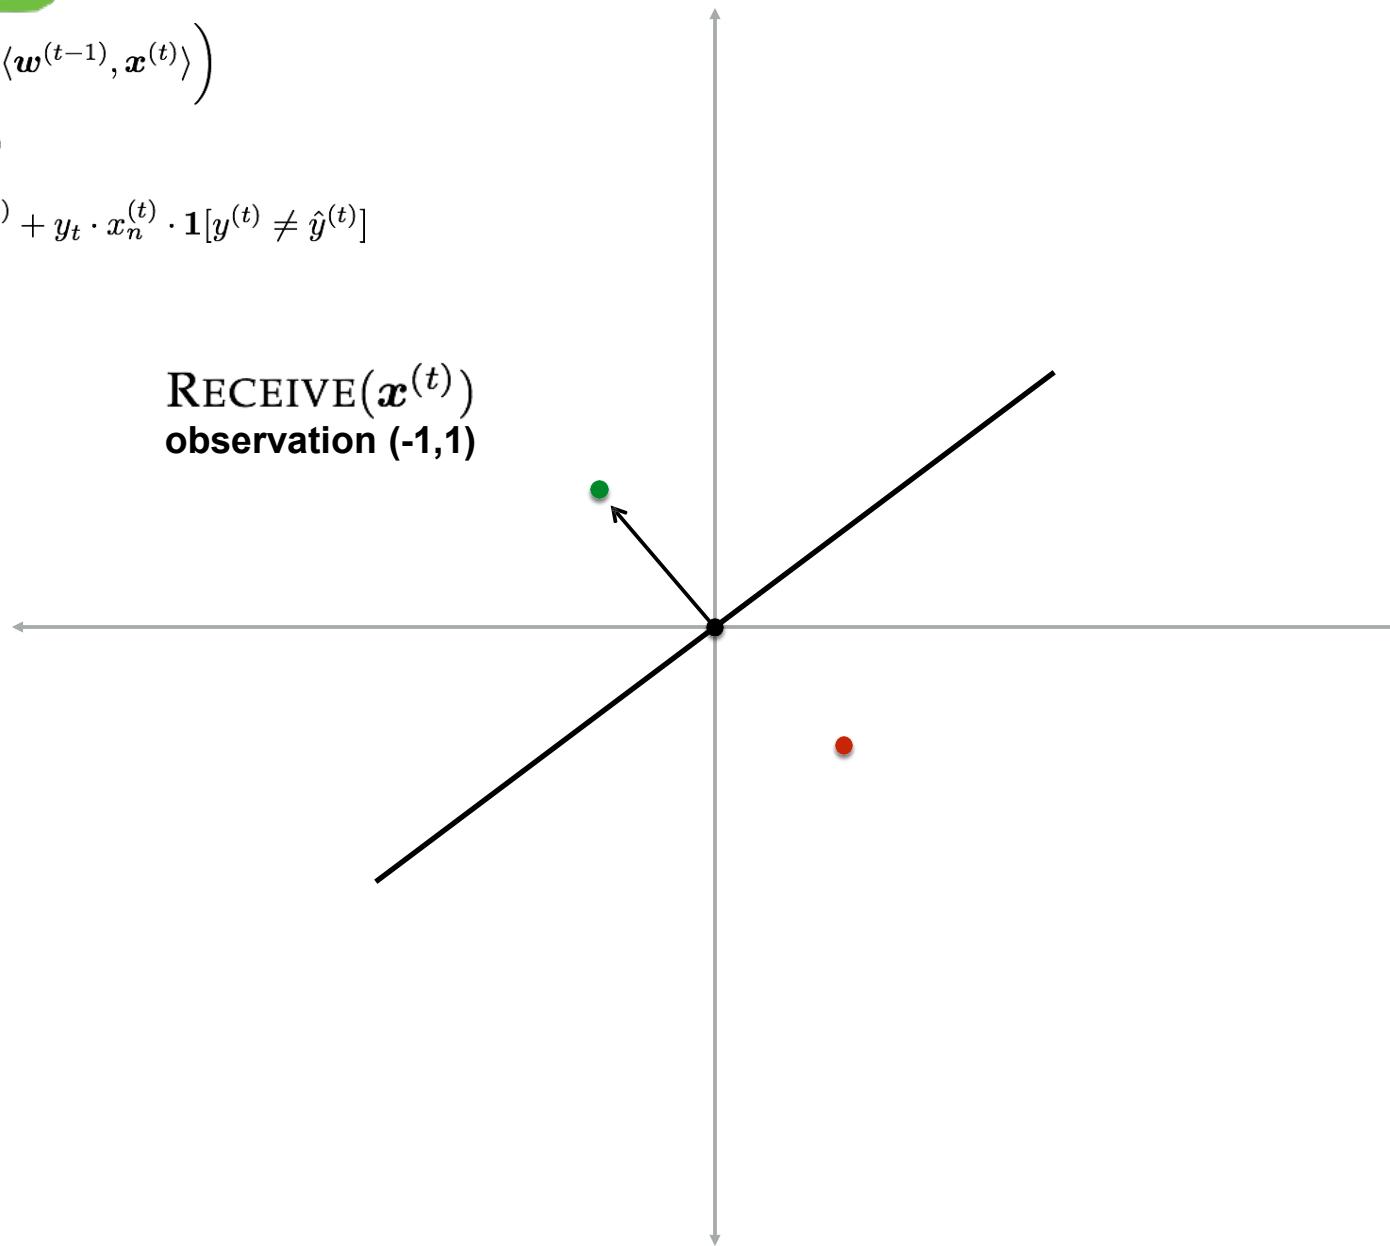
$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

(-1,1)

RECEIVE( $\mathbf{x}^{(t)}$ )  
observation (-1,1)



RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

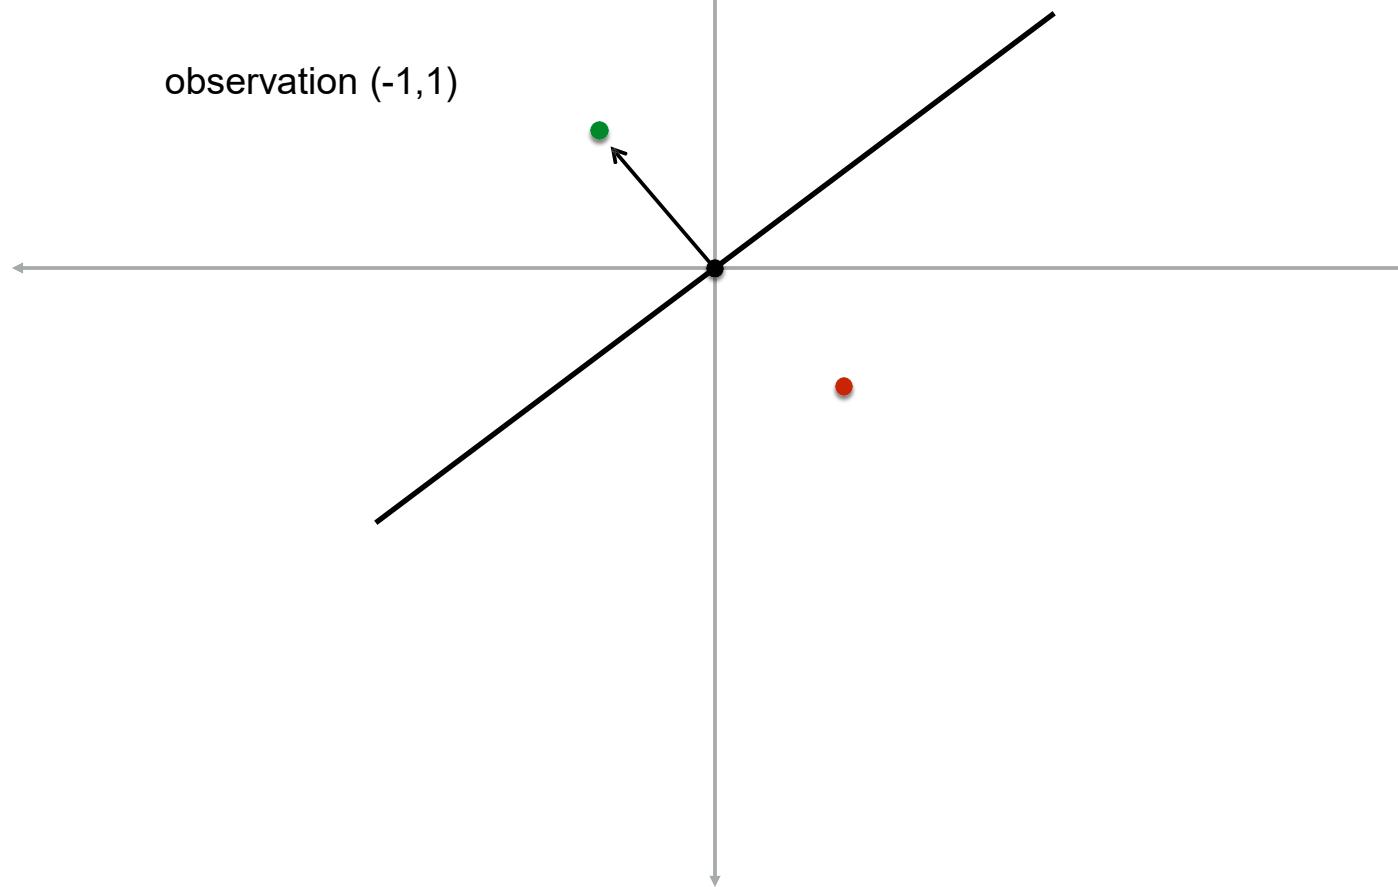
RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

(-1,1)      (-1,1)

$$\begin{aligned}\hat{y}_A^{(t)} &= \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right) \\ &= 1\end{aligned}$$

observation (-1,1)



RECEIVE( $\mathbf{x}^{(t)}$ )

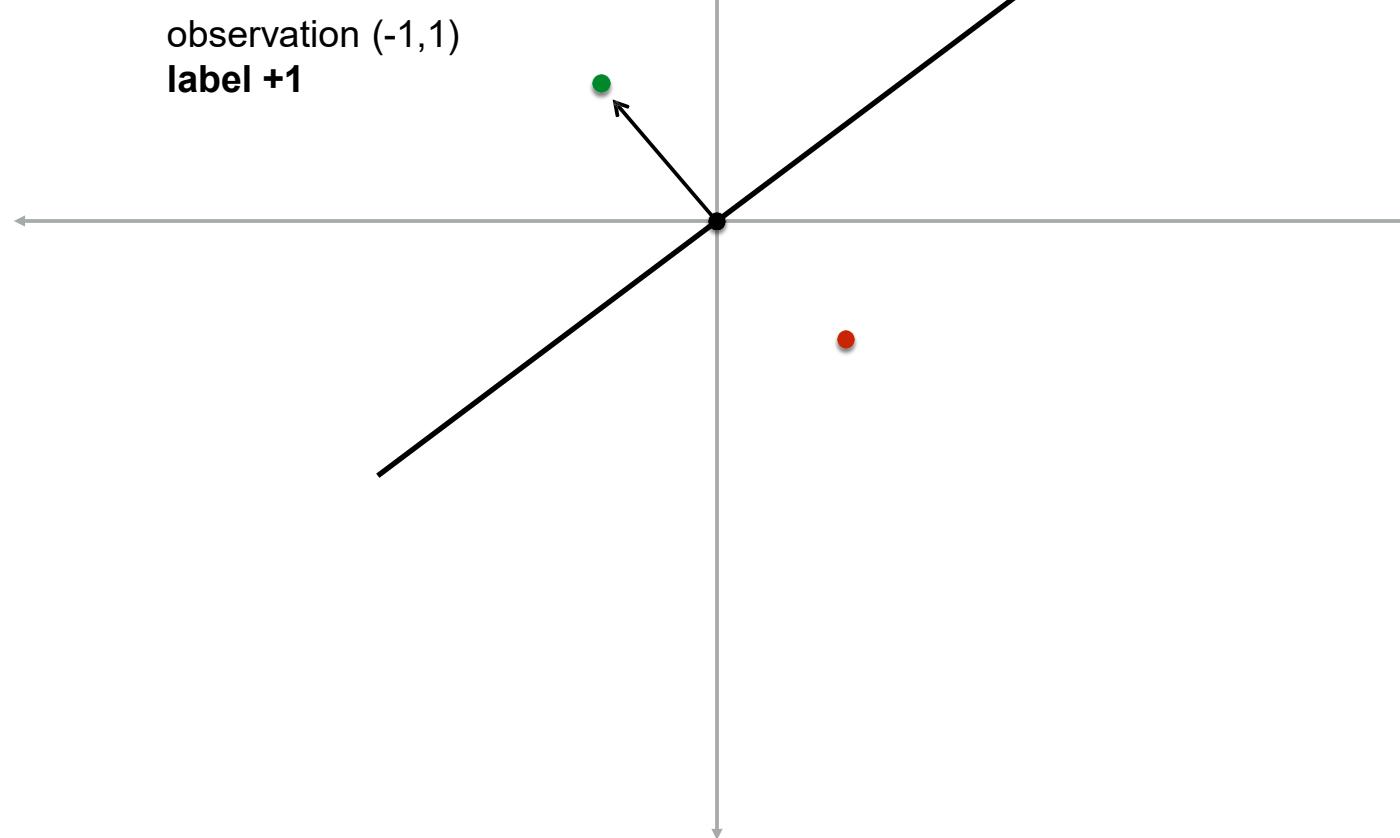
$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

(-1,1)      (-1,1)

$$\begin{aligned}\hat{y}_A^{(t)} &= \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right) \\ &= 1\end{aligned}$$



RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

update  $\mathbf{w}$

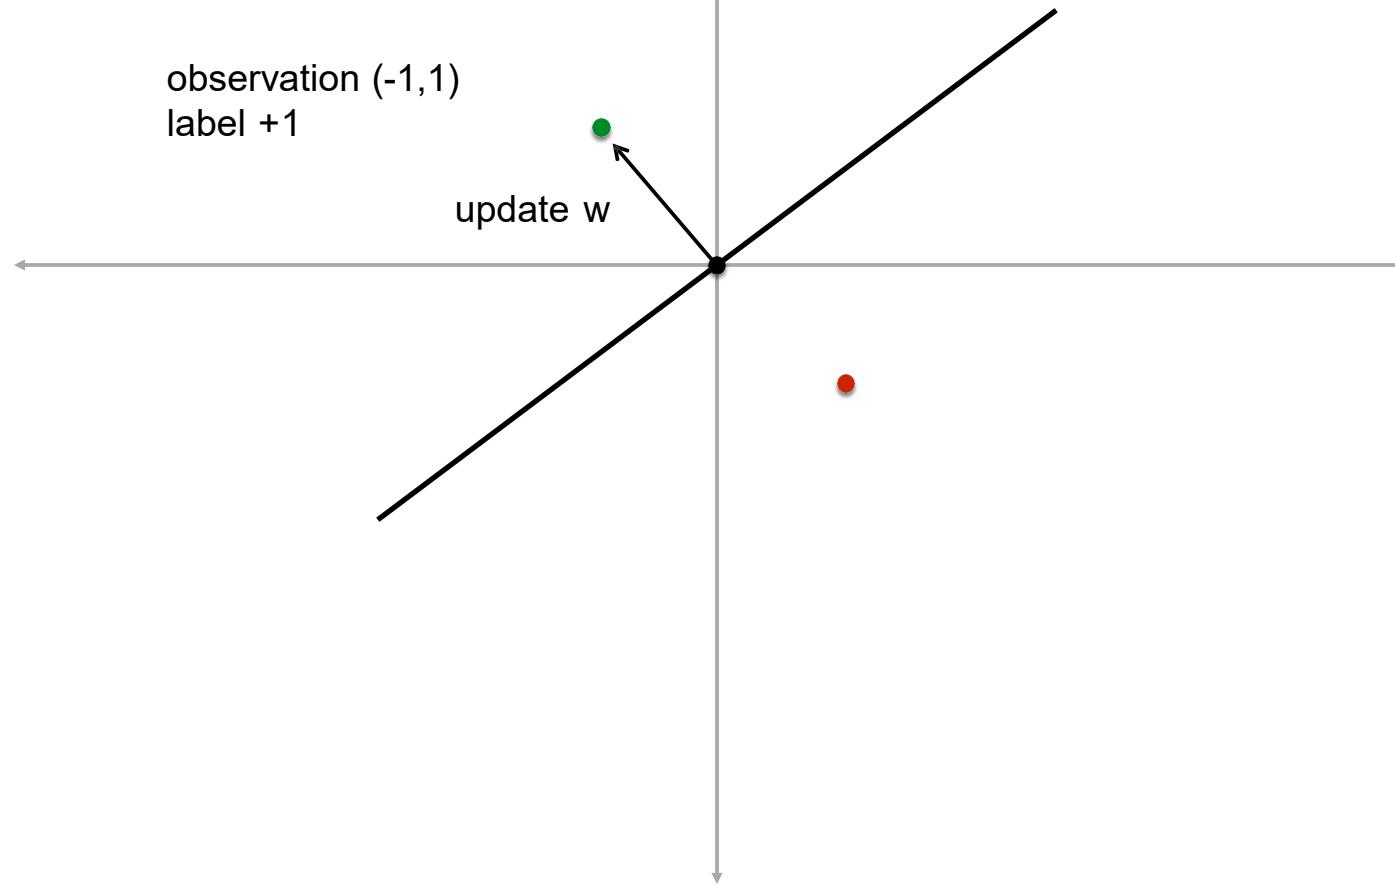
$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

match!

$$(-1,1) \quad (-1,1) \quad +1 \quad (-1,1) \quad 0$$

observation (-1,1)  
label +1

update  $\mathbf{w}$

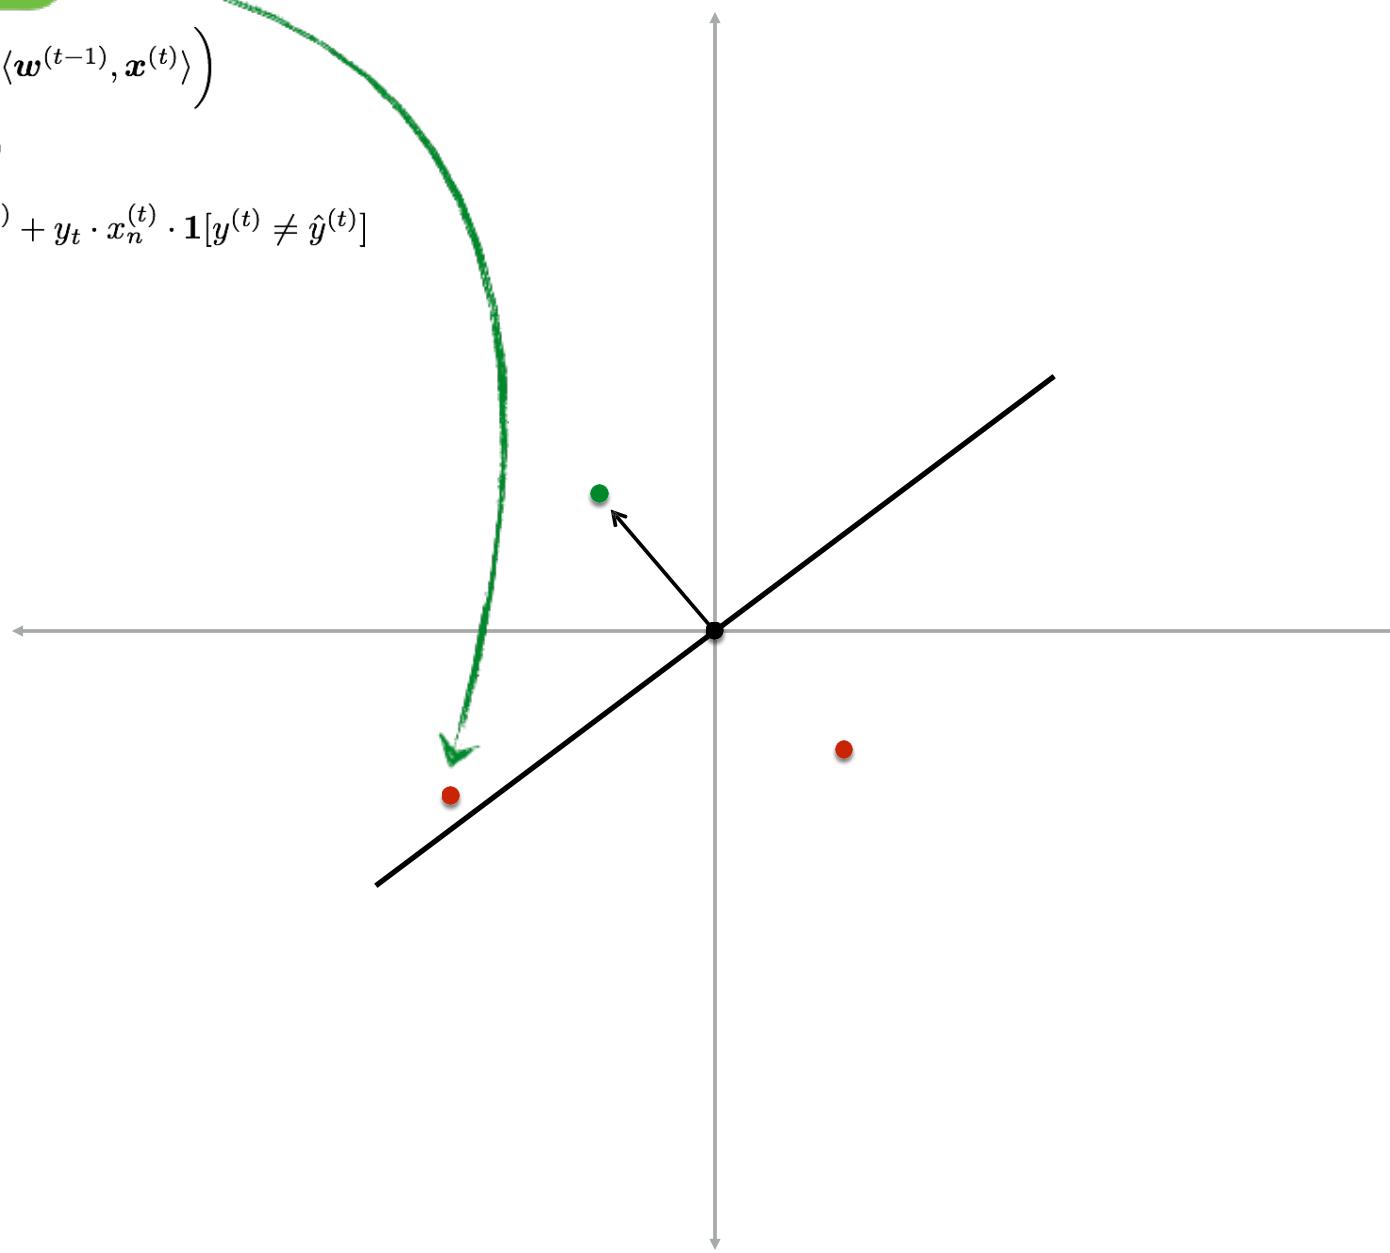


RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

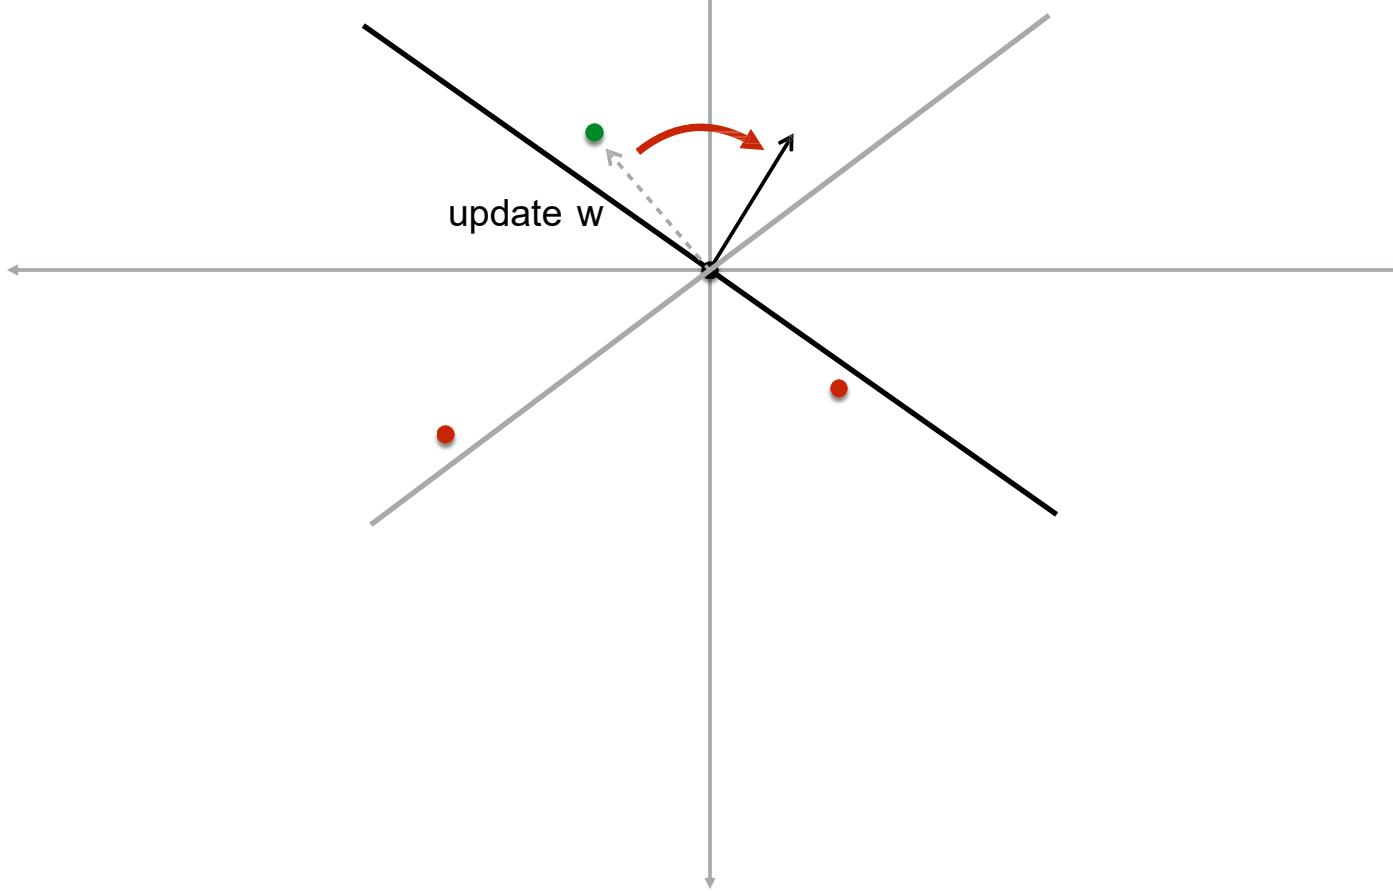


RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

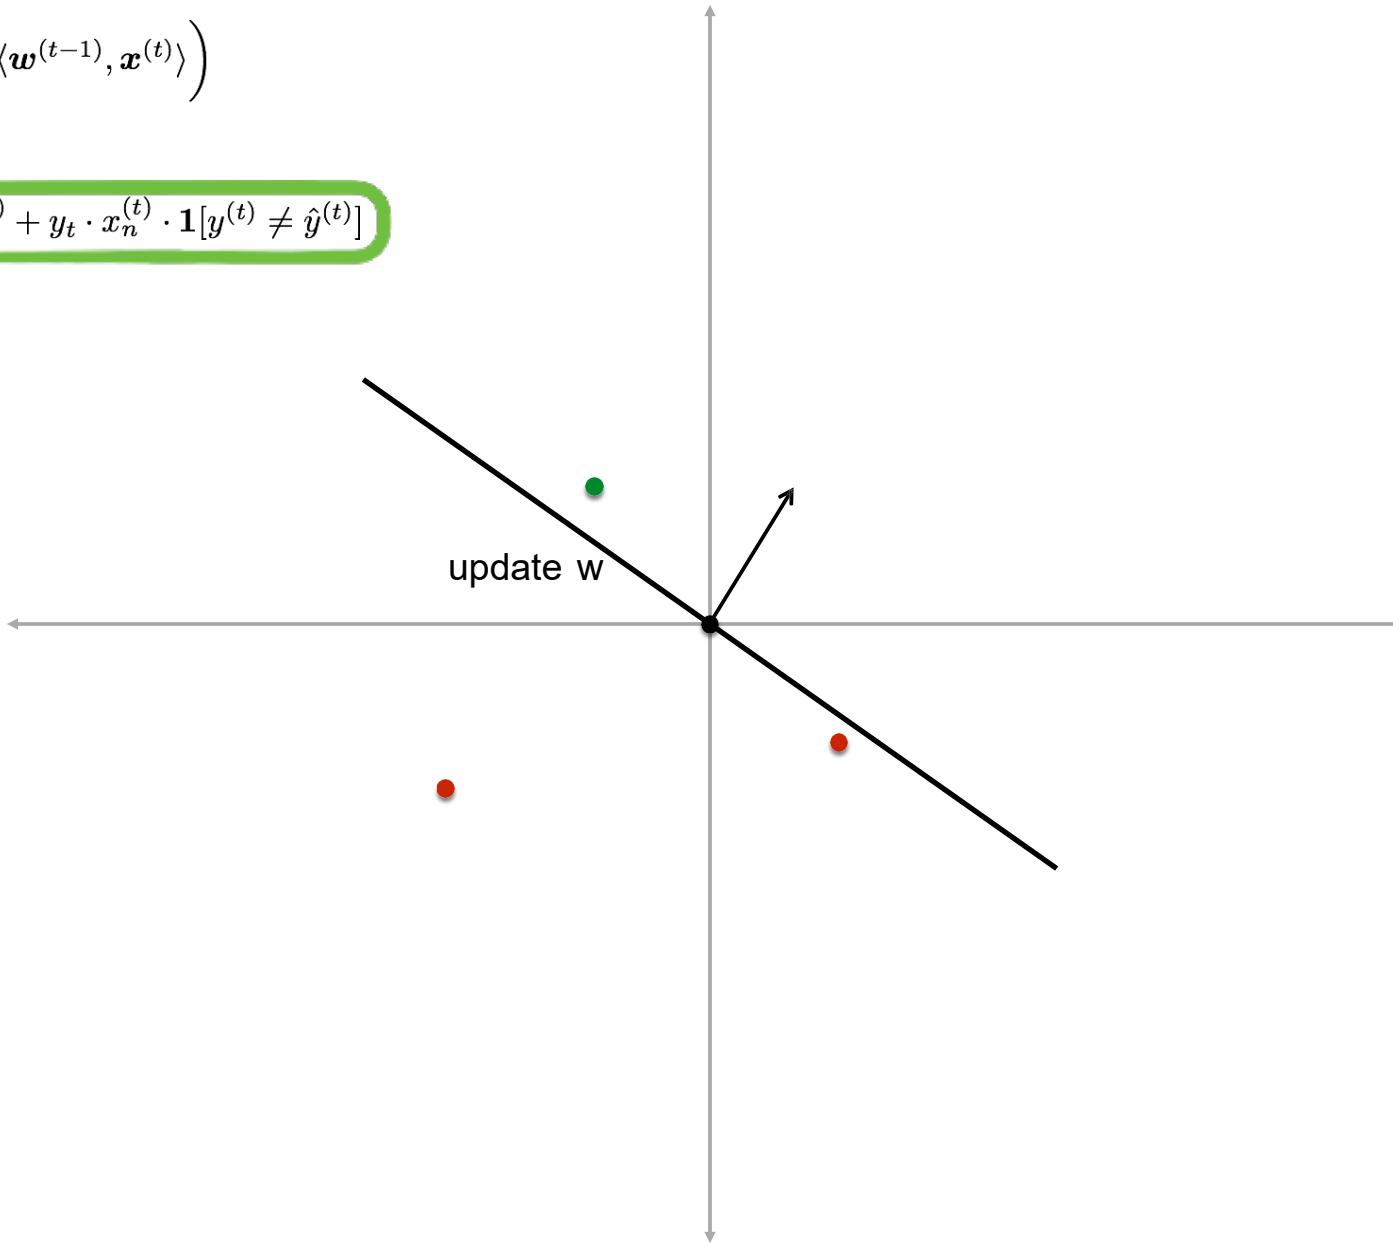


RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

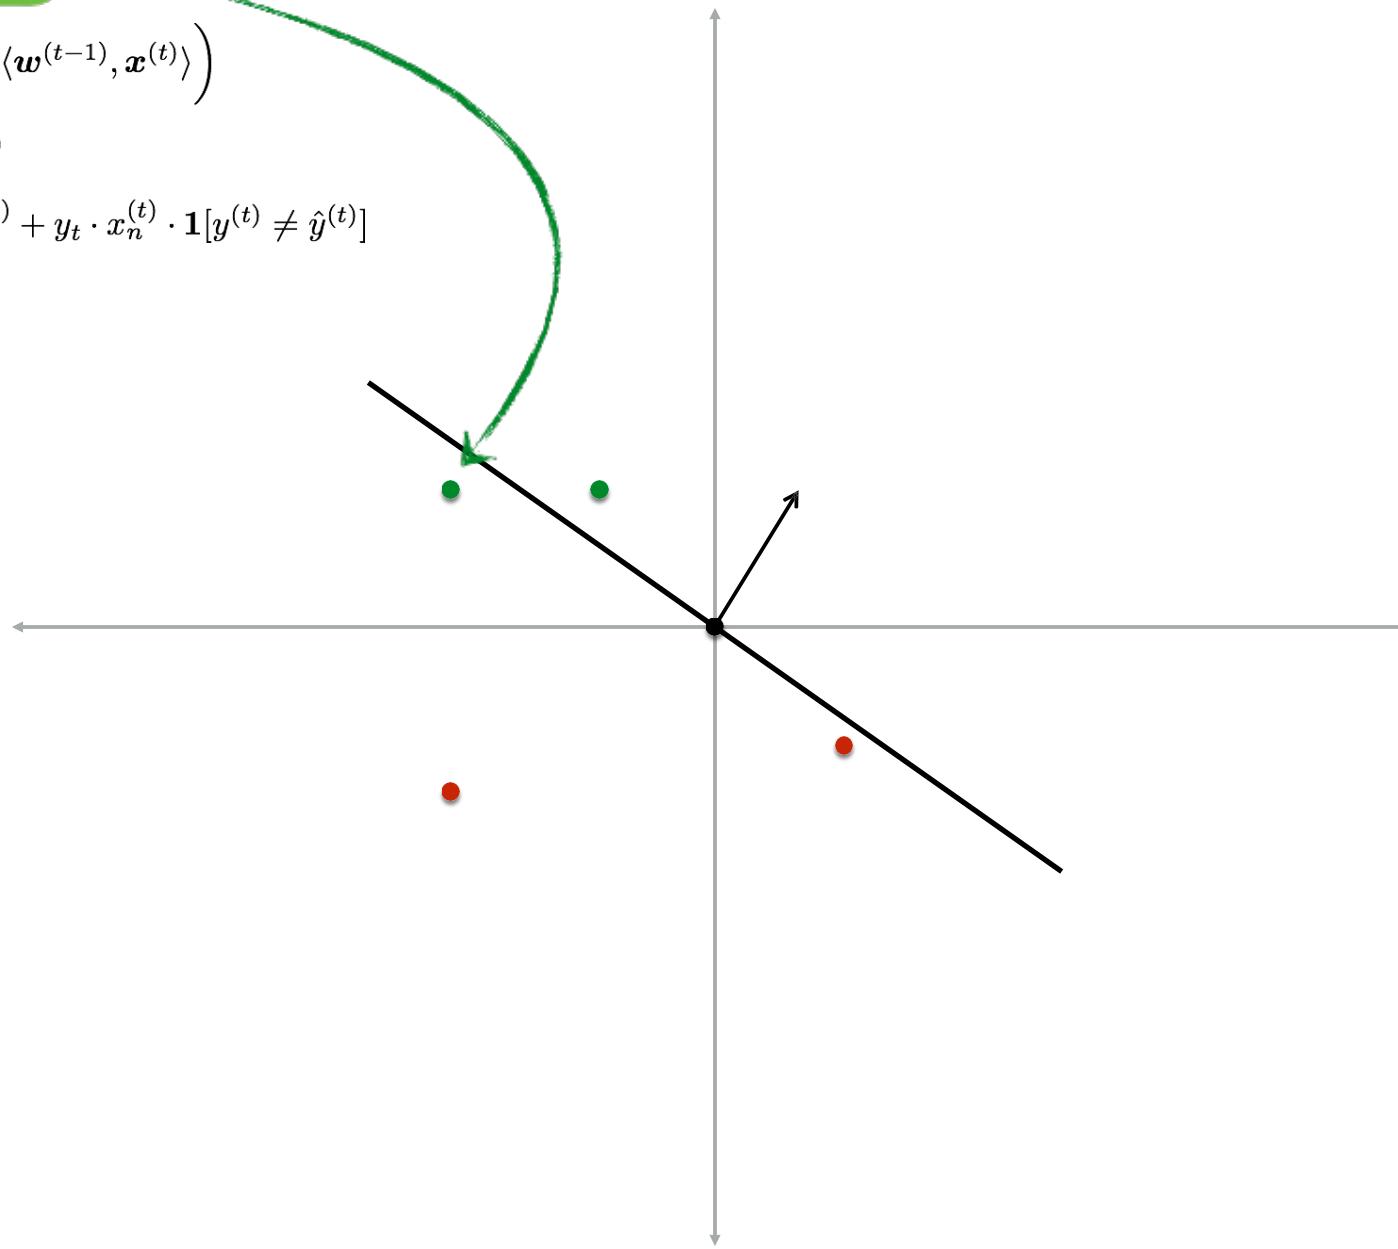


RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

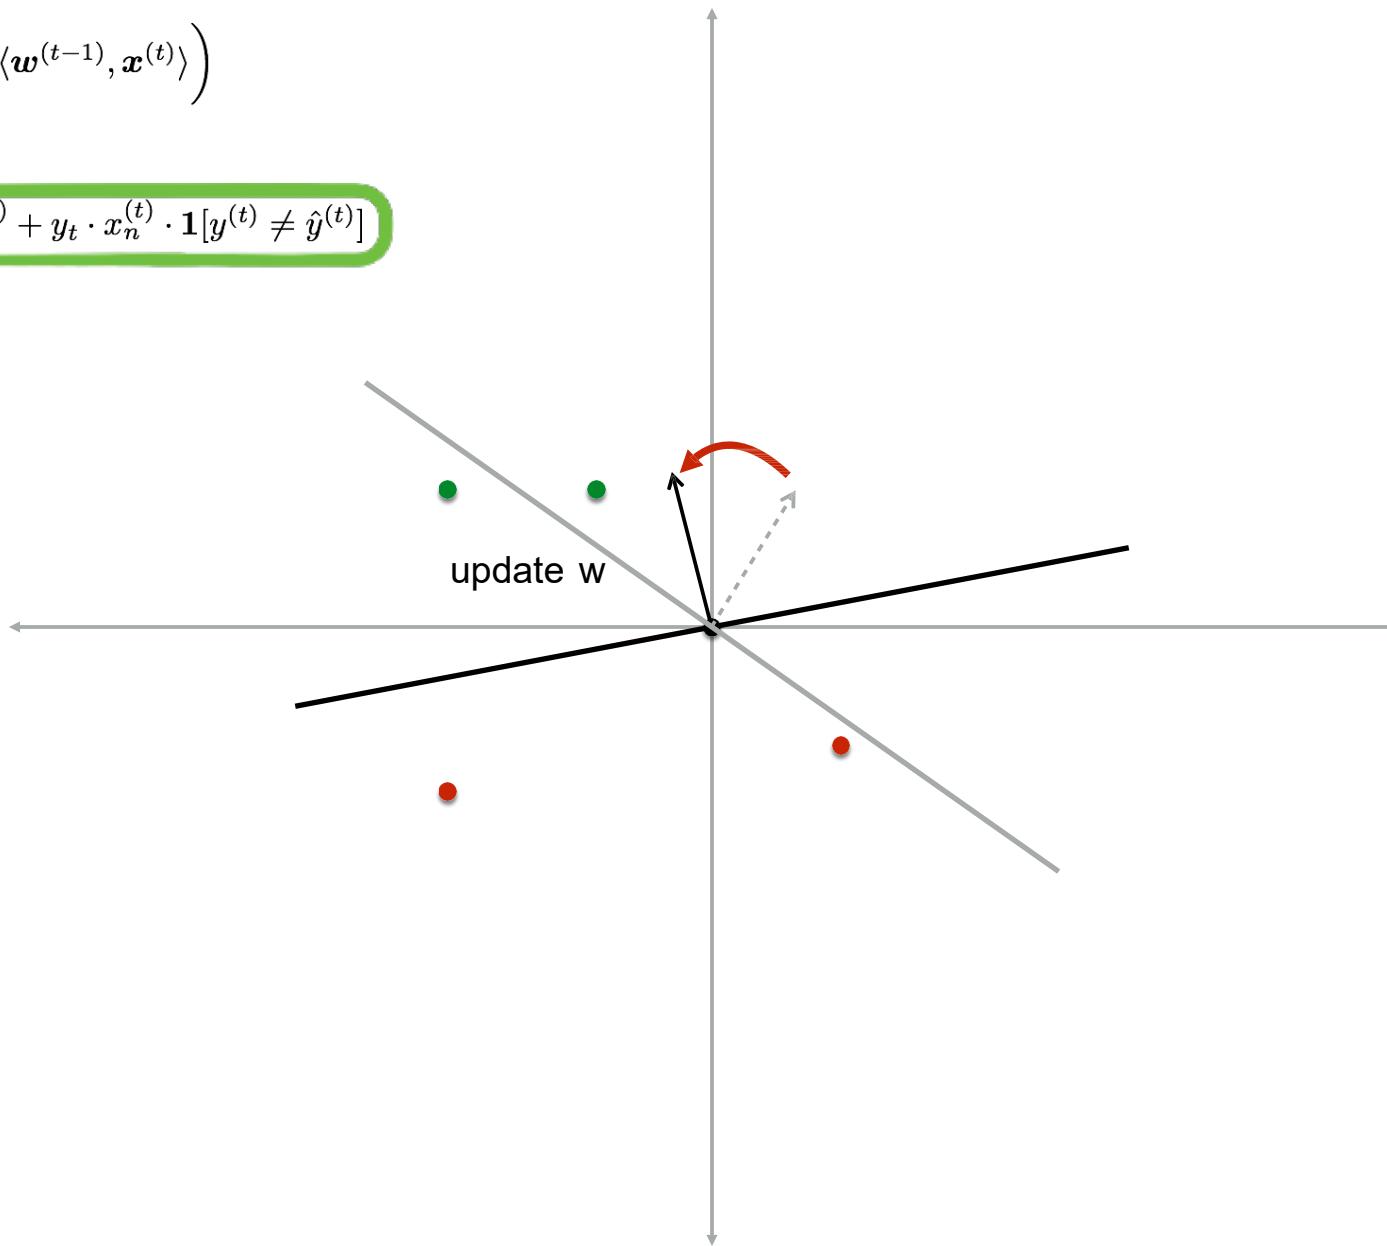


RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

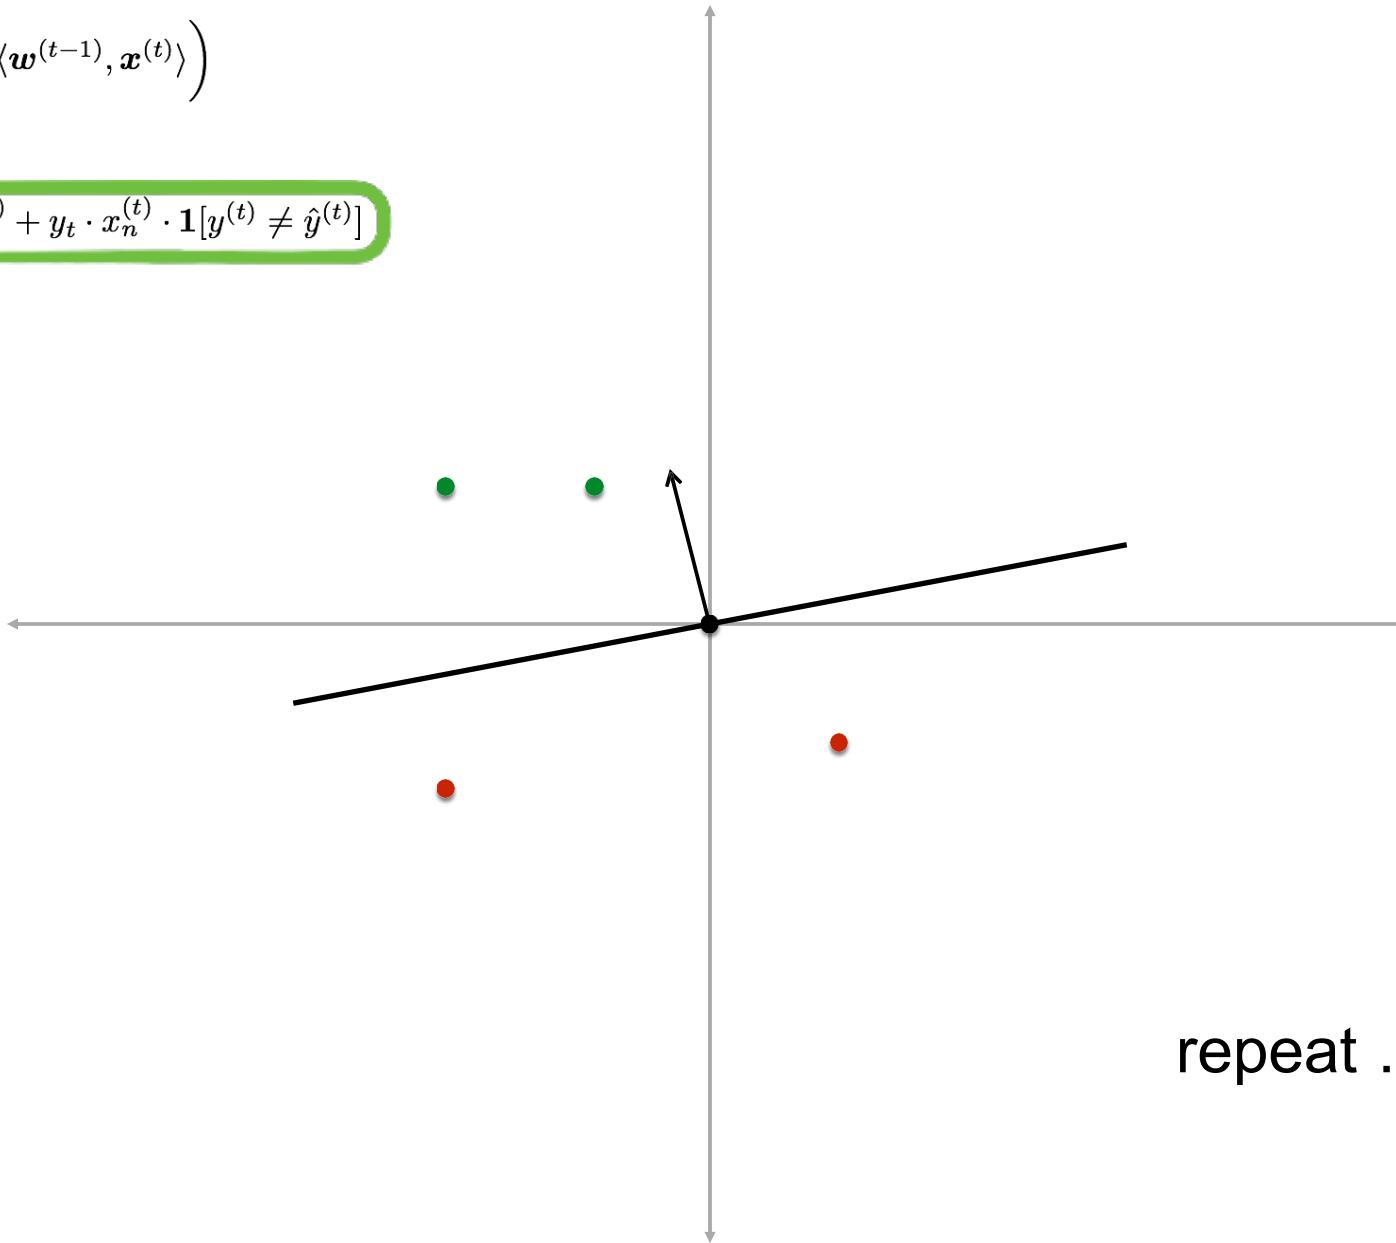


RECEIVE( $\mathbf{x}^{(t)}$ )

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

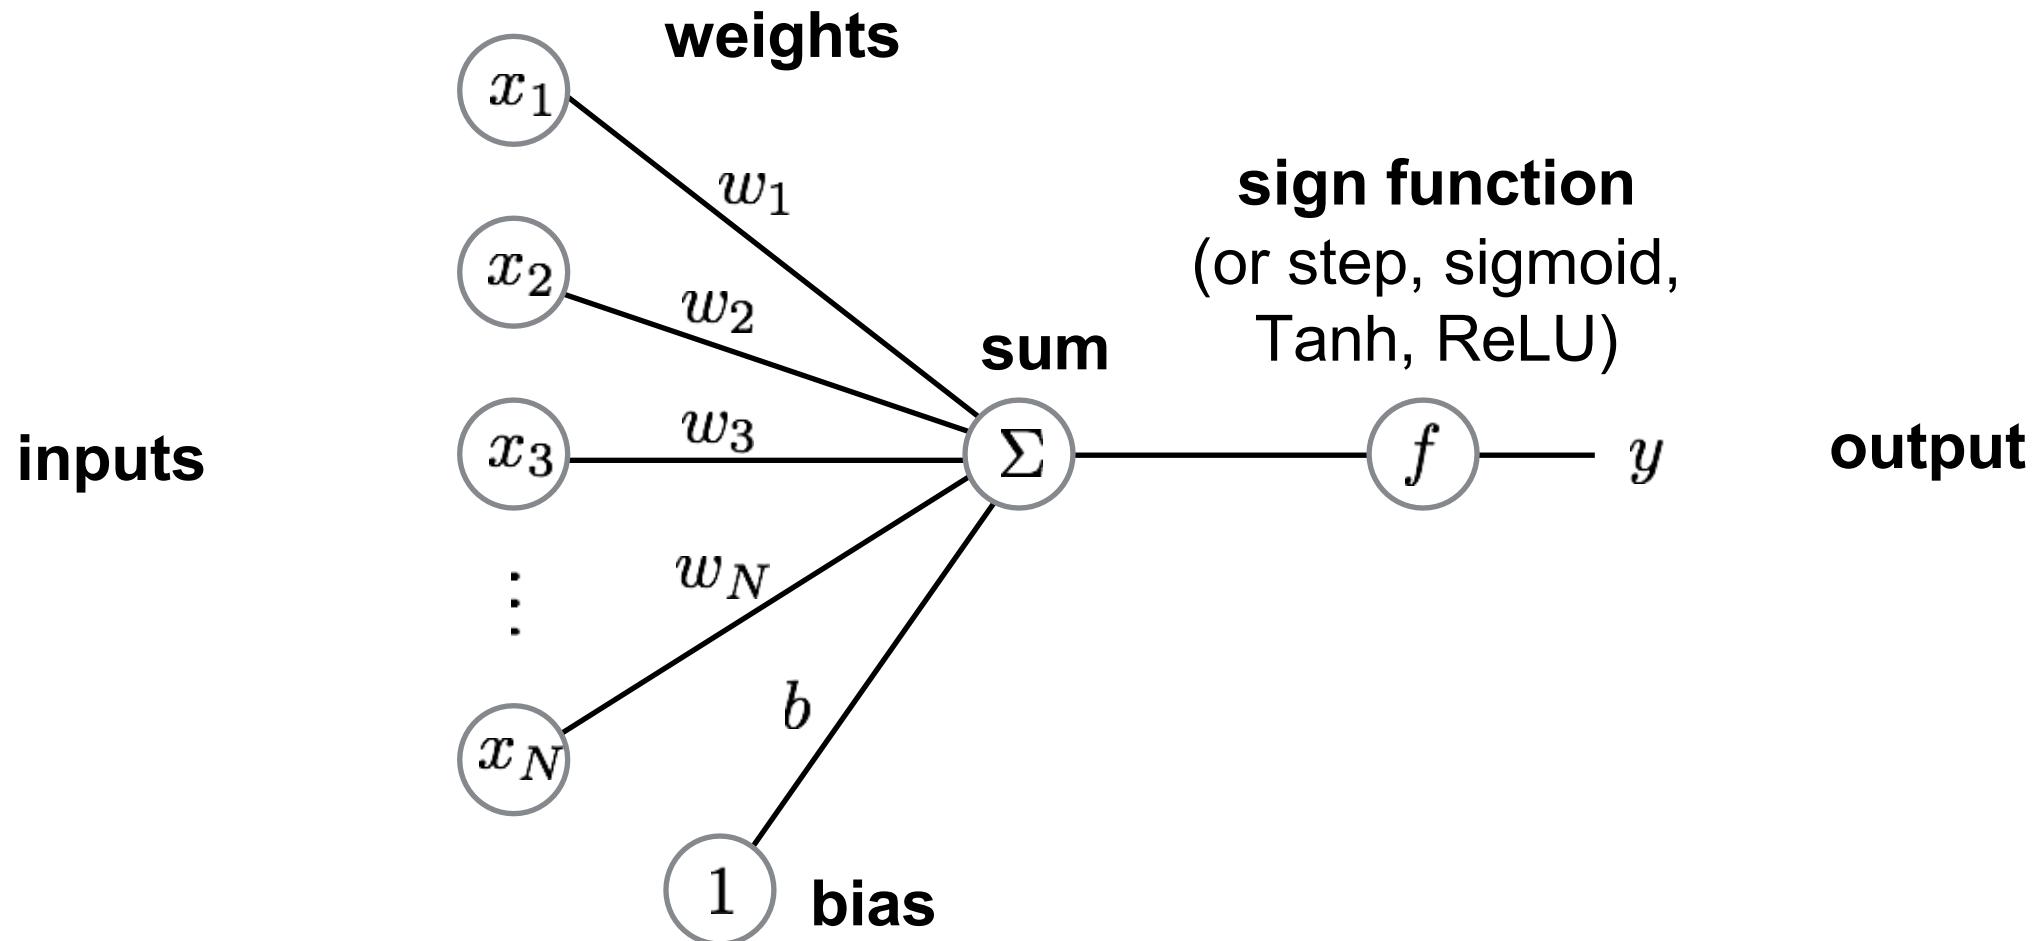
RECEIVE( $y^t$ )

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

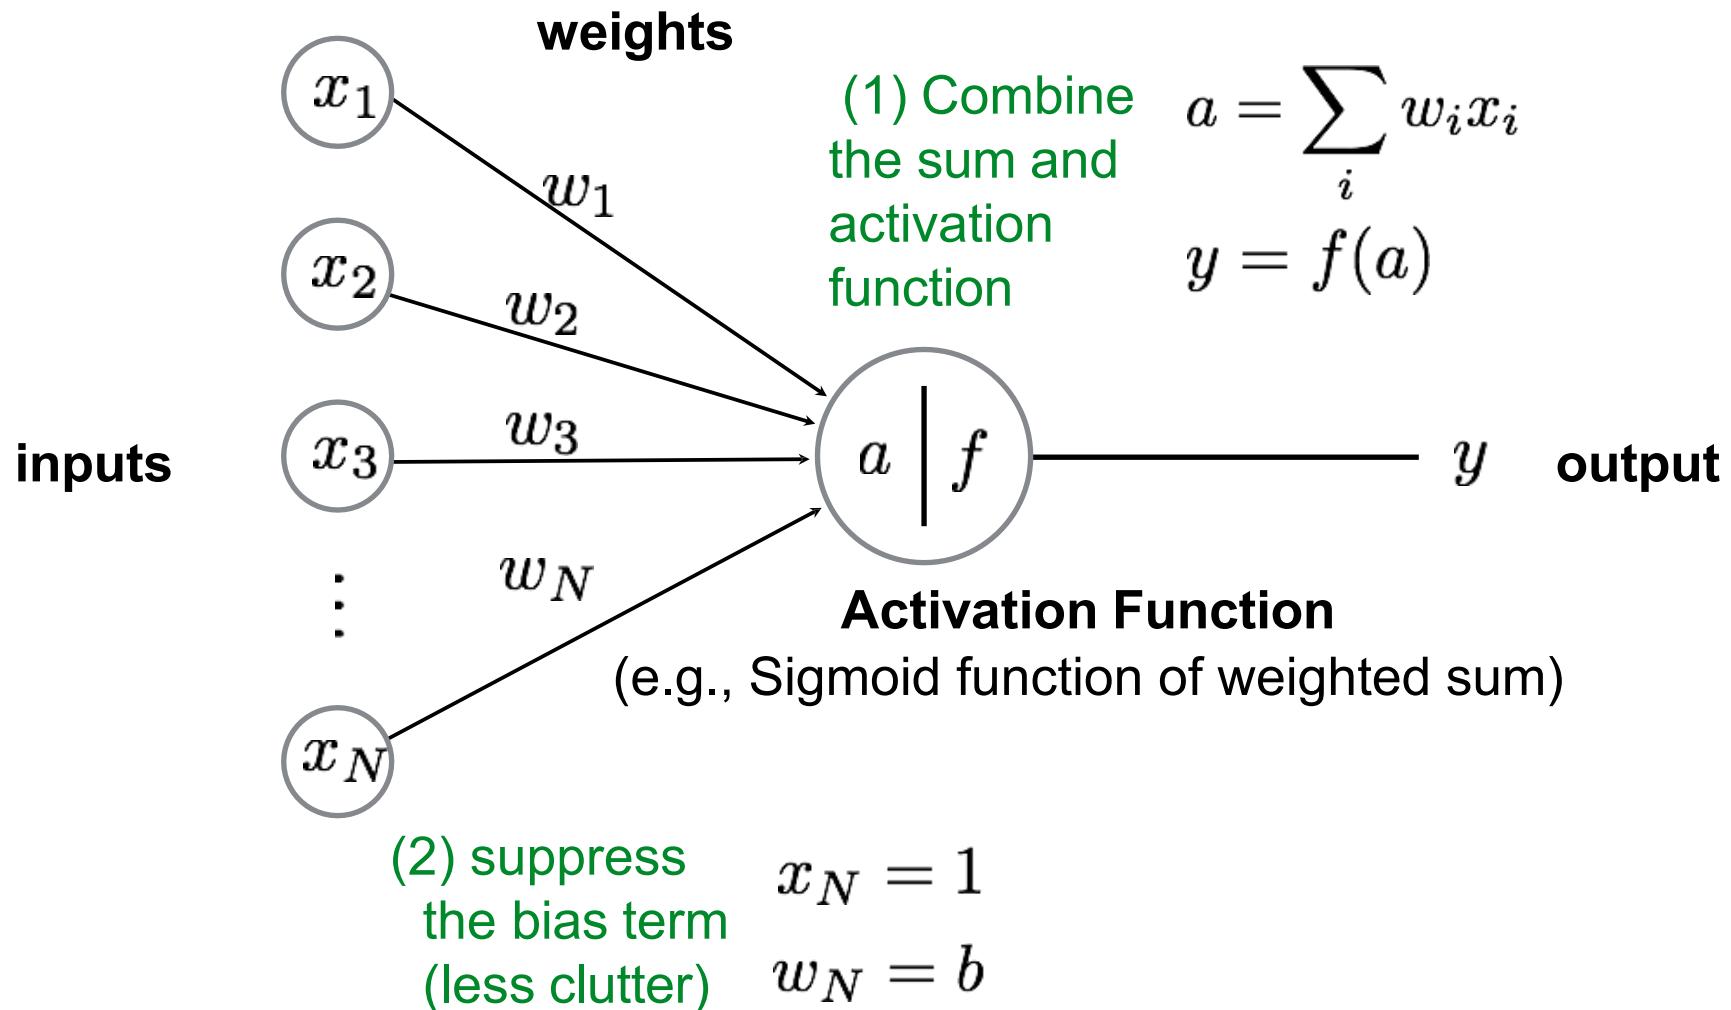


repeat ...

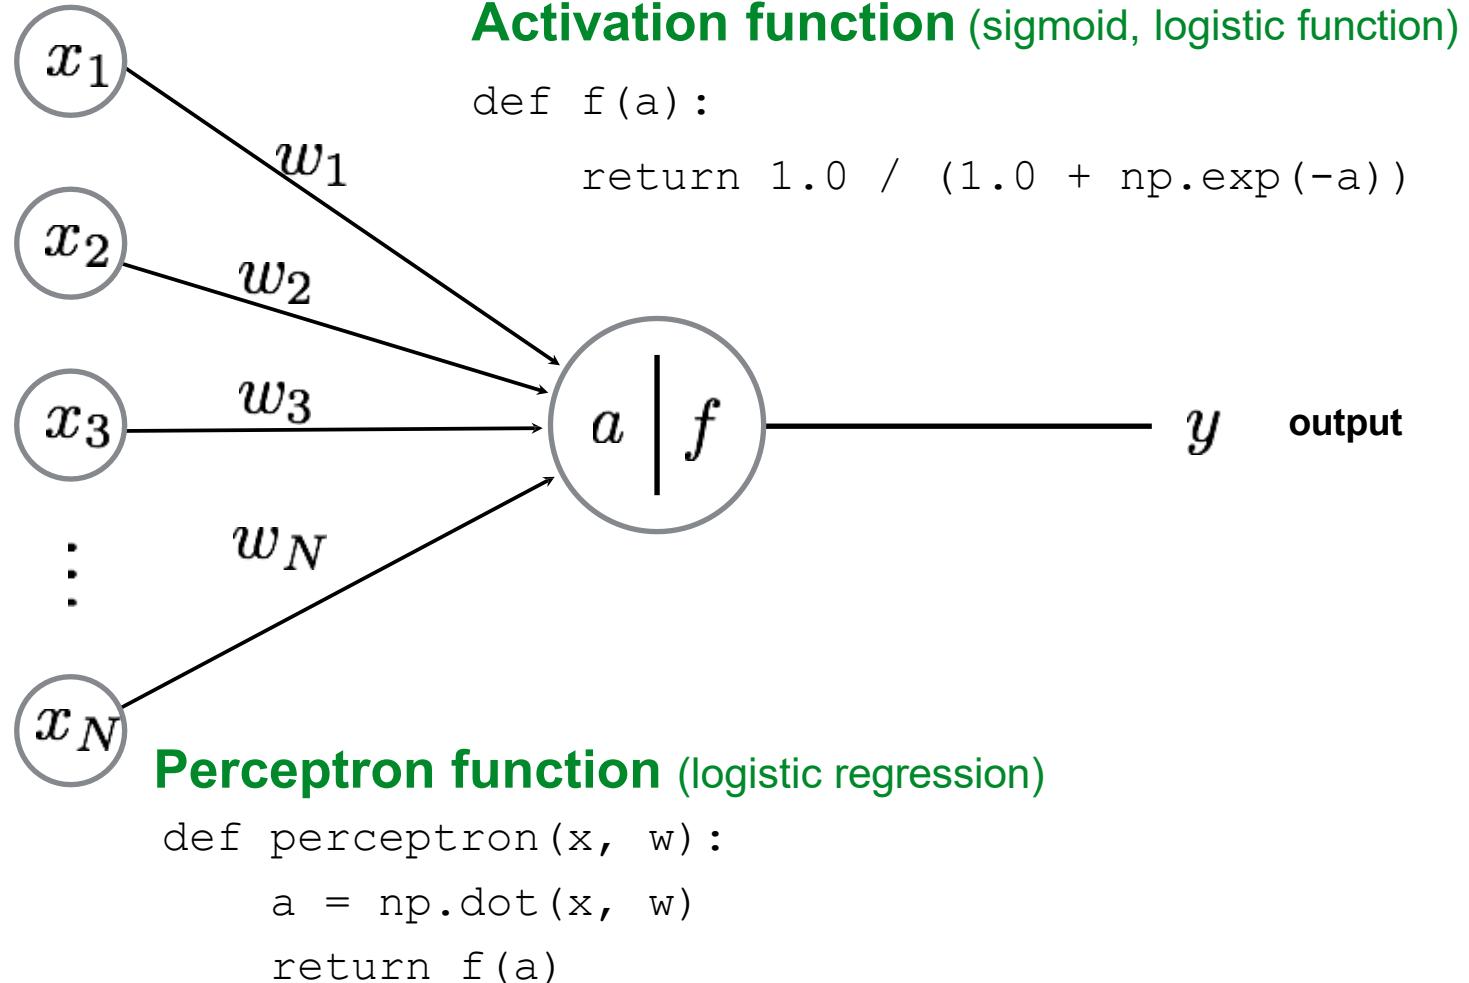
# The Perceptron



# Another way to draw it...



# Programming the 'forward pass'

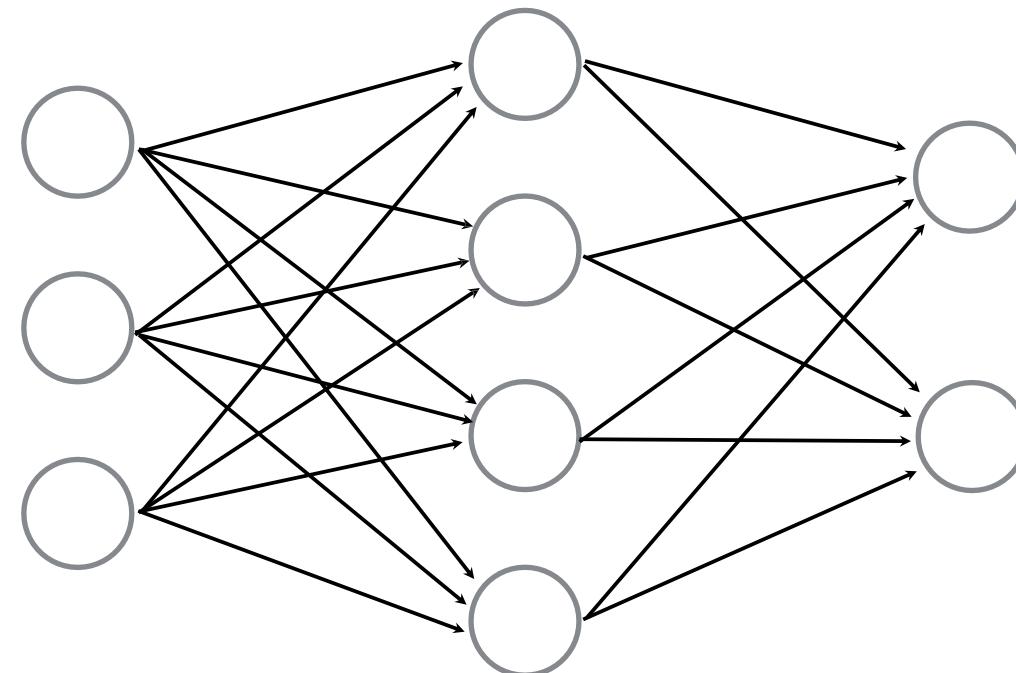


# Neural networks

Connect a bunch of perceptrons together ...

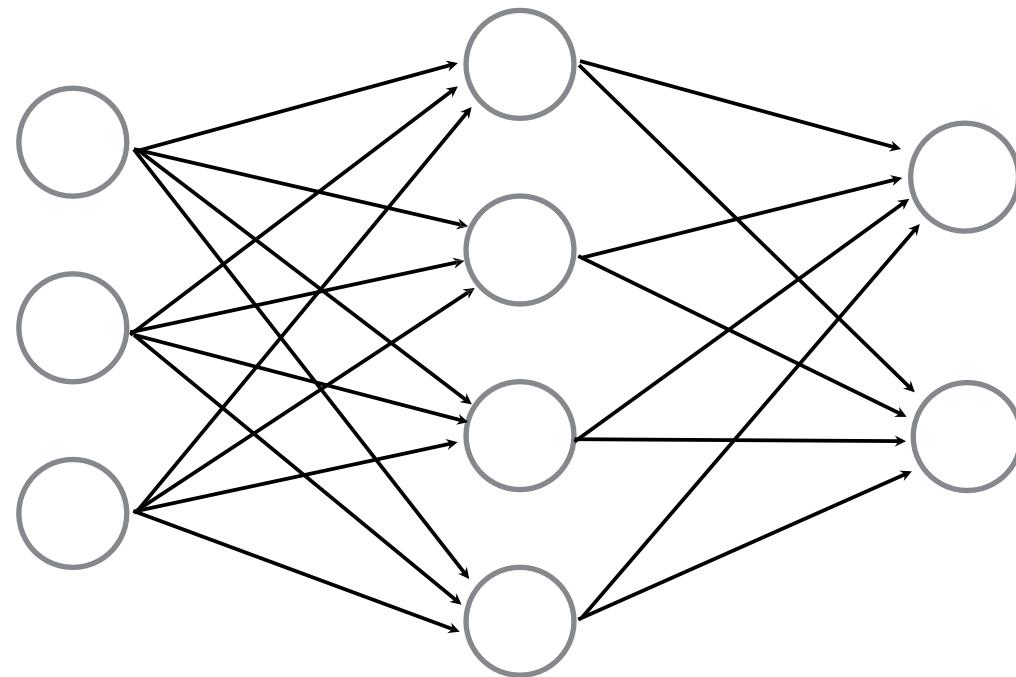
# Neural Network

a collection of connected perceptrons



# Neural Network

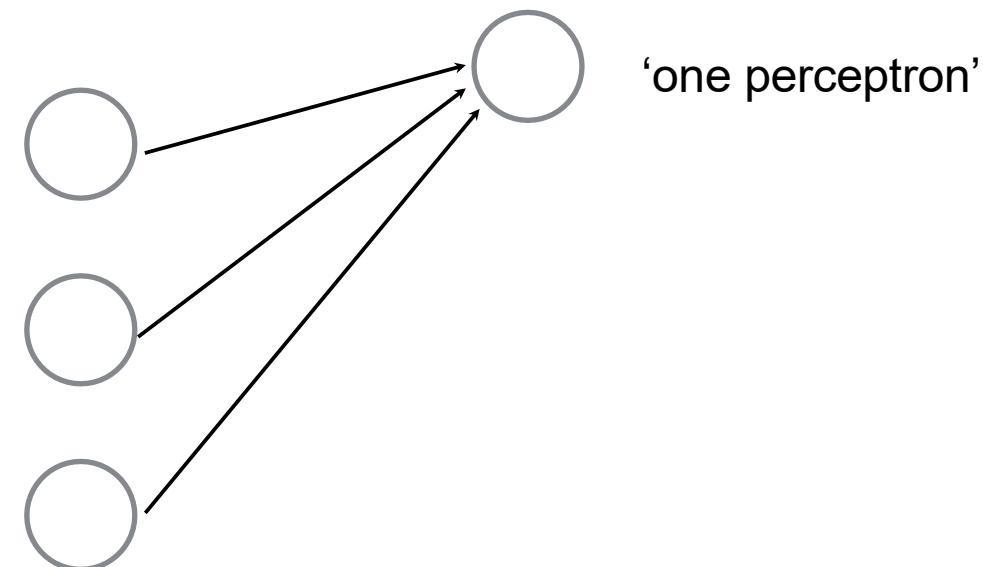
a collection of connected perceptrons



*How many perceptrons in this neural network?*

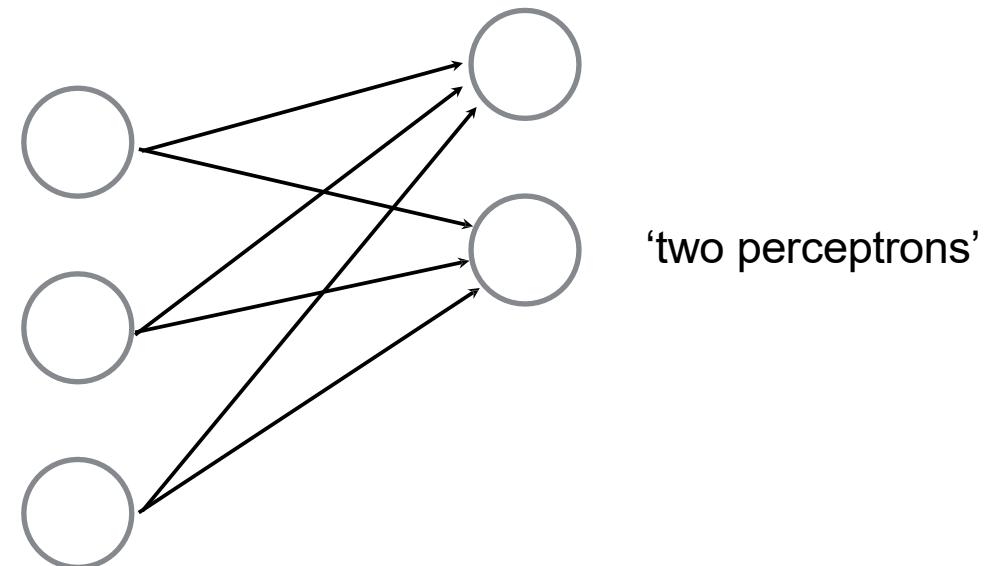
# Neural Network

a collection of connected perceptrons



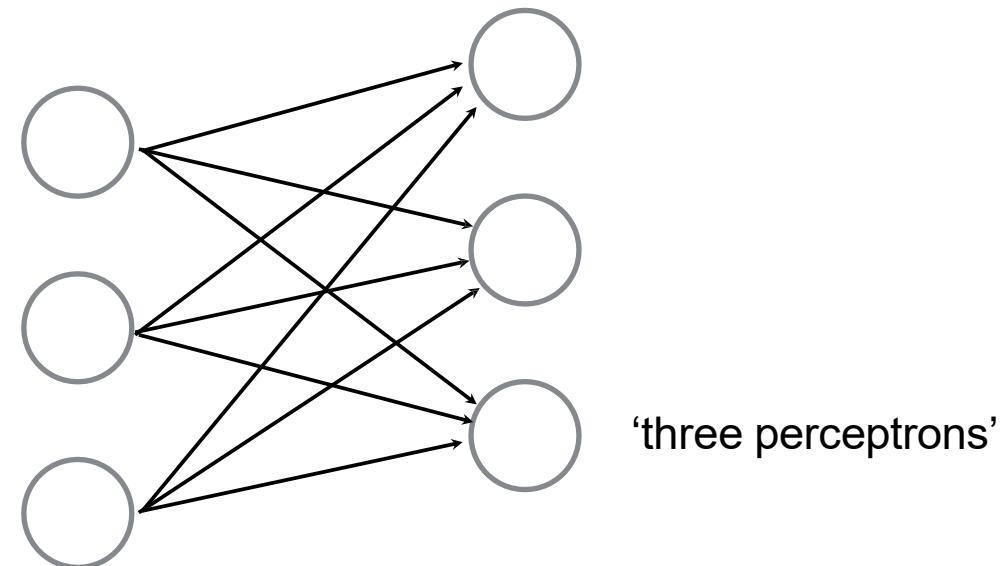
# Neural Network

a collection of connected perceptrons



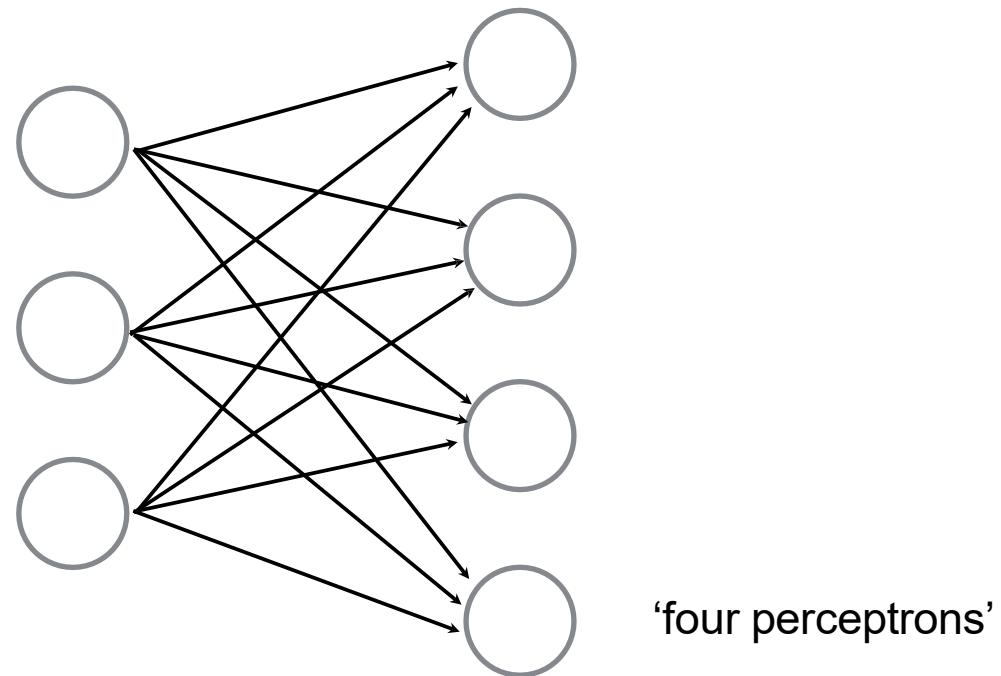
# Neural Network

a collection of connected perceptrons



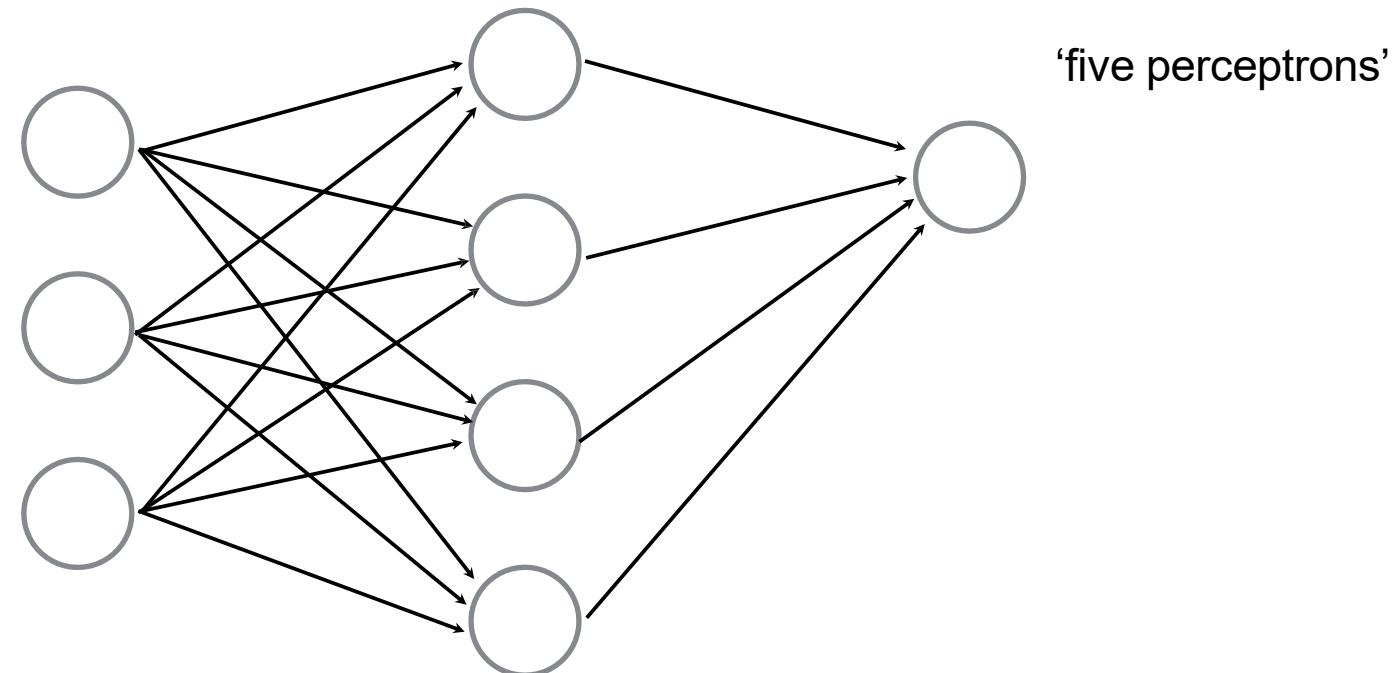
# Neural Network

a collection of connected perceptrons



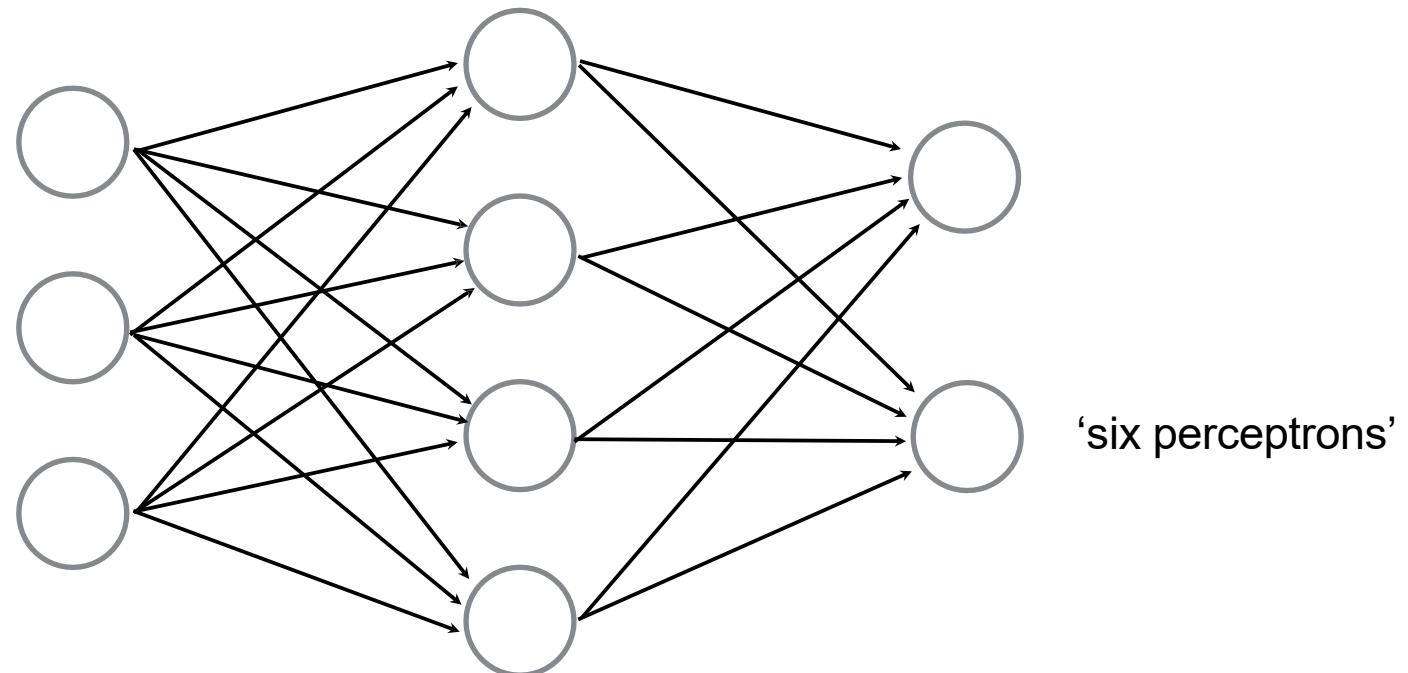
# Neural Network

a collection of connected perceptrons



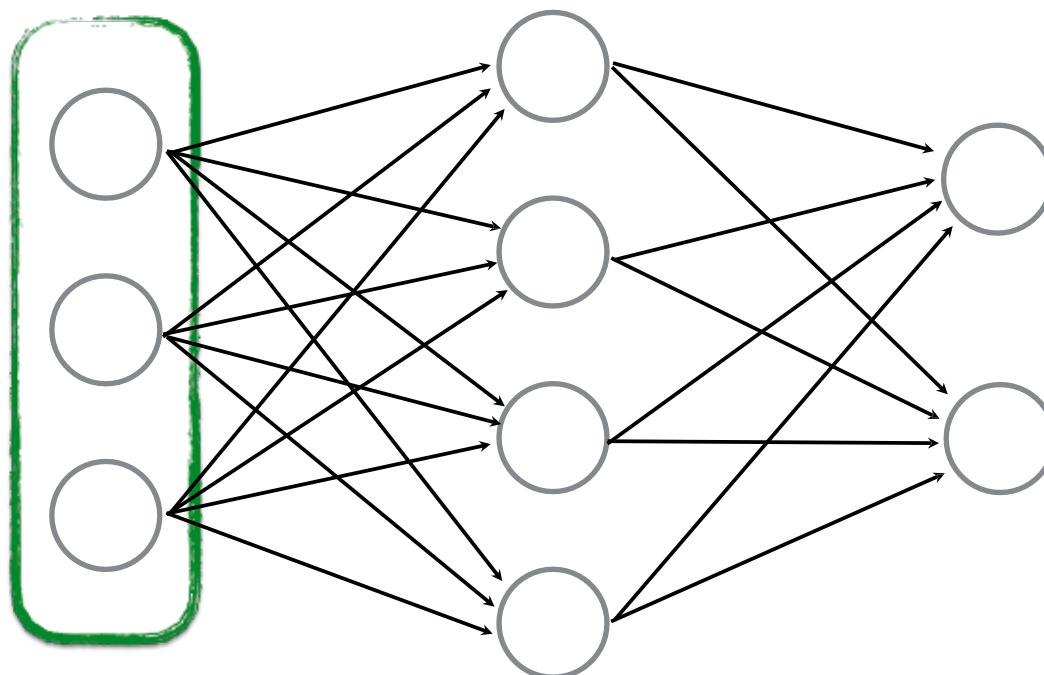
# Neural Network

a collection of connected perceptrons



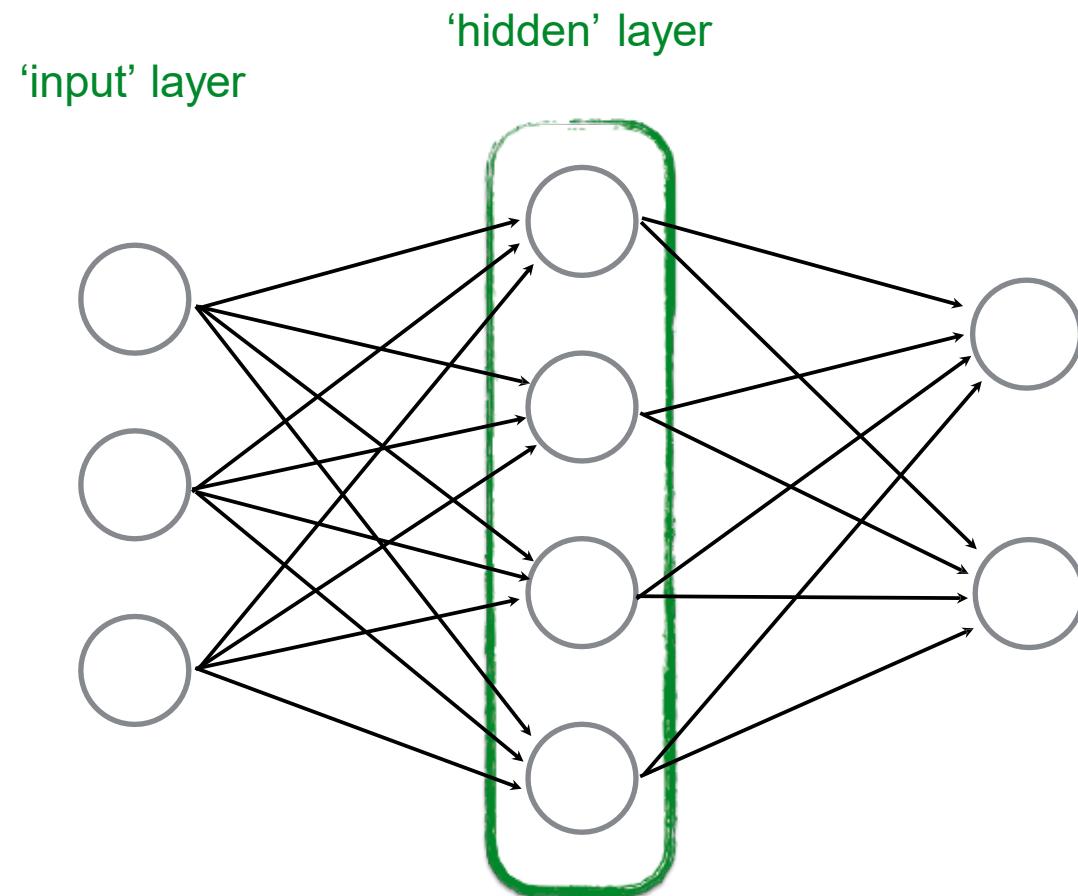
# Some terminology...

'input' layer



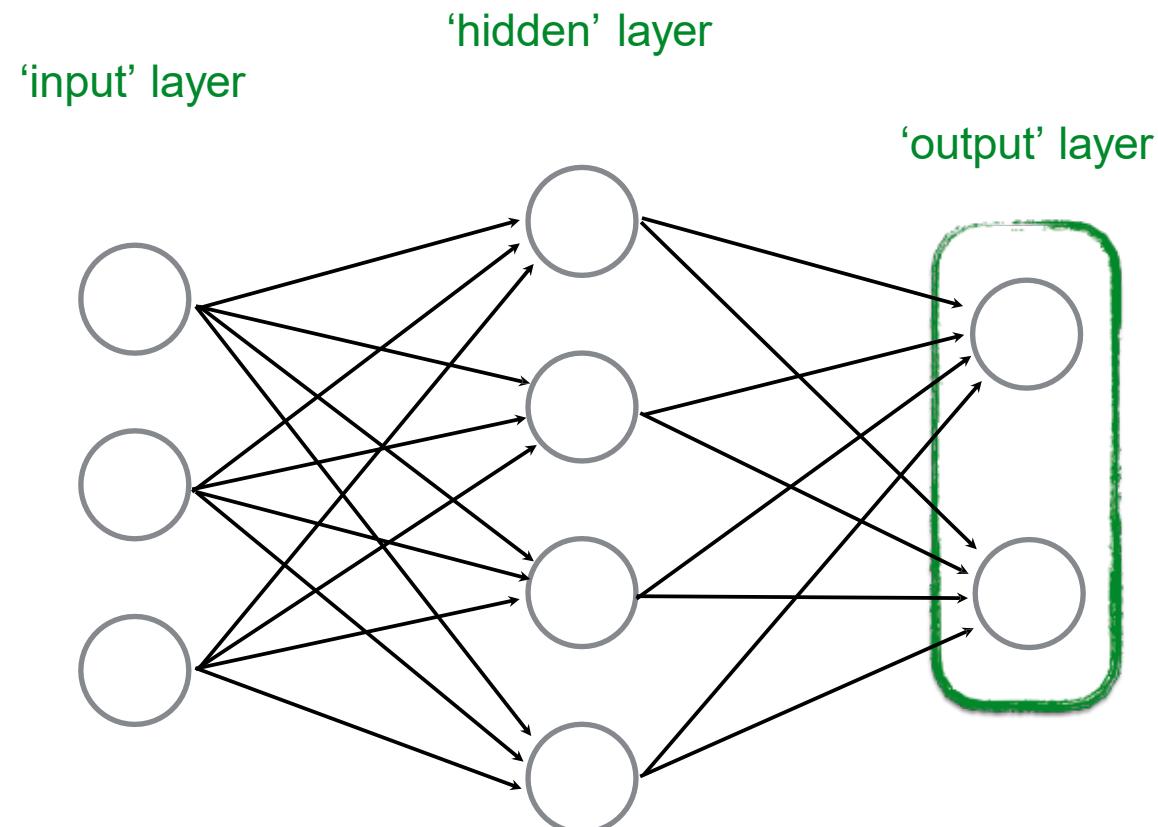
...also called a **Multi-layer Perceptron (MLP)**

# Some terminology...



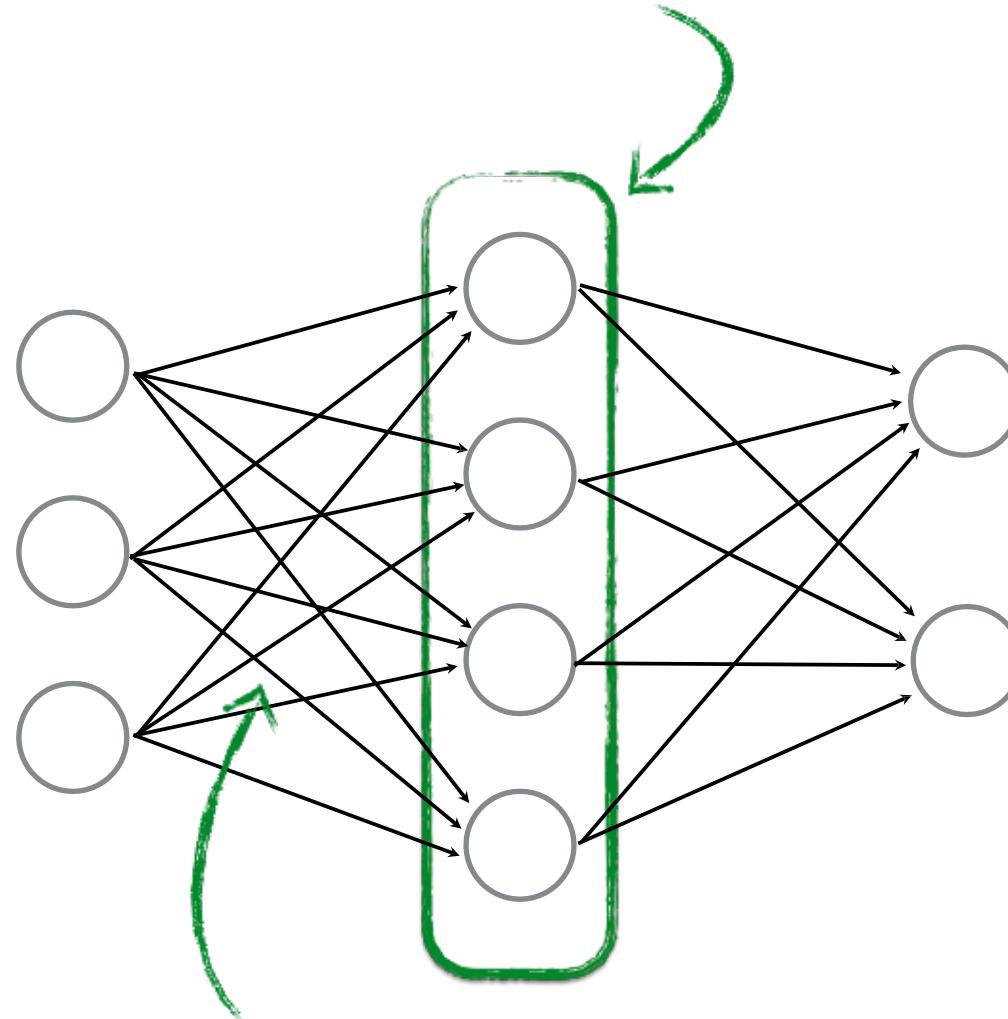
...also called a **Multi-layer Perceptron (MLP)**

# Some terminology...



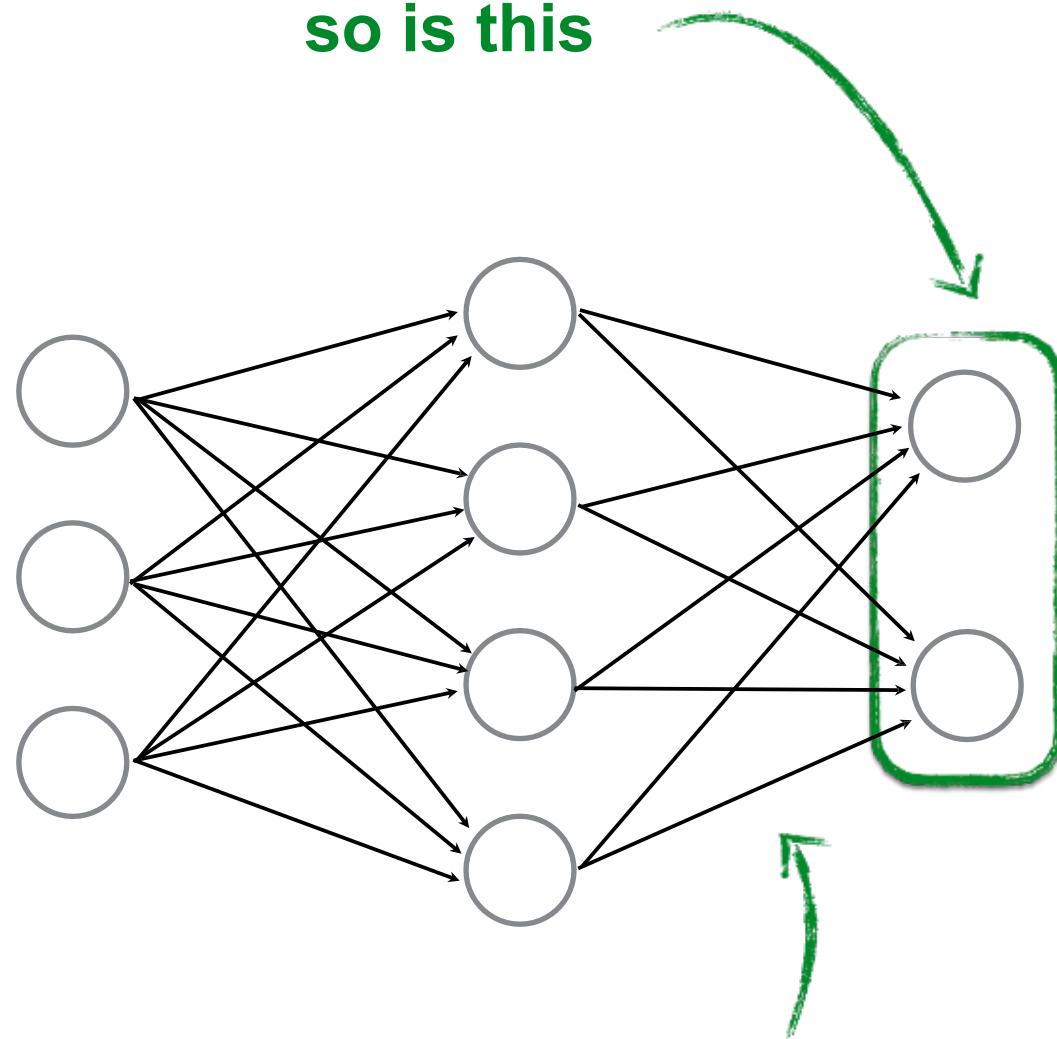
...also called a **Multi-layer Perceptron (MLP)**

this layer is a ‘fully connected layer’



all pairwise neurons between layers are connected

so is this

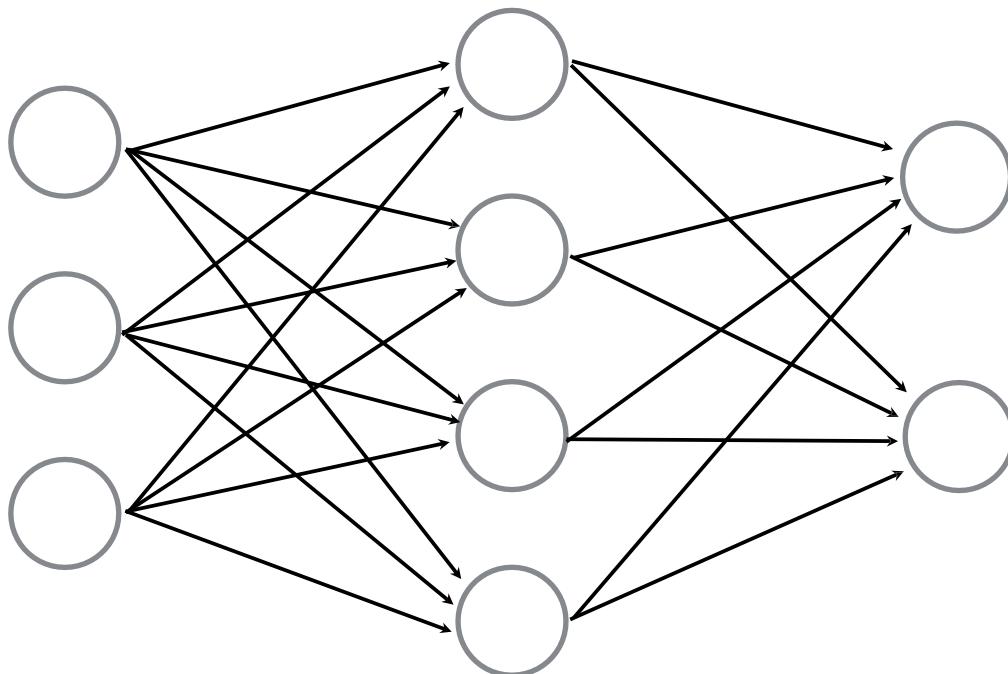


all pairwise neurons between layers are connected

# Neural Network

*How many neurons (perceptrons)?*

*How many weights (edges)?*



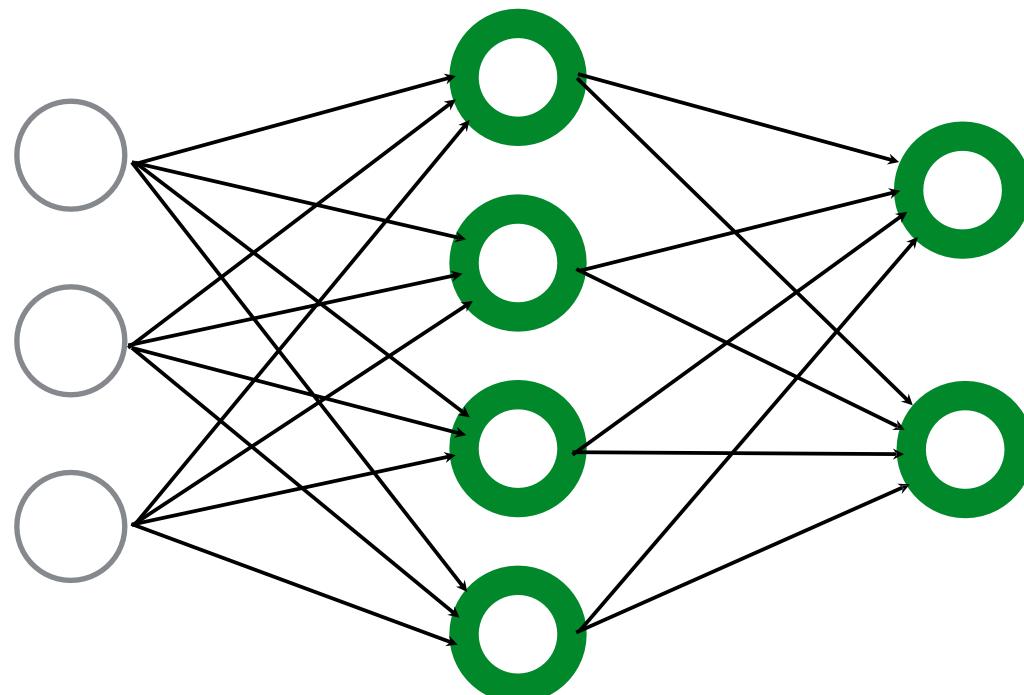
*How many learnable parameters total?*

# Neural Network

*How many neurons (perceptrons)?*

$$4 + 2 = 6$$

*How many weights (edges)?*



*How many learnable parameters total?*

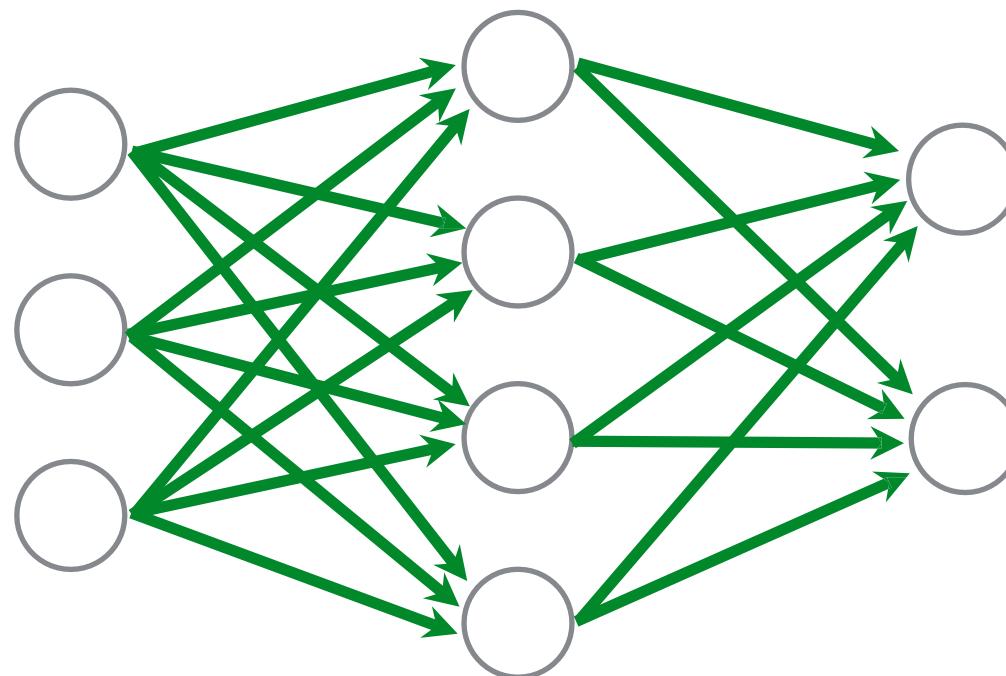
# Neural Network

*How many neurons (perceptrons)?*

$$4 + 2 = 6$$

*How many weights (edges)?*

$$(3 \times 4) + (4 \times 2) = 20$$



*How many learnable parameters total?*

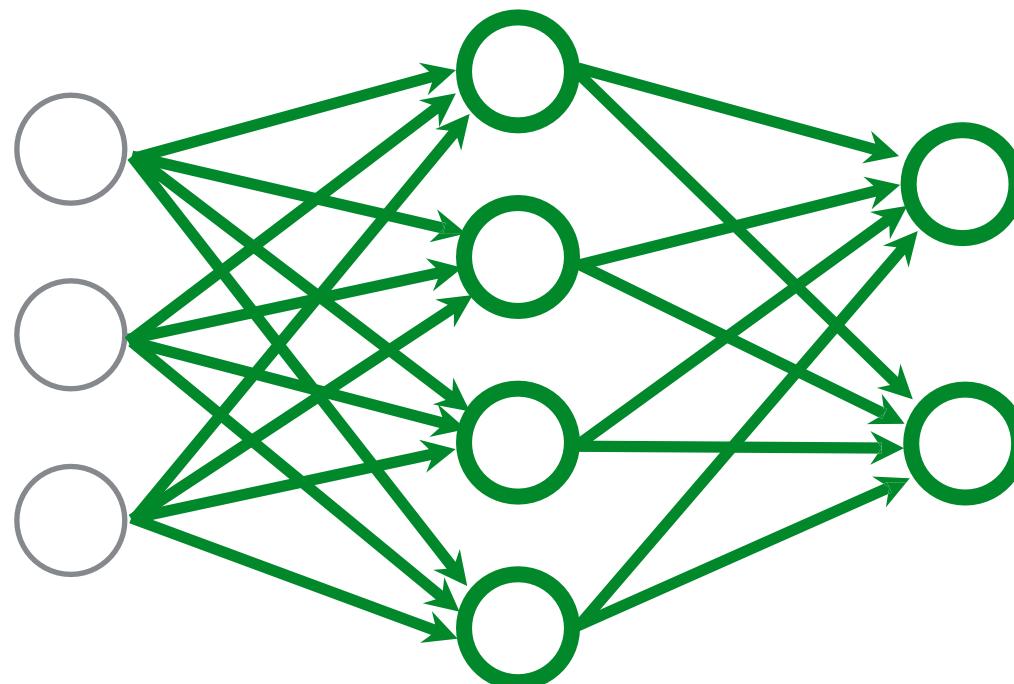
# Neural Network

*How many neurons (perceptrons)?*

$$4 + 2 = 6$$

*How many weights (edges)?*

$$(3 \times 4) + (4 \times 2) = 20$$



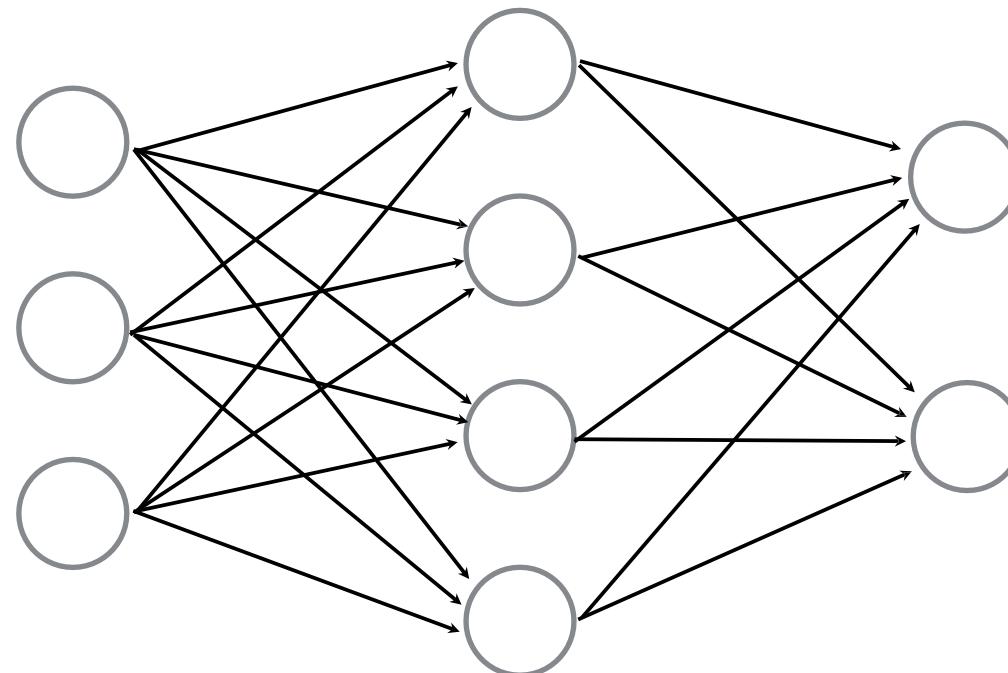
*How many learnable parameters total?*

$$20 + 4 + 2 = 26$$

bias terms

# Neural Network

performance usually tops out at 2-3 layers, deeper networks  
don't really improve performance...



...with the exception of **convolutional** networks for images