

# **Introduction to Computer Vision**

---

**Kaveh Fathian**


Assistant Professor

Computer Science Department

Colorado School of Mines


**Lecture 19**

# Training Perceptrons




# SVMs vs. Perceptrons

- What is the relation between **SVMs & perceptrons**?
- **SVMs** attempt to learn the support vectors which **maximize the margin between classes**




# SVMs vs. Perceptrons

- What is the relation between **SVMs & perceptrons**?
- **SVMs** attempt to learn the support vectors which **maximize the margin between classes**
- A **perceptron** does not




# World's Smallest Perceptron!



$$y = wx$$

What does this look like?

# World's Smallest Perceptron!



$$y = wx$$

(a.k.a. line equation, linear regression)

# Learning a Perceptron

Given a set of samples and a Perceptron

$$\{x_i, y_i\}$$

$$y = f_{\text{PER}}(x; w)$$

Estimate the parameters of the Perceptron

$$w$$

Given training data:

$x$	$y$
10	10.1
2	1.9
3.5	3.4
1	1.1

*What do you think the weight parameter is?*

$$y = wx$$

Given training data:

$x$	$y$
10	10.1
2	1.9
3.5	3.4
1	1.1

*What do you think the weight parameter is?*

$$y = wx$$

not so obvious as the network gets more complicate!

# An Incremental Learning Strategy

(gradient descent)

Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$\hat{y} = w\mathbf{x}$$

# An Incremental Learning Strategy

(gradient descent)

Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$\hat{y} = w x$$

Modify weight  $w$  such that  $\hat{y}$  gets ‘closer’ to  $y$

# An Incremental Learning Strategy

(gradient descent)

Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$\hat{y} = wx$$

Modify weight  $w$  such that  $\hat{y}$  gets ‘closer’ to  $y$

perceptron  
parameter

perceptron  
output

true  
label

# An Incremental Learning Strategy

(gradient descent)

Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$\hat{y} = wx$$

Modify weight  $w$  such that  $\hat{y}$  gets ‘closer’ to  $y$

perceptron  
parameter

perceptron  
output

*what does  
this mean?*

true  
label

Before diving into gradient descent, we need to understand ...


**Loss Function**  
defines what it means to be  
**close** to the true solution

**YOU get to chose the loss function!**  
(some are better than others depending on what you want to do)

# Loss Function


## Squared Error (L2)

(a popular loss function)




## L1 Loss

$$\ell(\hat{y}, y) = |\hat{y} - y|$$




## L2 Loss

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$




## Zero-One Loss

$$\ell(\hat{y}, y) = \mathbf{1}[\hat{y} = y]$$




## Hinge Loss

$$\ell(\hat{y}, y) = \max(0, 1 - y \cdot \hat{y})$$



back to the...

# World's Smallest Perceptron!



$$y = wx$$

(a.k.a. line equation, linear regression)

function of **ONE** parameter!

# Learning a Perceptron

Given a set of samples and a Perceptron

$$\{x_i, y_i\}$$

$$y = f_{\text{PER}}(x; w)$$

*what is this  
activation function?*

Estimate the parameter of the Perceptron

$$w$$

# Learning a Perceptron

Given a set of samples and a Perceptron

$$\{x_i, y_i\}$$

$$y = f_{\text{PER}}(x; w)$$

*what is this  
activation function?*  linear function!  $f(x) = wx$

Estimate the parameter of the Perceptron

$$w$$

# Learning Strategy

(gradient descent)

Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$\hat{y} = wx$$

Modify weight  $w$  such that  $\hat{y}$  gets ‘closer’ to  $y$

perceptron  
parameter

perceptron  
output

true  
label

Let's demystify this process first...

Code to train your perceptron:

Let's demystify this process first...

Code to train your perceptron:

for  $n = 1 \dots N$

$w = w + (y_n - \hat{y})x_i;$


just one line of code!

Now where does this come from?


# Gradient Descent

**(partial) derivatives** tell us how  
much one variable affects another

Two ways to think about them:




Slope of a function



Knobs on a machine

# 1. Slope of a function:


$$\frac{\partial f(\mathbf{x})}{\partial \mathbf{x}} = \left[ \frac{\partial f(\mathbf{x})}{\partial x}, \frac{\partial f(\mathbf{x})}{\partial y} \right] \quad \text{describes the slope around a point}$$

## 2. Knobs on a machine:




describes how each  
'knob' affects the  
output

$$\frac{\partial f(x)}{\partial w_1} \quad \frac{\partial f(x)}{\partial w_2} \quad \frac{\partial f(x)}{\partial w_3}$$


small change in parameter  $\Delta w_1$   $\rightarrow$  output will change by  $\frac{\partial f(x)}{\partial w_1} \Delta w_1$

# Gradient descent:



Given a  
fixed-point on a function,  
move in the direction  
opposite of the gradient


# Gradient descent:



# **Backpropagation**

back to the...

# World's Smallest Perceptron!



$$y = wx$$

(a.k.a. line equation, linear regression)

function of **ONE** parameter!

## Training the world's smallest perceptron

for  $n = 1 \dots N$

$$w = w + (y_n - \hat{y})x_i;$$

This is just gradient descent, that means...

this should be the gradient of the loss function

Now where does this come from?

$$\frac{d\mathcal{L}}{dw}$$

...is the rate at which **this** will change...

$$\mathcal{L} = \frac{1}{2}(y - \hat{y})^2$$

the loss function



... per unit change of **this**

$$y = w x$$

the weight parameter



Let's compute the derivative...

Compute the derivative

$$\begin{aligned}\frac{d\mathcal{L}}{dw} &= \frac{d}{dw} \left\{ \frac{1}{2}(y - \hat{y})^2 \right\} \\ &= -(y - \hat{y}) \frac{dwx}{dw} \\ &= -(y - \hat{y})x = \nabla w \quad \text{just shorthand}\end{aligned}$$

That means the weight update for **gradient descent** is:

$$\begin{aligned}w &= w - \nabla w \quad \text{move in direction of negative gradient} \\ &= w + (y - \hat{y})x\end{aligned}$$

## **Gradient Descent** (world's smallest perceptron)

For each sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

$$\hat{y} = wx_i$$

b. Compute Loss

$$\mathcal{L}_i = \frac{1}{2}(y_i - \hat{y})^2$$

2. Update

a. Back Propagation

$$\frac{d\mathcal{L}_i}{dw} = -(y_i - \hat{y})x_i = \nabla w$$

b. Gradient update


$$w = w - \nabla w$$

## Training the world's smallest perceptron

**for**  $n = 1 \dots N$

$$w = w + (y_n - \hat{y})x_i;$$

# World's (second) smallest **perceptron!**



function of **two** parameters!

# Gradient Descent

For each sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

b. Compute Loss

2. Update

a. Back Propagation

b. Gradient update

we just need to compute partial derivatives for this network



# Derivative computation

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_1} &= \frac{\partial}{\partial w_1} \left\{ \frac{1}{2}(y - \hat{y})^2 \right\} \\&= -(y - \hat{y}) \frac{\partial \hat{y}}{\partial w_1} \\&= -(y - \hat{y}) \frac{\partial \sum_i w_i x_i}{\partial w_1} \\&= -(y - \hat{y}) \frac{\partial w_1 x_1}{\partial w_1} \\&= -(y - \hat{y}) x_1 = \nabla w_1\end{aligned}$$

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_2} &= \frac{\partial}{\partial w_2} \left\{ \frac{1}{2}(y - \hat{y})^2 \right\} \\&= -(y - \hat{y}) \frac{\partial \hat{y}}{\partial w_2} \\&= -(y - \hat{y}) \frac{\partial \sum_i w_i x_i}{\partial w_2} \\&= -(y - \hat{y}) \frac{\partial w_2 x_2}{\partial w_2} \\&= -(y - \hat{y}) x_2 = \nabla w_2\end{aligned}$$

*Why do we have partial derivatives now?*

# Derivative computation

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_1} &= \frac{\partial}{\partial w_1} \left\{ \frac{1}{2}(y - \hat{y})^2 \right\} \\&= -(y - \hat{y}) \frac{\partial \hat{y}}{\partial w_1} \\&= -(y - \hat{y}) \frac{\partial \sum_i w_i x_i}{\partial w_1} \\&= -(y - \hat{y}) \frac{\partial w_1 x_1}{\partial w_1} \\&= -(y - \hat{y}) x_1 = \nabla w_1\end{aligned}$$

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_2} &= \frac{\partial}{\partial w_2} \left\{ \frac{1}{2}(y - \hat{y})^2 \right\} \\&= -(y - \hat{y}) \frac{\partial \hat{y}}{\partial w_2} \\&= -(y - \hat{y}) \frac{\partial \sum_i w_i x_i}{\partial w_2} \\&= -(y - \hat{y}) \frac{\partial w_2 x_2}{\partial w_2} \\&= -(y - \hat{y}) x_2 = \nabla w_2\end{aligned}$$

## Gradient Update

$$\begin{aligned}w_1 &= w_1 - \eta \nabla w_1 \\&= w_1 + \eta(y - \hat{y}) x_1\end{aligned}$$

$$\begin{aligned}w_2 &= w_2 - \eta \nabla w_2 \\&= w_2 + \eta(y - \hat{y}) x_2\end{aligned}$$

# Gradient Descent

For each sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass  $\hat{y} = f_{\text{MLP}}(x_i; \theta)$

b. Compute Loss  $\mathcal{L}_i = \frac{1}{2}(y_i - \hat{y})$  (side computation to track loss.  
not needed for backprop)

2. Update

a. Back Propagation

b. Gradient update

two lines now

$$\nabla w_{1i} = -(y_i - \hat{y})x_{1i}$$
$$\nabla w_{2i} = -(y_i - \hat{y})x_{2i}$$


$$w_{1i} = w_{1i} + \eta(y - \hat{y})x_{1i}$$
$$w_{2i} = w_{2i} + \eta(y - \hat{y})x_{2i}$$

(adjustable step size)





We haven't seen a lot of 'propagation' yet  
because our perceptrons only had one layer...


# Multi-Layer Perceptron




function of **FOUR** parameters and **FOUR** layers!








$$a_1 = w_1 \cdot x + b_1$$




$$a_1 = w_1 \cdot x + b_1$$




$$a_1 = w_1 \cdot x + b_1$$

$$a_2 = w_2 \cdot f_1(w_1 \cdot x + b_1)$$



$$a_1 = w_1 \cdot x + b_1$$


$$a_2 = w_2 \cdot f_1(w_1 \cdot x + b_1)$$



$$a_1 = w_1 \cdot x + b_1$$

$$a_2 = w_2 \cdot f_1(w_1 \cdot x + b_1)$$


$$a_3 = w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1))$$



$$a_1 = w_1 \cdot x + b_1$$

$$a_2 = w_2 \cdot f_1(w_1 \cdot x + b_1)$$

$$a_3 = w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1))$$



$$a_1 = w_1 \cdot x + b_1$$

$$a_2 = w_2 \cdot f_1(w_1 \cdot x + b_1)$$

$$a_3 = w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1))$$

$$y = f_3(w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1)))$$

Entire network can be written out as one long equation

$$y = f_3(w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1)))$$

We need to train the network:

*What is known? What is unknown?*

Entire network can be written out as one long equation

$$y = f_3(w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1)))$$



We need to train the network:

*What is known? What is unknown?*

Entire network can be written out as one long equation

$$y = f_3(w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1)))$$



We need to train the network:

*What is known? What is unknown?*

# Learning an MLP (Multi-Layer Perceptron)

Given a set of samples and a MLP

$$\{x_i, y_i\}$$

$$y = f_{\text{MLP}}(x; \theta)$$

Estimate the parameters of the MLP

$$\theta = \{f, w, b\}$$

# Gradient Descent

For each **random** sample  $\{x_i, y_i\}$

1. Predict

a. Forward pass

$$\hat{y} = f_{\text{MLP}}(x_i; \theta)$$

b. Compute Loss

2. Update

a. Back Propagation

$$\frac{\partial \mathcal{L}}{\partial \theta}$$

vector of parameter partial derivatives


b. Gradient update

$$\theta \leftarrow \theta - \eta \nabla \theta$$

vector of parameter update equations


So we need to compute the partial derivatives

$$\frac{\partial \mathcal{L}}{\partial \theta} = \left[ \frac{\partial \mathcal{L}}{\partial w_3} \frac{\partial \mathcal{L}}{\partial w_2} \frac{\partial \mathcal{L}}{\partial w_1} \frac{\partial \mathcal{L}}{\partial b} \right]$$




Remember,

Partial derivative  $\frac{\partial L}{\partial w_1}$  describes...



So, how do you compute it?


# The Chain Rule




According to the chain rule...

$$\frac{\partial L}{\partial w_3} = \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}$$

Intuitively, the effect of weight on loss function :  $\frac{\partial L}{\partial w_3}$






$$L(y, \hat{y}) = \frac{\eta}{2} (y - \hat{y})^2$$

$$\frac{\partial L}{\partial w_3} = \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}$$

Chain Rule!




$$L(y, \hat{y}) = \frac{\eta}{2} (y - \hat{y})^2$$

$$\begin{aligned}\frac{\partial L}{\partial w_3} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\ &= -\eta(y - \hat{y}) \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}\end{aligned}$$


Just the partial  
derivative of L2 loss






$$L(y, \hat{y}) = \frac{\eta}{2} (y - \hat{y})^2$$

$$\begin{aligned}\frac{\partial L}{\partial w_3} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\ &= -\eta(y - \hat{y}) \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}\end{aligned}$$




Let's use a Sigmoid function

$$\frac{ds(x)}{dx} = s(x)(1 - s(x))$$




$$s(x) = \frac{1}{1 + e^{-x}}$$




$$L(y, \hat{y}) = \frac{\eta}{2} (y - \hat{y})^2$$

$$\begin{aligned}\frac{\partial L}{\partial w_3} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\ &= -\eta(y - \hat{y}) \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\ &= -\eta(y - \hat{y}) f_3(1 - f_3) \frac{\partial a_3}{\partial w_3}\end{aligned}$$




Let's use a Sigmoid function

$$\frac{ds(x)}{dx} = s(x)(1 - s(x))$$




$$s(x) = \frac{1}{1 + e^{-x}}$$




$$L(y, \hat{y}) = \frac{\eta}{2} (y - \hat{y})^2$$

$$\begin{aligned}
 \frac{\partial L}{\partial w_3} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\
 &= -\eta(y - \hat{y}) \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\
 &= -\eta(y - \hat{y}) f_3(1 - f_3) \frac{\partial a_3}{\partial w_3} \\
 &= -\eta(y - \hat{y}) f_3(1 - f_3) f_2
 \end{aligned}$$




$$\frac{\partial L}{\partial w_2} = \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial w_2}$$



$$\frac{\partial L}{\partial w_2} = \boxed{\frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial w_2}}$$


already computed.  
re-use (propagate)!

# The Chain Rule




a.k.a. backpropagation

The chain rule says...




$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial w_1}$$

The chain rule says...



$$\frac{\partial L}{\partial w_1} = \boxed{\frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2}} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial w_1}$$

already computed.  
re-use (propagate)!




$$\frac{\partial \mathcal{L}}{\partial w_3} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}$$

$$\frac{\partial \mathcal{L}}{\partial w_2} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial w_2}$$

$$\frac{\partial \mathcal{L}}{\partial w_1} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial w_1}$$

$$\frac{\partial \mathcal{L}}{\partial b} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial b}$$




$$\frac{\partial \mathcal{L}}{\partial w_3} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}$$

$$\frac{\partial \mathcal{L}}{\partial w_2} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial w_2}$$

$$\frac{\partial \mathcal{L}}{\partial w_1} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial w_1}$$

$$\frac{\partial \mathcal{L}}{\partial b} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial b}$$



$$\frac{\partial \mathcal{L}}{\partial w_3} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}$$

$$\frac{\partial \mathcal{L}}{\partial w_2} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial w_2}$$

$$\frac{\partial \mathcal{L}}{\partial w_1} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial w_1}$$

$$\frac{\partial \mathcal{L}}{\partial b} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial b}$$

# Gradient Descent

For each example sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

$$\hat{y} = f_{\text{MLP}}(x_i; \theta)$$

b. Compute Loss

$$\mathcal{L}_i$$

2. Update

a. Back Propagation

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_3} &= \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\ \frac{\partial \mathcal{L}}{\partial w_2} &= \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial w_2} \\ \frac{\partial \mathcal{L}}{\partial w_1} &= \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial w_1} \\ \frac{\partial \mathcal{L}}{\partial b} &= \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial b}\end{aligned}$$

$$w_3 = w_3 - \eta \nabla w_3$$

$$w_2 = w_2 - \eta \nabla w_2$$

$$w_1 = w_1 - \eta \nabla w_1$$

$$b = b - \eta \nabla b$$

# Gradient Descent

For each example sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

$$\hat{y} = f_{\text{MLP}}(x_i; \theta)$$

b. Compute Loss

$$\mathcal{L}_i$$

2. Update

a. Back Propagation

$$\frac{\partial \mathcal{L}}{\partial \theta}$$

vector of parameter partial derivatives

b. Gradient update

$$\theta \leftarrow \theta + \eta \frac{\partial \mathcal{L}}{\partial \theta}$$

vector of parameter update equations

# **Stochastic Gradient Descent**

# What we are truly minimizing:

$$\min_{\theta} \sum_{i=1}^N L(y_i, f_{MLP}(x_i))$$

# **What we are truly minimizing:**

$$\min_{\theta} \sum_{i=1}^N L(y_i, f_{MLP}(x_i))$$

**The gradient is:**

$$\sum_{i=1}^N \frac{\partial L(y_i, f_{MLP}(x_i))}{\partial \theta}$$

# **What we are truly minimizing:**

$$\min_{\theta} \sum_{i=1}^N L(y_i, f_{MLP}(x_i))$$

## **The gradient is:**

$$\sum_{i=1}^N \frac{\partial L(y_i, f_{MLP}(x_i))}{\partial \theta}$$

## **What we use for gradient update is:**

$$\frac{\partial L(y_i, f_{MLP}(x_i))}{\partial \theta} \quad \text{for some } i$$

# Stochastic Gradient Descent

- For each example sample  $\{x_i, y_i\}$

1. Predict  $\hat{y} = f_{\text{MLP}}(x_i; \theta)$

a. Forward pass  $\mathcal{L}_i$

b. Compute Loss

2. Update

a. Back Propagation  
b. Gradient update

$$\frac{\partial \mathcal{L}}{\partial \theta}$$

vector of parameter partial derivatives

$$\theta \leftarrow \theta + \eta \frac{\partial \mathcal{L}}{\partial \theta}$$

vector of parameter update equations

# **How do we select which sample?**

# **How do we select which sample?**

- Select randomly!

## **Do we need to use only one sample?**

## **How do we select which sample?**

- Select randomly!

## **Do we need to use only one sample?**

- You can use a *minibatch* of size  $B < N$ .

## **Why not do gradient descent with all samples?**

## **How do we select which sample?**

- Select randomly!

## **Do we need to use only one sample?**

- You can use a *minibatch* of size  $B < N$ .

## **Why not do gradient descent with all samples?**

- It's very expensive when  $N$  is large (big data).

## **Do I lose anything by using stochastic GD?**

How do we select which


- Select randomly!

Do we need to use only

- You can use a *minibatch*

Why not do gradient des


- It's very expensive whe




## Do I lose anything by using stochastic GD?

- Same convergence guarantees and complexity!
- Better generalization

# Notes on Optimization



# Learning rates



$$\theta \leftarrow \theta \boxed{-} \eta \frac{\partial \mathcal{L}}{\partial \theta}$$

Step size: learning rate  
Too big: will miss the minimum  
Too small: slow convergence

# Learning rate scheduling

- Use different **learning rate** at each iteration
- Most common choice:

$$\eta_t = \frac{\eta_0}{\sqrt{t}}$$

- Need to select initial learning rate  $\eta_0$ , important!!!
- More modern choice: **Adaptive** learning rates


$$\eta_t = G \left( \left\{ \frac{\partial L}{\partial \theta} \right\}_{i=0}^t \right)$$

- Many choices for G (**Adam**, **Adagrad**, **Adadelta**)

# Momentum Update


$$\theta \leftarrow \theta - \eta \frac{\partial \mathcal{L}}{\partial \theta}$$

$$\Delta\theta \leftarrow w \frac{\partial L}{\partial \theta} + (1 - w)\Delta\theta$$




Take direction history into account!

```
weights_grad = evaluate_gradient(loss_fun, data, weights)
vel = vel * 0.9 - step_size * weights_grad
weights += vel
```



(Fig. 2a)



(Fig. 2b)

## Many other ways to perform optimization...

- Second order methods that use the Hessian (or its approximation): BFGS, **LBFGS**, etc.
- Currently, the lesson from the trenches is that well-tuned SGD+Momentum is very hard to beat for CNNs.
- No consensus on Adam etc.: Seem to give faster performance to worse local minima

# Derivatives

- Given  $f(x)$ , where  $x$  is vector of inputs
  - Compute gradient of  $f$  at  $x$ :  $\nabla f(x)$

How do we do differentiation?

# Numerical differentiation

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

$$f(x + h) = f(x) + h \frac{df(x)}{dx}$$

# Numerical differentiation

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

$$f(x + h) = f(x) + h \frac{df(x)}{dx}$$

Numerical differentiation is:

- Approximate
- Slow
- Numerically unstable
- Easy to write

# Symbolic differentiation

- What Mathematica does: Automatically derive *analytical* expressions for derivative.

# Symbolic differentiation

- What Mathematica does: Automatically derive *analytical* expressions for derivative.
- Often results in very redundant (and expensive to evaluate) expressions.

```
D[Log[1 + Exp[w*x + b]], w]
```

$$\text{Out}[11]= \frac{e^{b+w x} w}{1 + e^{b+w x}}$$

```
In[19]:= D[Log[1 + Exp[w2 * Log[1 + Exp[w1 * x + b1]] + b2]], w1]
```

$$\text{Out}[19]= \frac{e^{b1+b2+w1 x+w2 \operatorname{Log}\left[1+e^{b1+w1 x}\right]} w2 x}{\left(1+e^{b1+w1 x}\right) \left(1+e^{b2+w2 \operatorname{Log}\left[1+e^{b1+w1 x}\right]}\right)}$$

- Often intractable.

# Automatic differentiation (autodiff)

- An autodiff system will convert the program into a sequence of **primitive operations** which have specified routines for computing derivatives.
- In this representation, backprop can be done in a completely mechanical way.

**Sequence of primitive operations:**

**Original program:**

$$z = wx + b$$

$$y = \frac{1}{1 + \exp(-z)}$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$

$$t_1 = wx$$

$$z = t_1 + b$$

$$t_3 = -z$$

$$t_4 = \exp(t_3)$$

$$t_5 = 1 + t_4$$

$$y = 1/t_5$$

$$t_6 = y - t$$

$$t_7 = t_6^2$$

$$\mathcal{L} = t_7/2$$

# In summary

- Numerical gradient: easy to implement, bad to use.
- Symbolic gradient: sometimes useful, often intractable.
- Automatic gradient: exact, fast, error-prone.

In practice: Use symbolic gradient for small/trivial programs.  
Almost always use analytic gradient, but check correctness of implementation with numerical gradient.

- This is called a gradient check.